Big Data and Hadoop

What is Big Data

"Big data" is a field that treats ways to analyze, systematically extract information from, or otherwise deal with data sets that are too large or complex to be dealt with by traditional data-processing application software

-- Wiki

Big Data is a term used to describe a collection of data that is huge in size and yet growing exponentially with time.

Difference Between RDBMS and Hadoop and

Feature	RDBMS	Hadoop
Data Variety	Mainly for Structured data.	Used for Structured, Semi-Structured and Unstructured data
Data Storage	Average size data (GBS)	Use for large data set (Tbs and Pbs)
Querying	SQL Language	HQL (Hive Query Language)
Schema	Required on write (static schema)	Required on reading (dynamic schema)
Speed	Reads are fast	Both reads and writes are fast
Cost	License	Free
Use Case	OLTP (Online transaction processing)	Analytics (Audio, video, logs etc), Data Discovery
Data Objects	Works on Relational Tables	Works on Key/Value Pair
Throughput	Low	High
Scalability	Vertical	Horizontal
Hardware Profile	High-End Servers	Commodity/Utility Hardware
Integrity	High (ACID)	Low

Attributes of Big Data

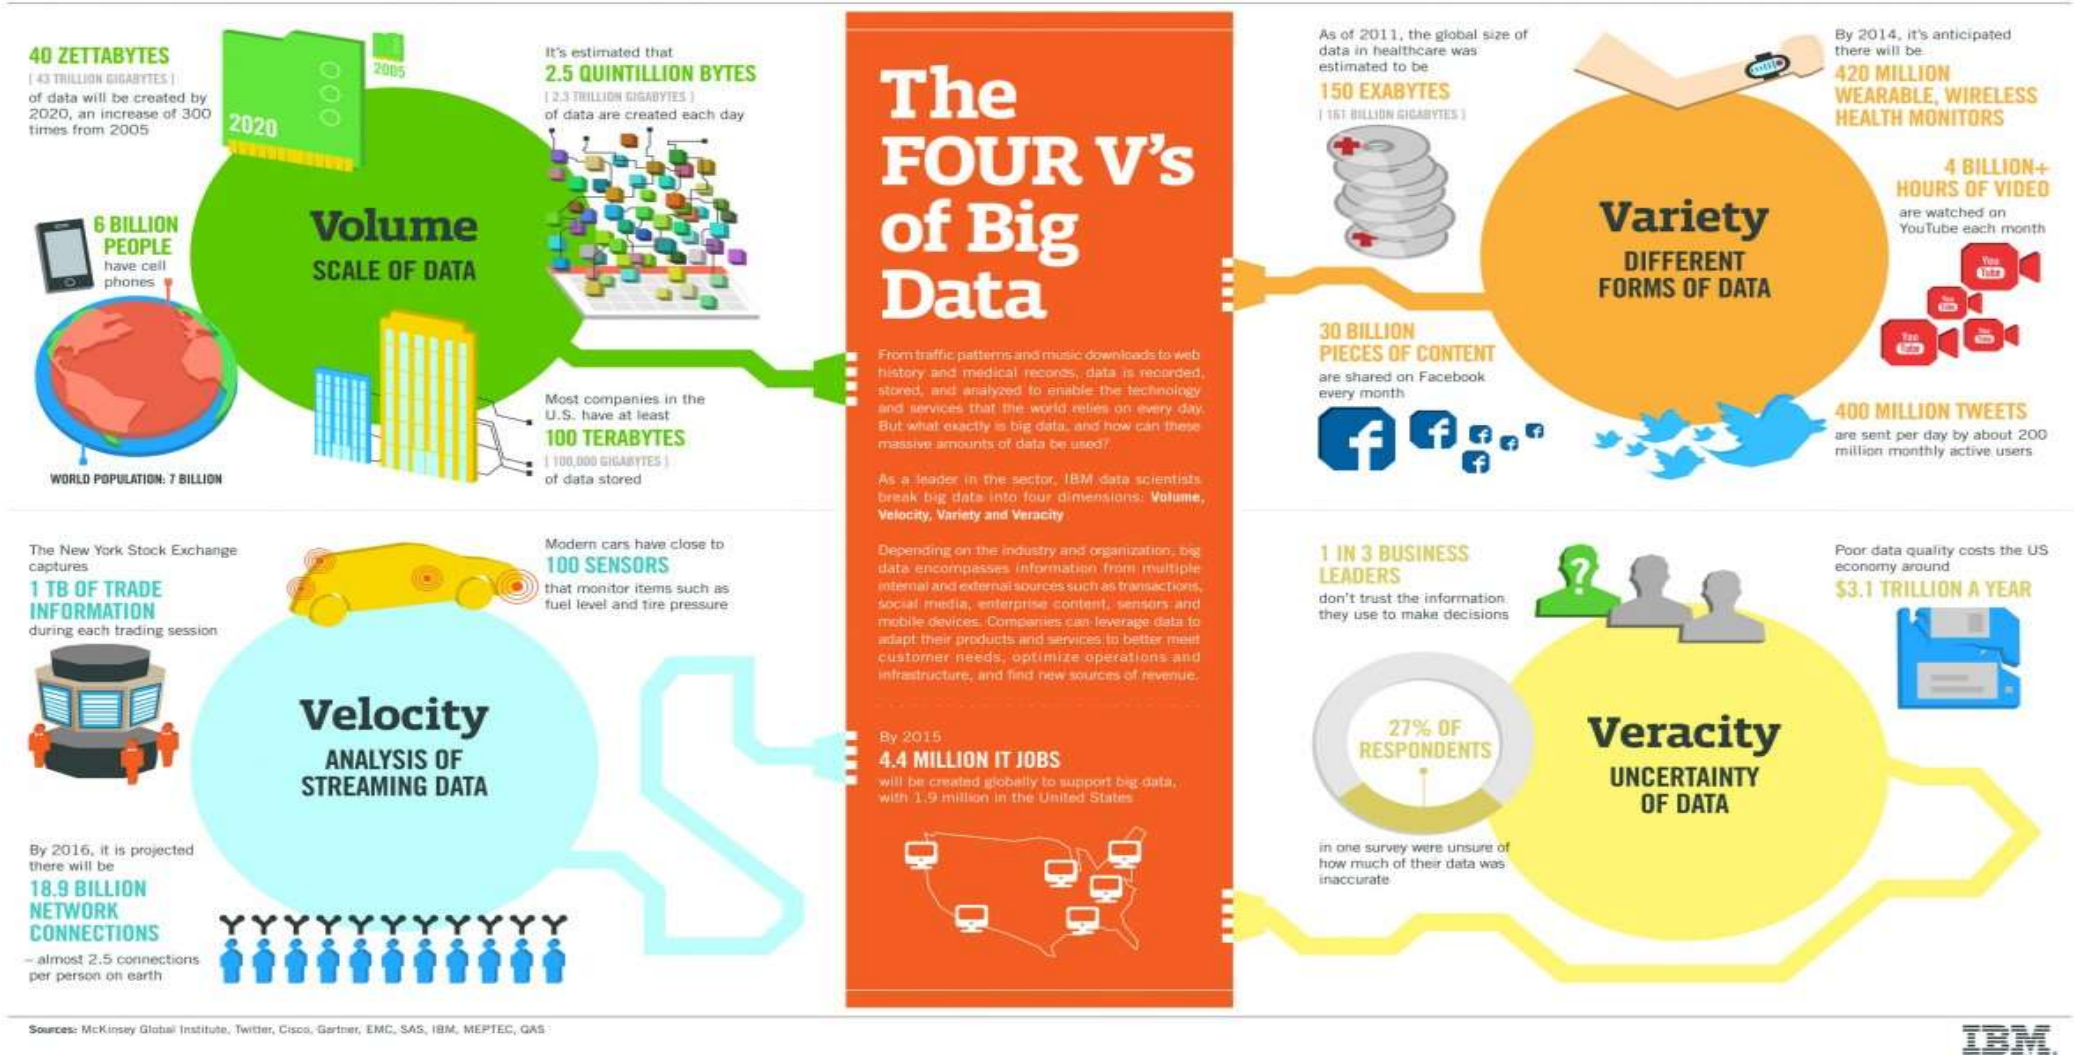

Applications of Big Data

NetApp Improves Customer Support by Deploying Cloudera Enterprise

Home > Government IT

Opinion

Barack Obama's Big Data won the US election

Joint Success Story: Major Retail Bank

CASE STUDY

Intel® Xeon® Processors

Intel® Distribution for Apache Hadoop* Software

Streamlining Healthcare Connectivity with Big Data

China Mobile Guangdong Gives Subscribers Real-Time Access to Billing and Call Data Records

What is Hadoop

- Big Data is a *Business Term* and *Hadoop* is the framework for implementation

- Hadoop
  - is an open-source software framework
  - used for distributed storage and
  - distributed processing of dataset of big data using the MapReduce programming model.
  - It consists of computer clusters built from commodity hardware

History of Hadoop

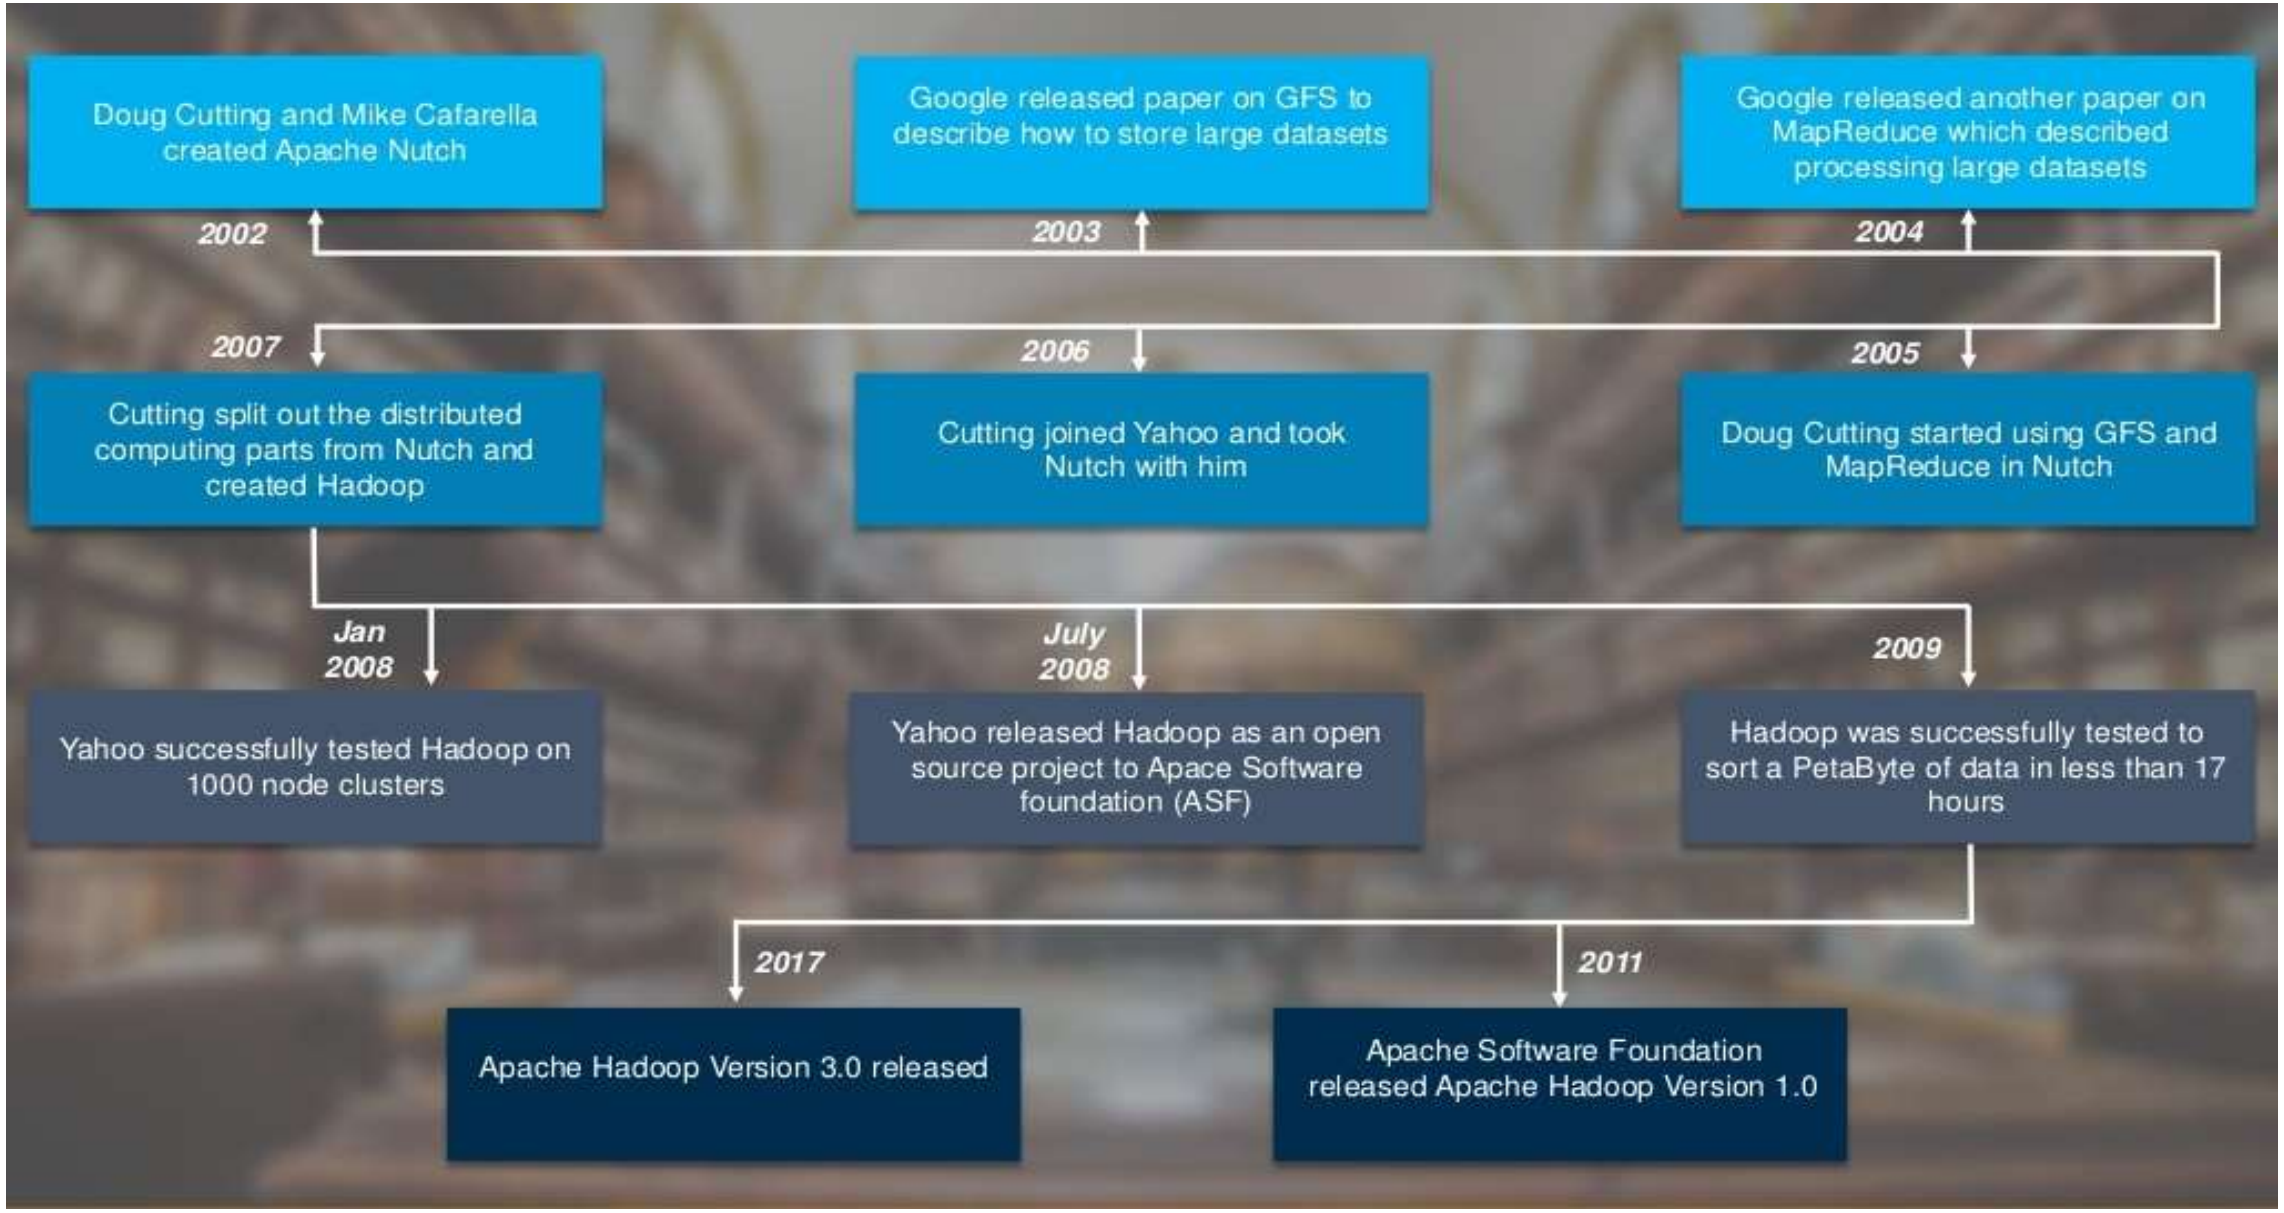

Hadoop 1.X Architecture

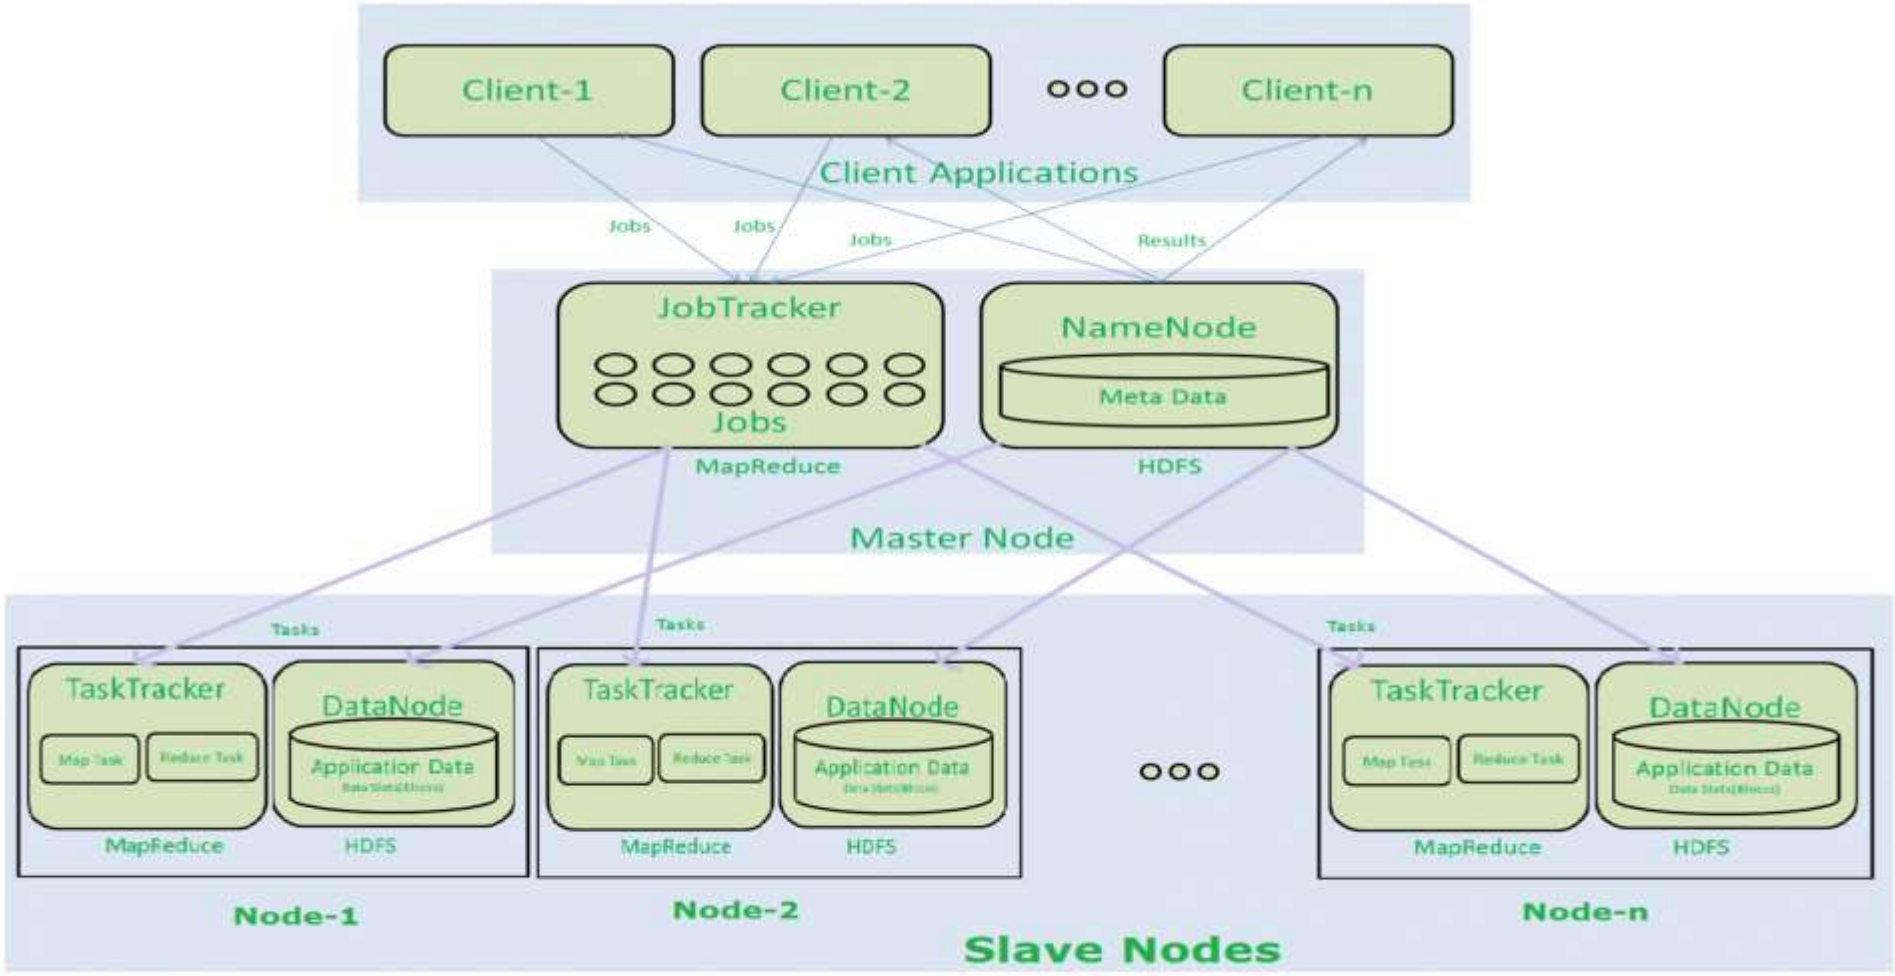

Hadoop 1.x Architecture

Hadoop Distributors

Difference between Hadoop 1 and Hadp2

Hadoop1	Hadoop2
Supports MapReduce (MR) processing model only. Does not support non-MR tools	Allows to work in MR as well as other distributed computing models like Spark, Hama, Giraph, Message Passing Interface) MPI & HBase coprocessors.
MR does both processing and cluster-resource management.	YARN (Yet Another Resource Negotiator) does cluster resource management and processing is done using different processing models.
Has limited scaling of nodes. Limited to 4000 nodes per cluster	Has better scalability. Scalable up to 10000 nodes per cluster
Works on concepts of slots – slots can run either a Map task or a Reduce task only.	Works on concepts of containers. Using containers can run generic tasks.
A single Namenode to manage the entire namespace.	Multiple Namenode servers manage multiple namespaces.
Has Single-Point-of-Failure (SPOF) – because of single Namenode- and in the case of Namenode failure, needs manual intervention to overcome.	Has to feature to overcome SPOF with a standby Namenode and in the case of Namenode failure, it is configured for automatic recovery.
Does not support Microsoft Windows	Added support for Microsoft windows

More on Master Slave

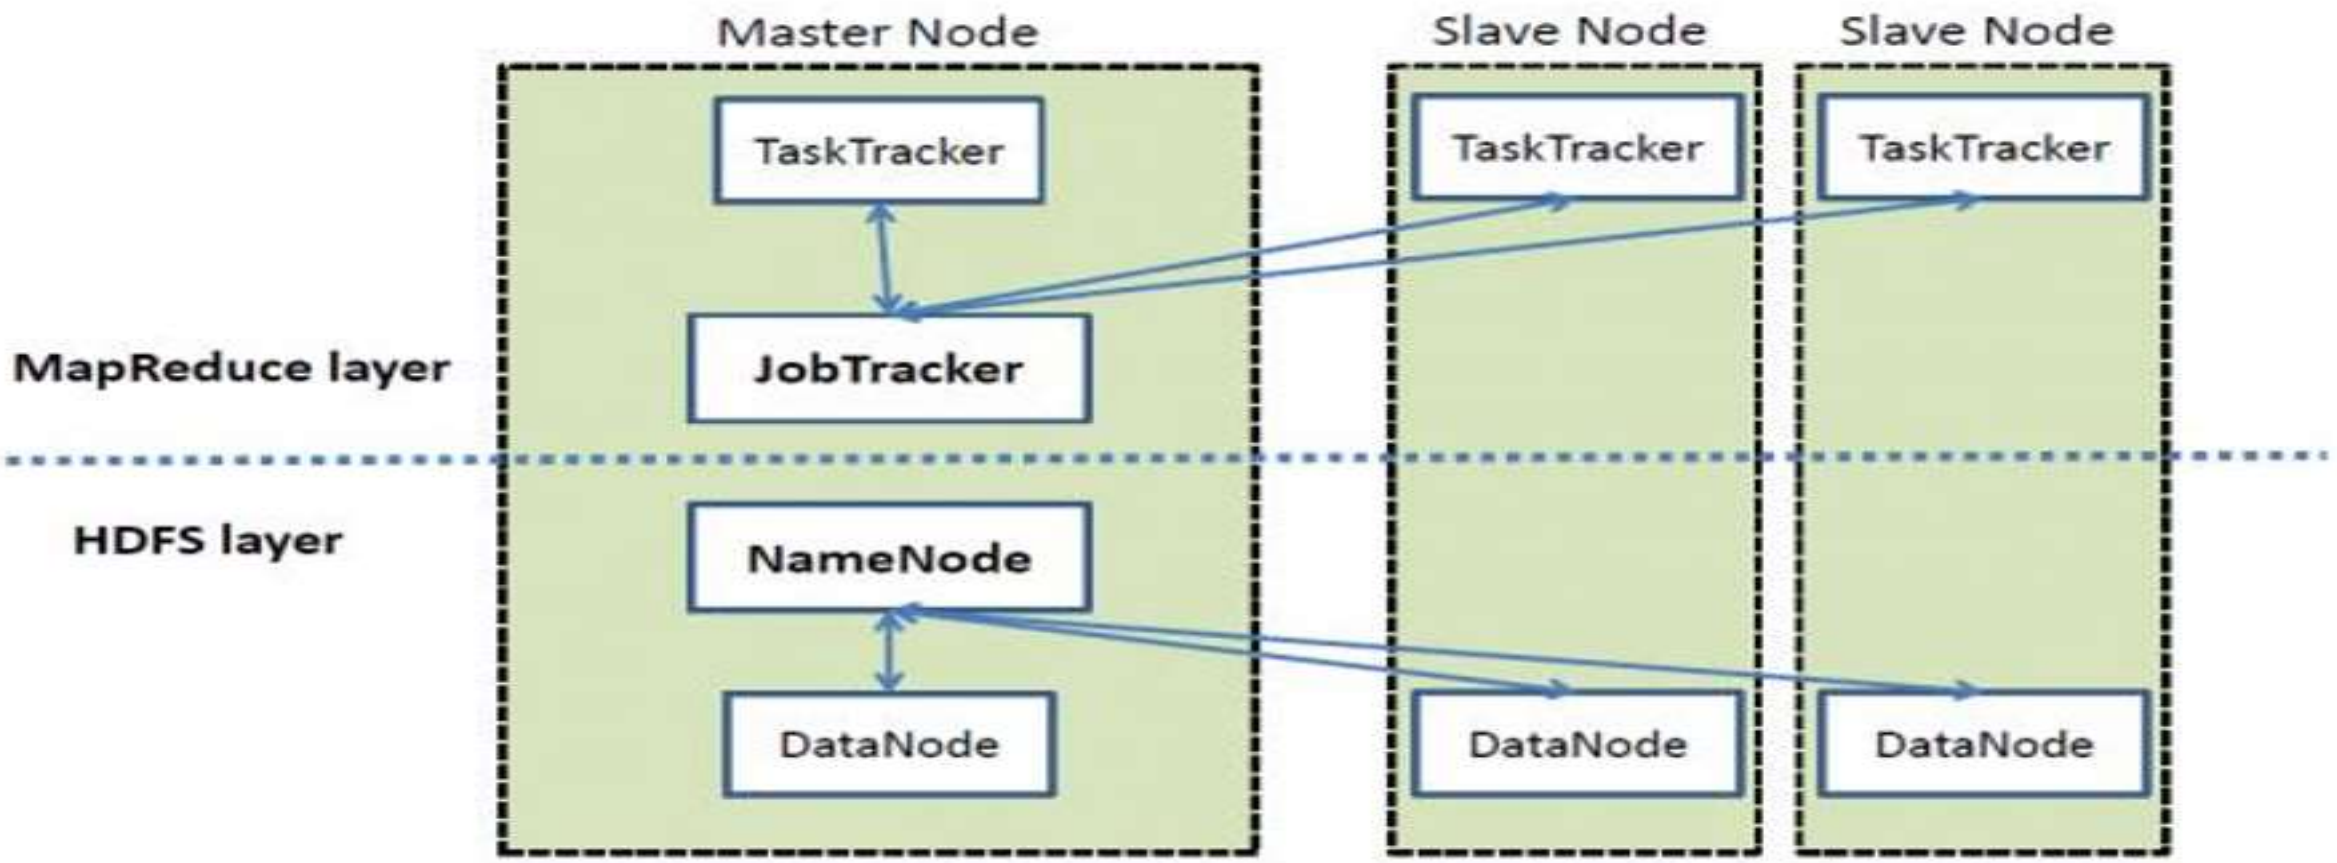

Hadoop 2.X Architecture

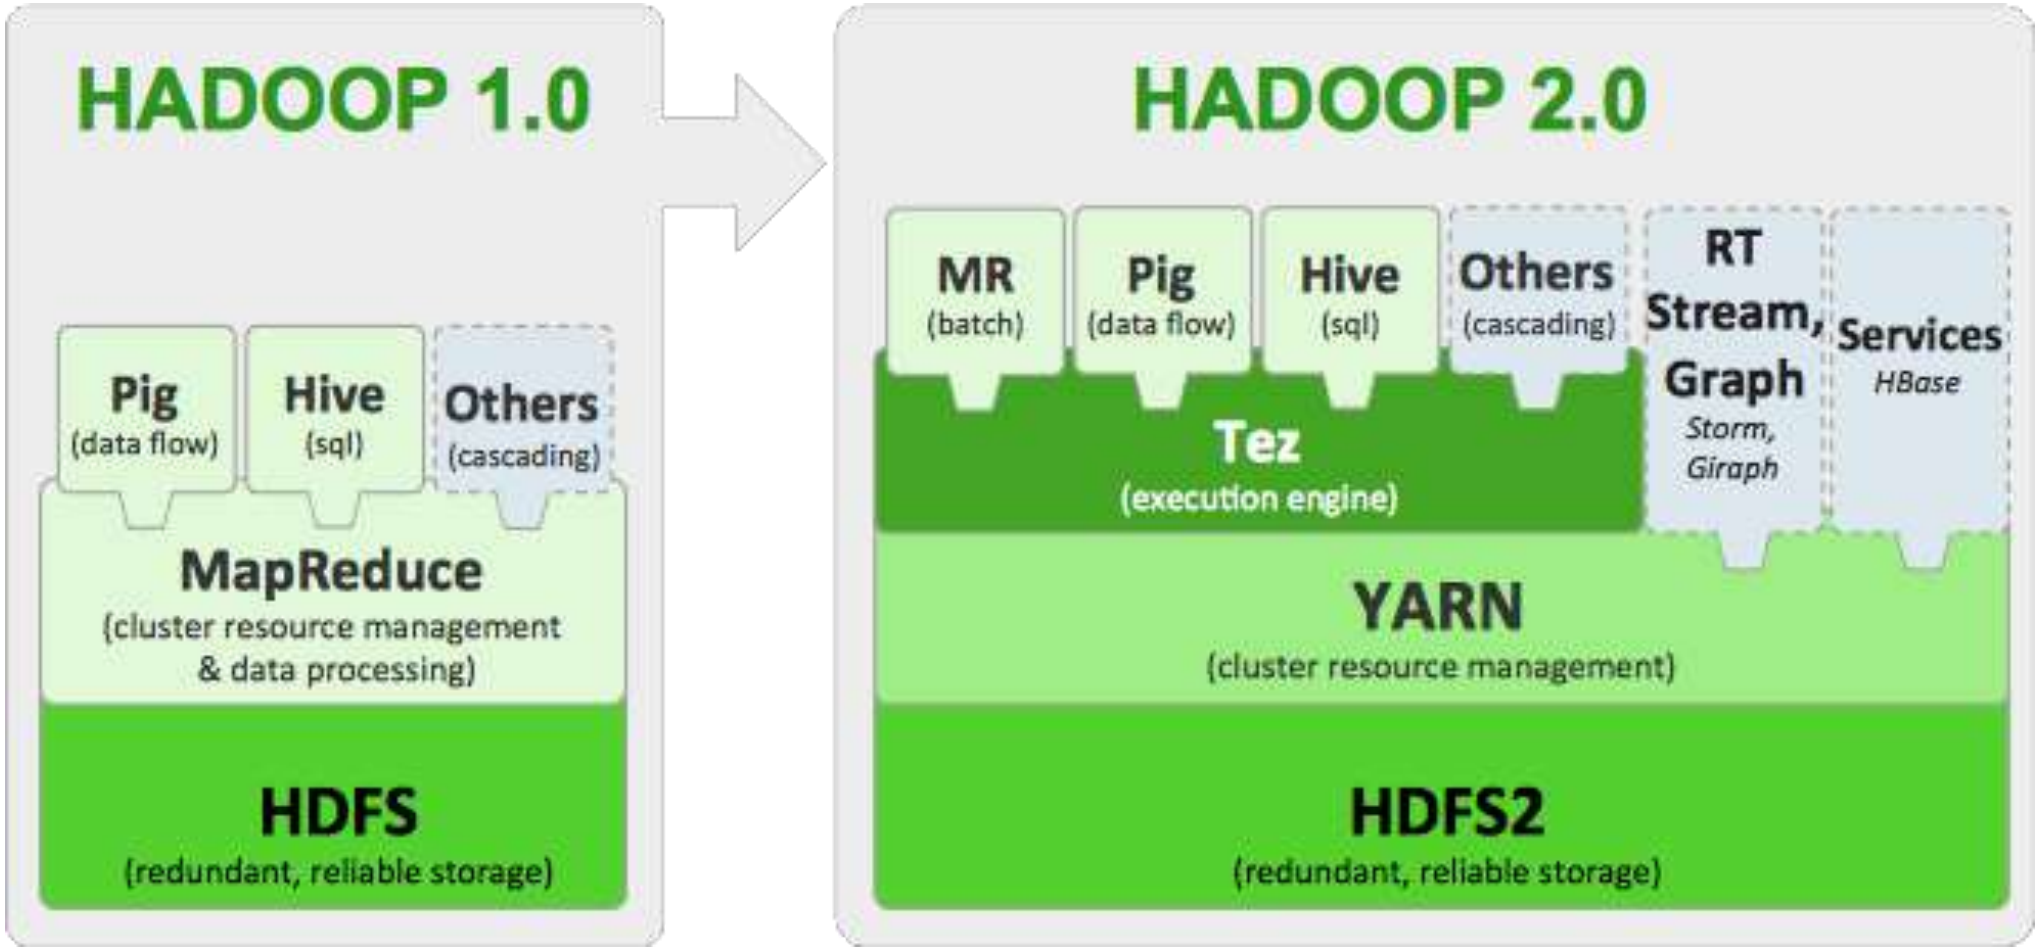

Hadoop 2.X Architecture Components

- **NameNode(HA) aka HDFS High Availability:** In Hadoop 1.0 NameNode was the single point of failure in a Cluster, resulting in data loss in case of a NameNode failure. Hadoop 2.0 Architecture supports multiple NameNodes to remove this bottleneck by using Standby NameNode.
- **HDFS :** enables support for multiple namespaces in the cluster to improve scalability and isolation
- **YARN(Yet Another Resource Negotiator) aka NextGen (MRv2):** This is next generation processing framework.

YARN Architecture

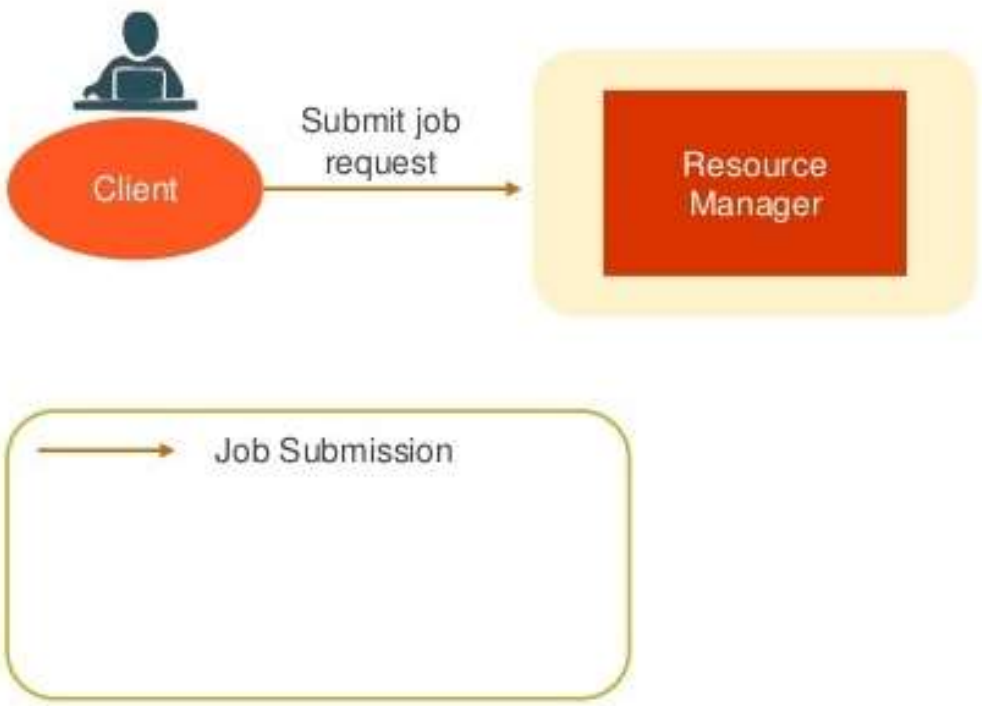

YARN Architecture

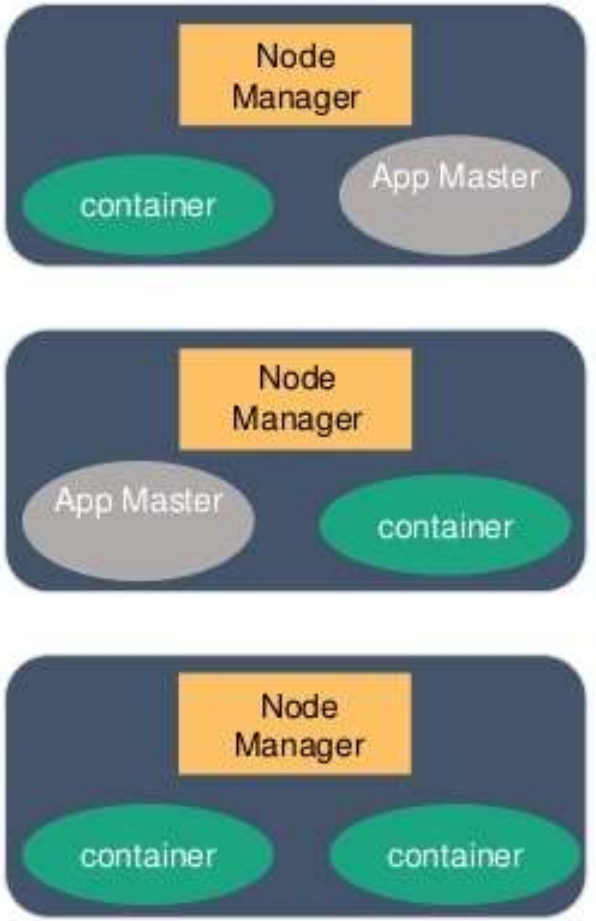

YARN Architecture

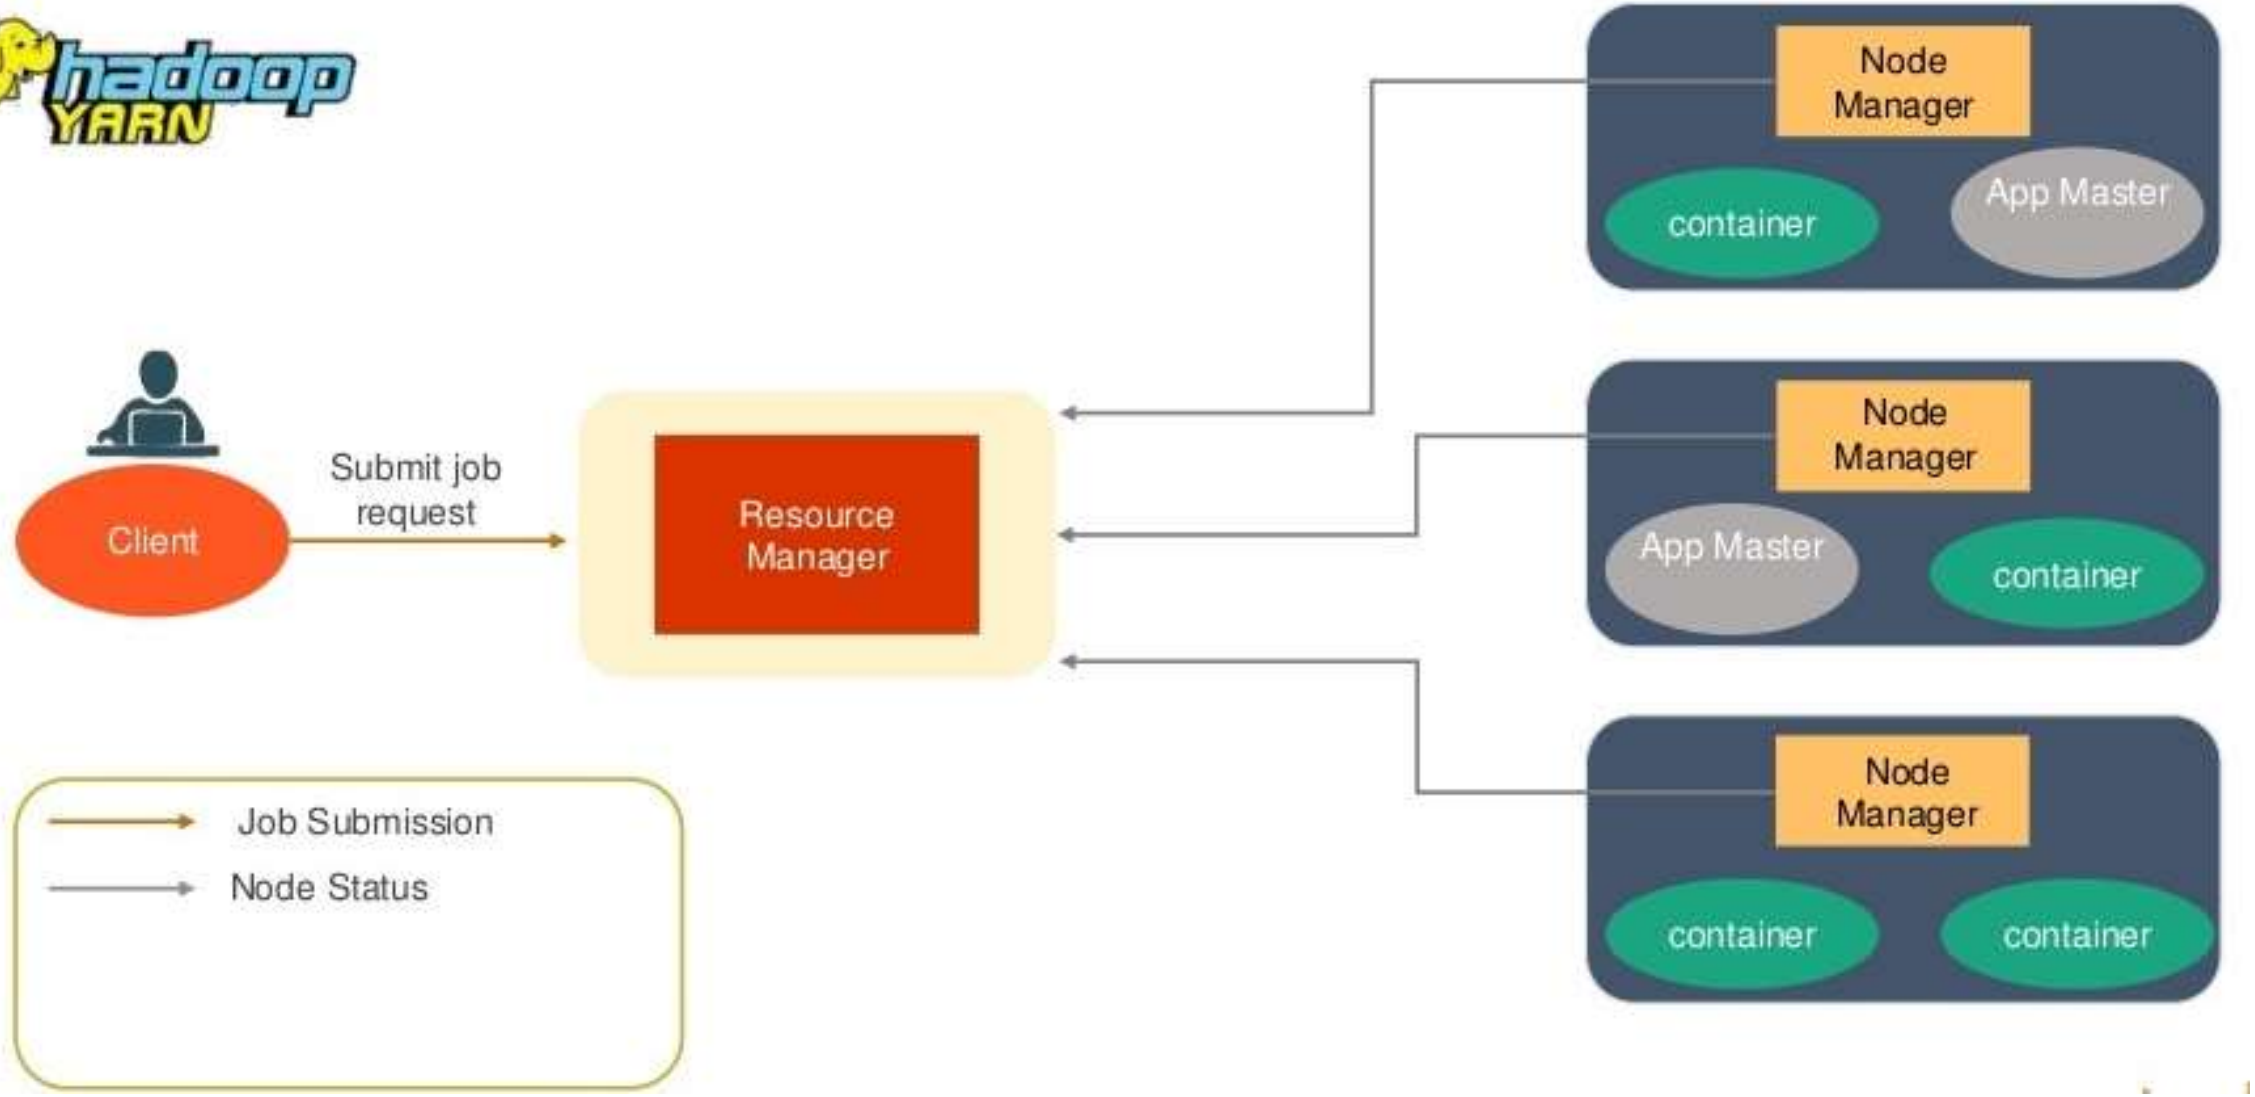

YARN Architecture

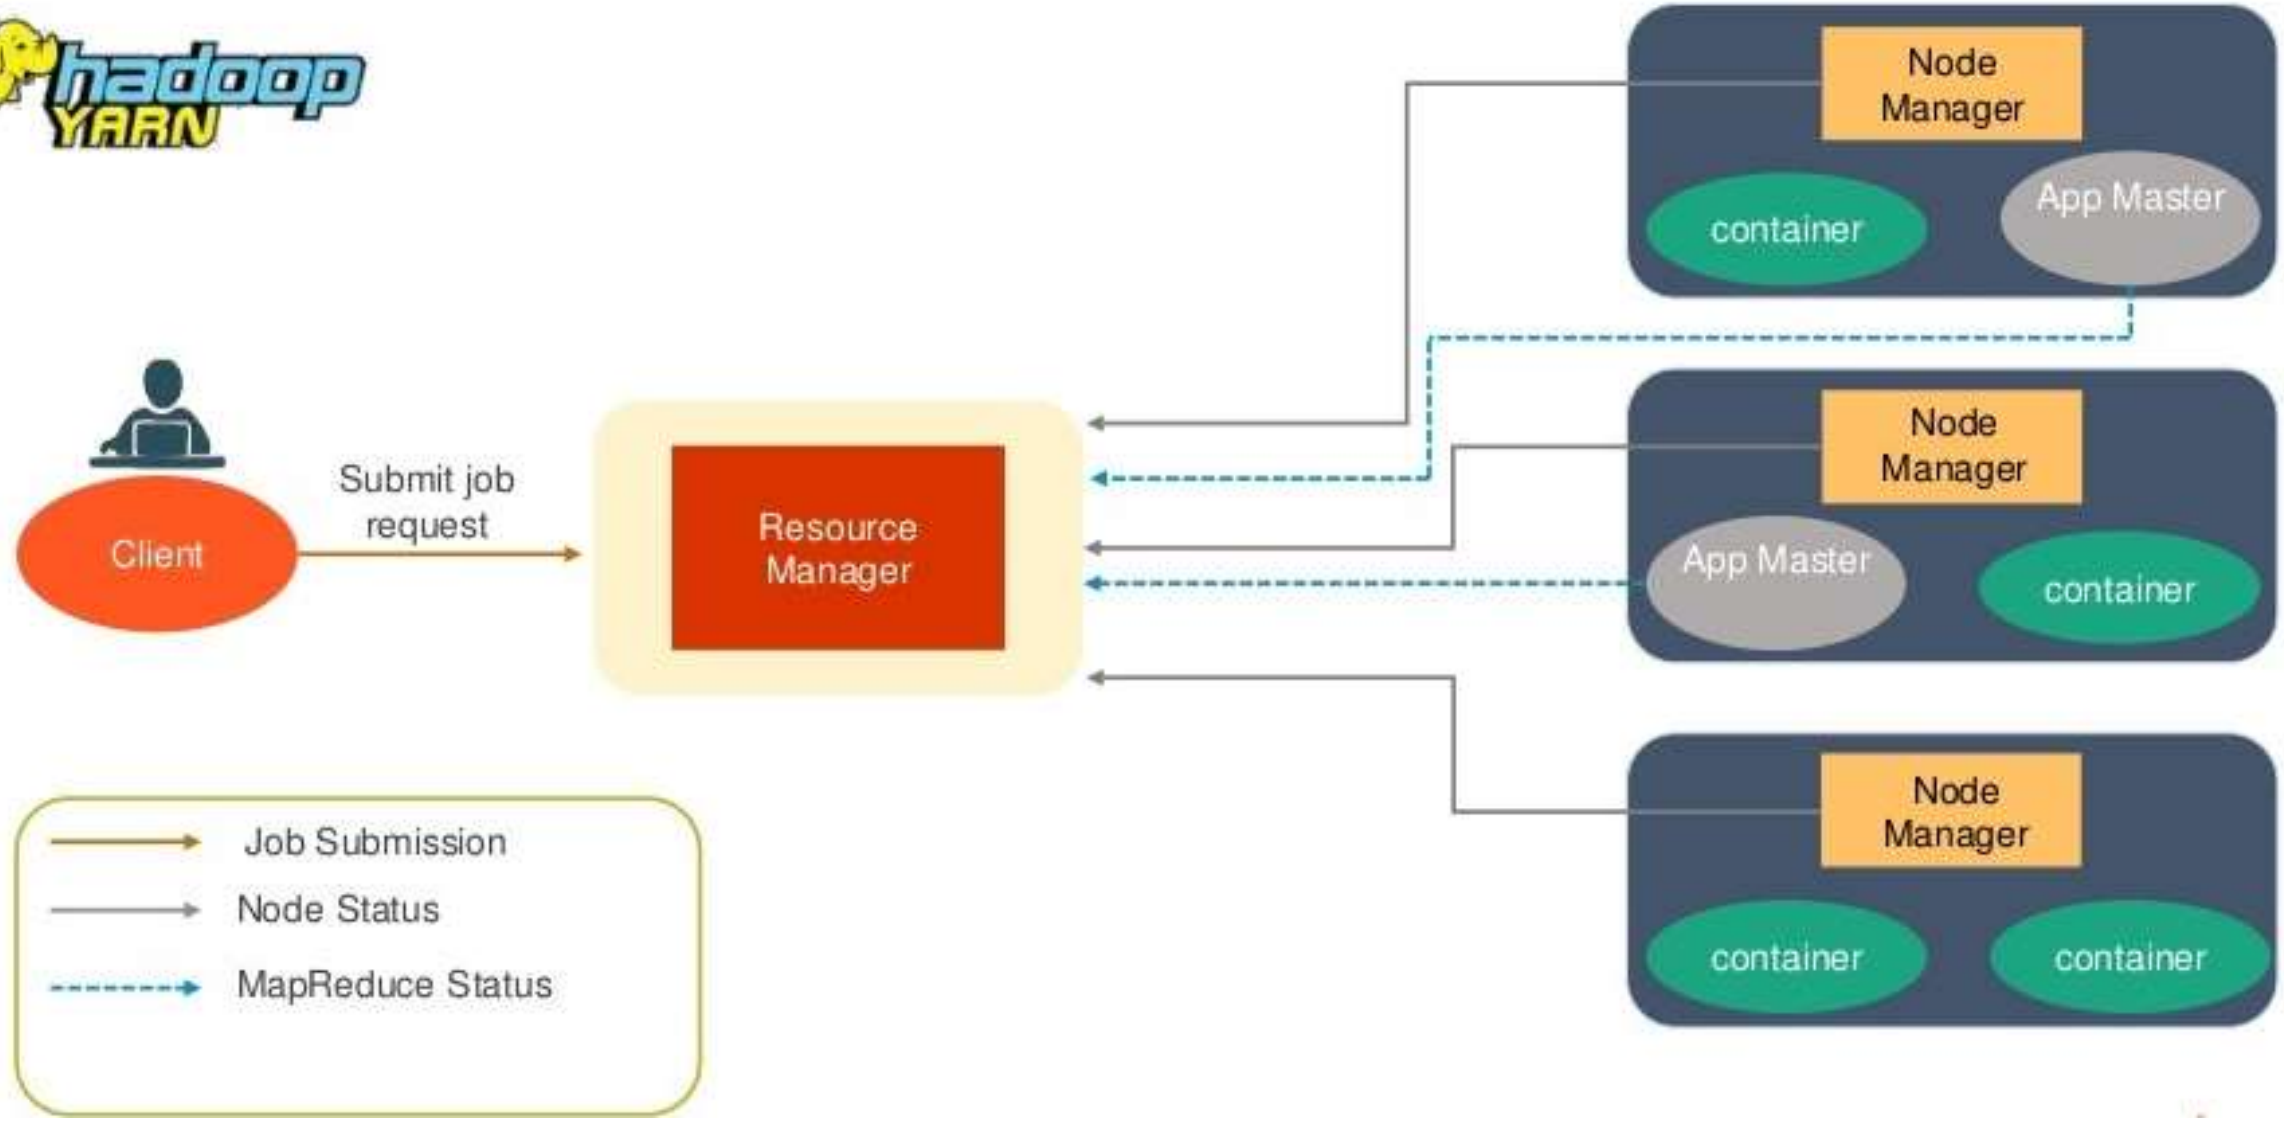

YARN Architecture

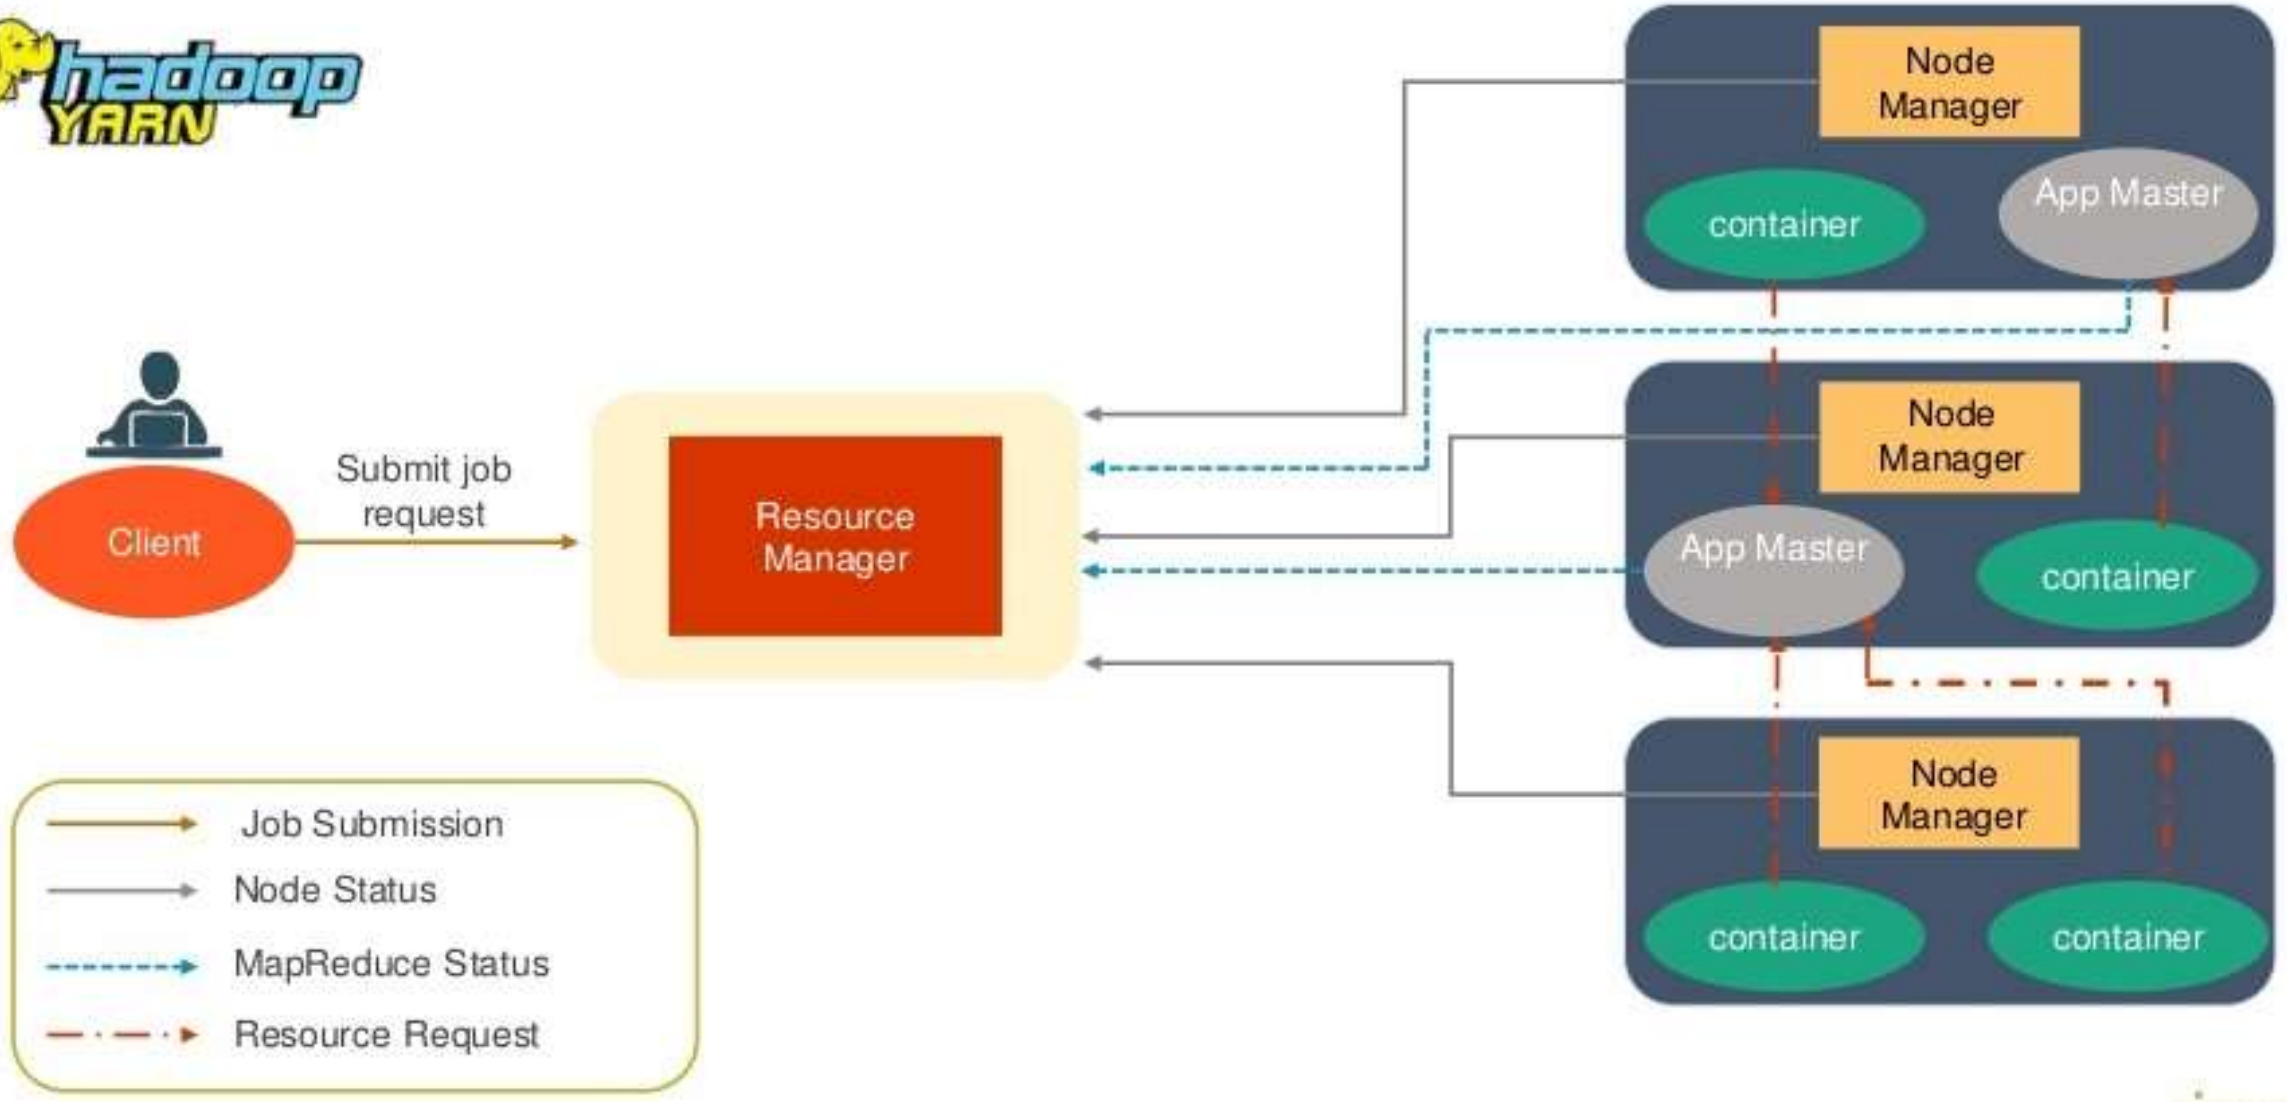

Hadoop Job Flow

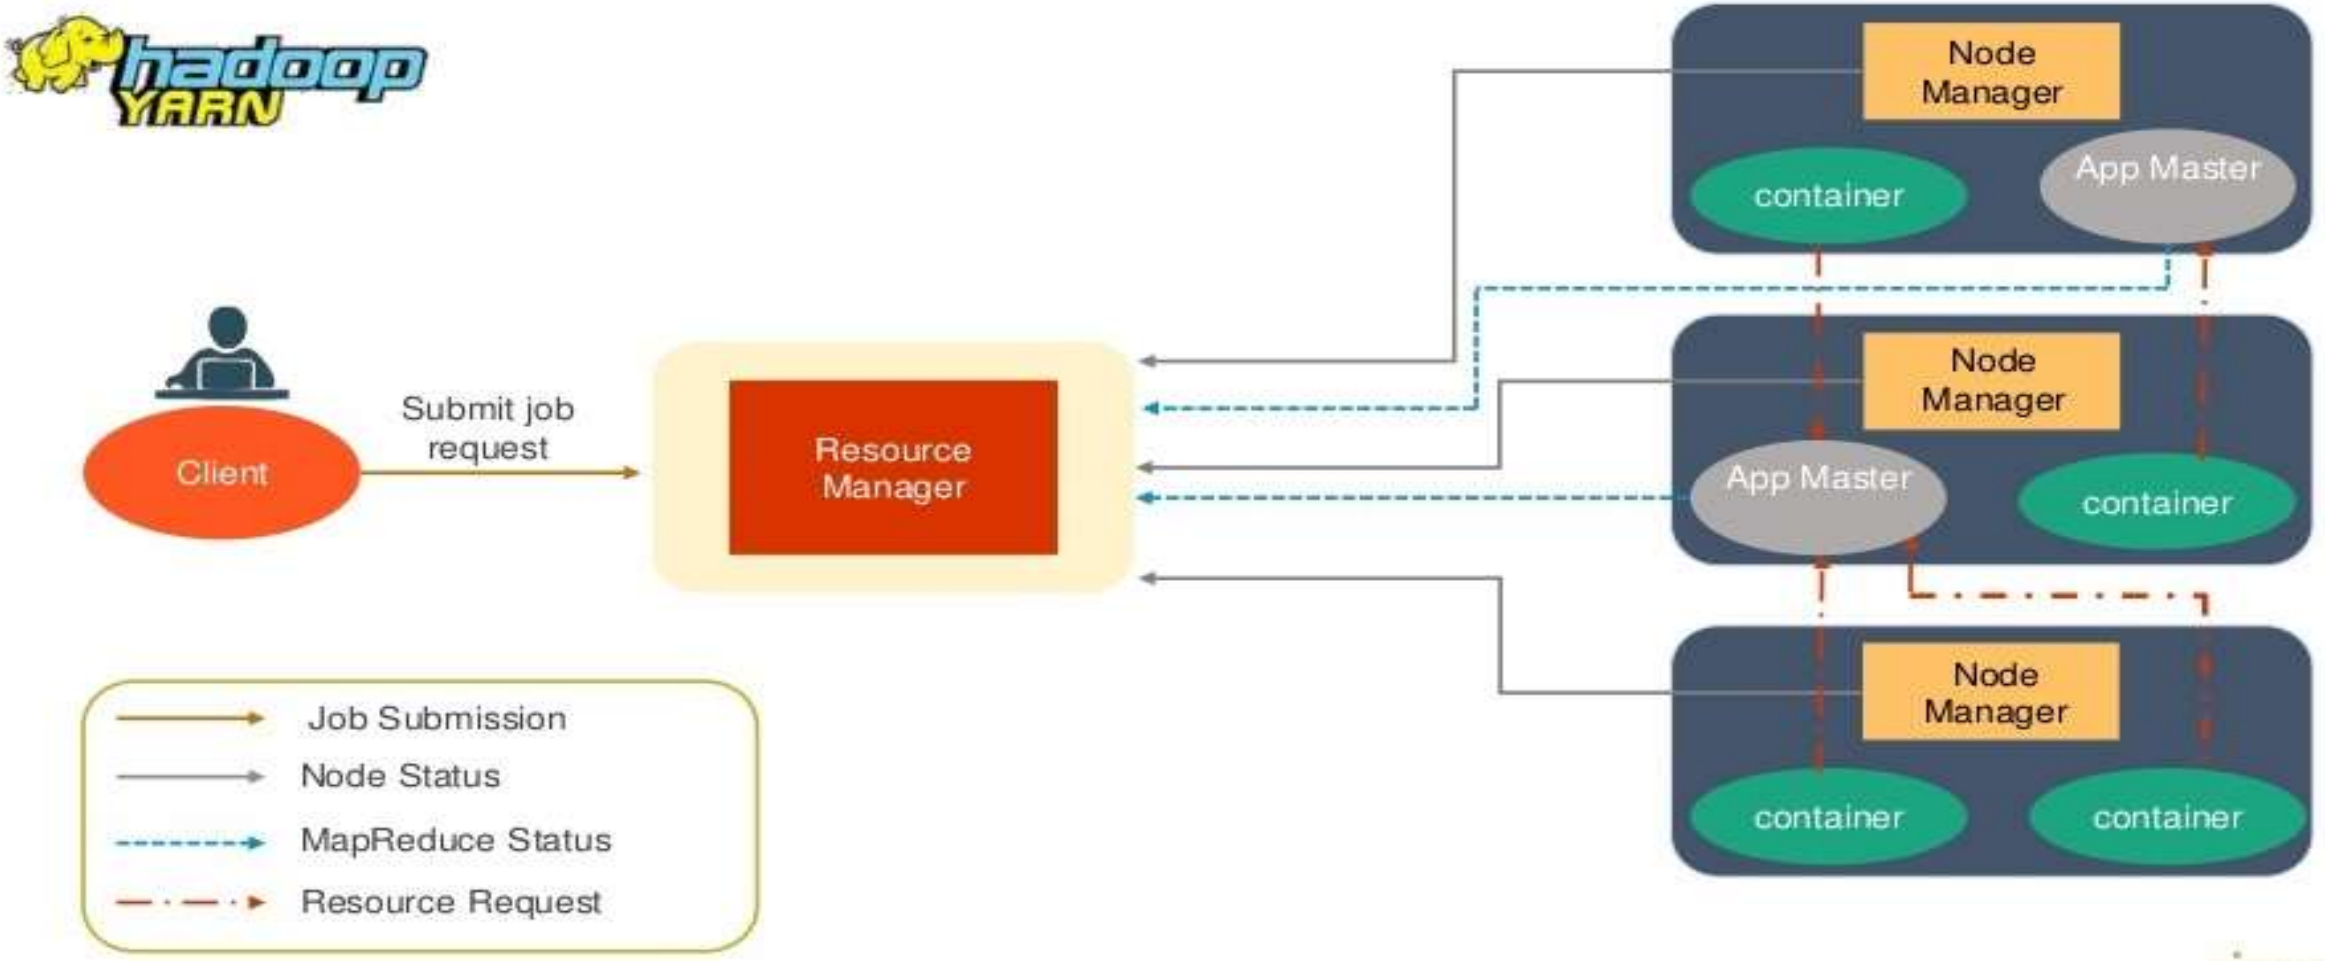

Hadoop Write Architecture

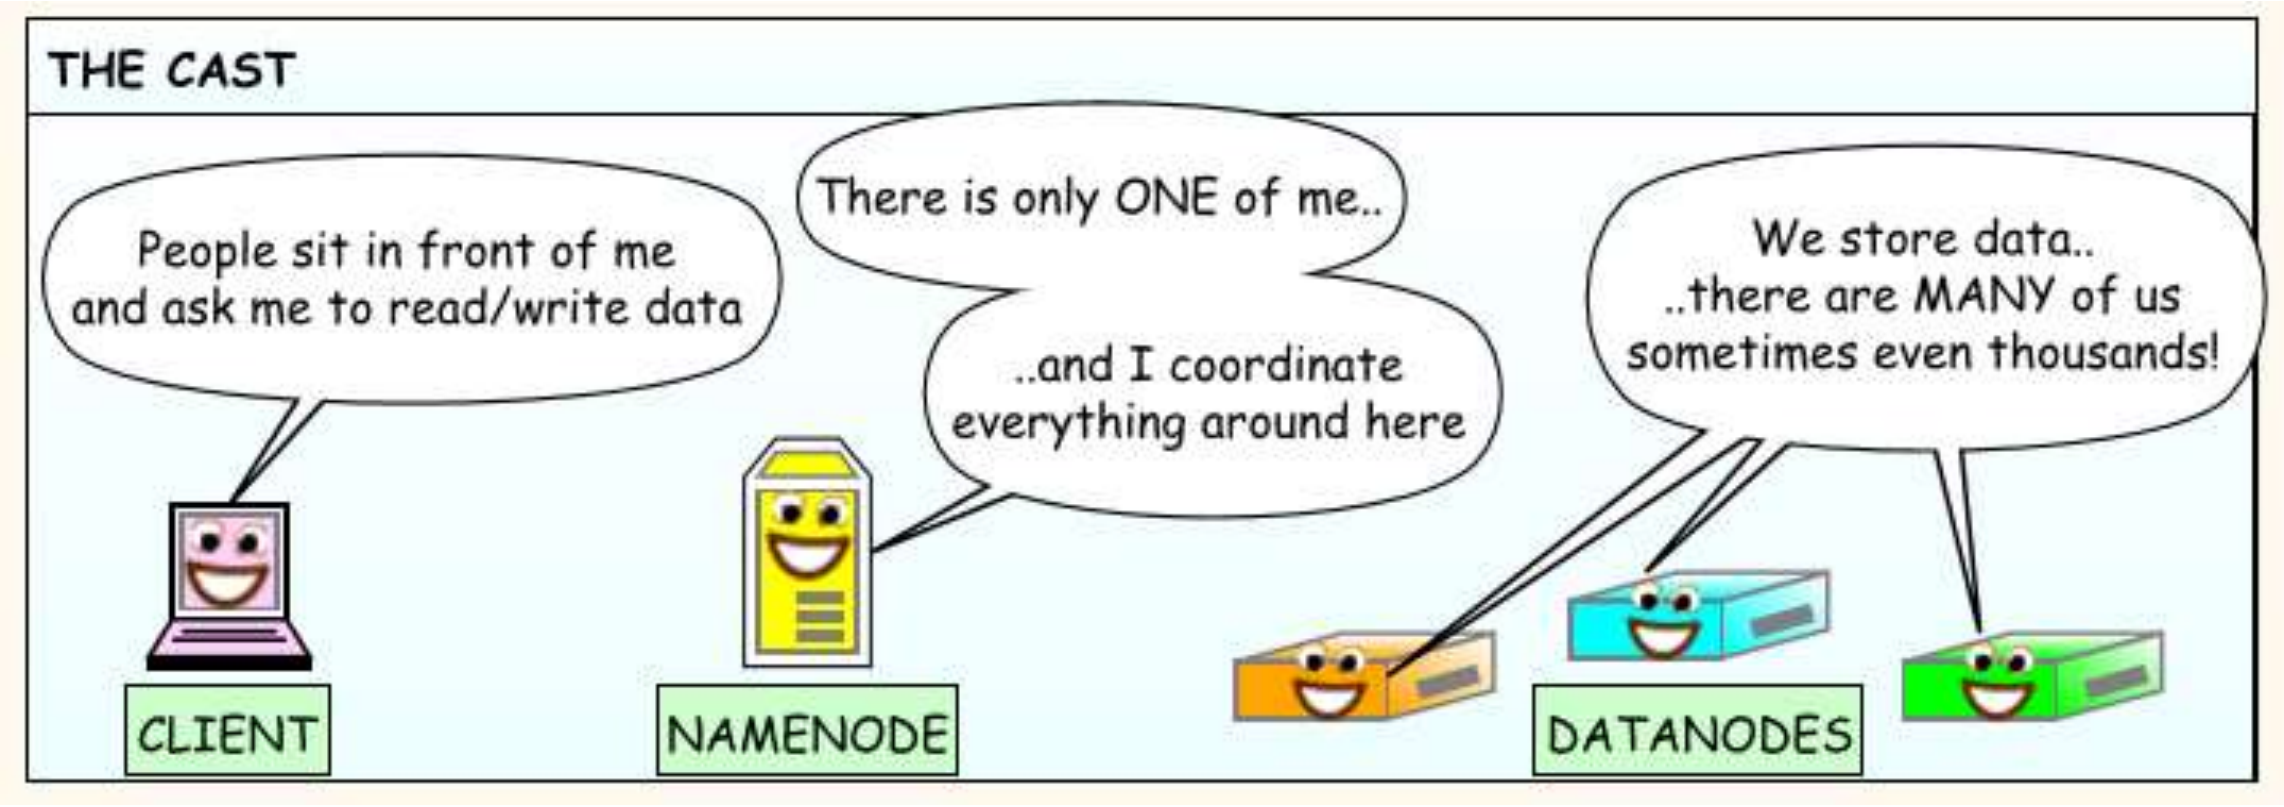

Hadoop Write Architecture

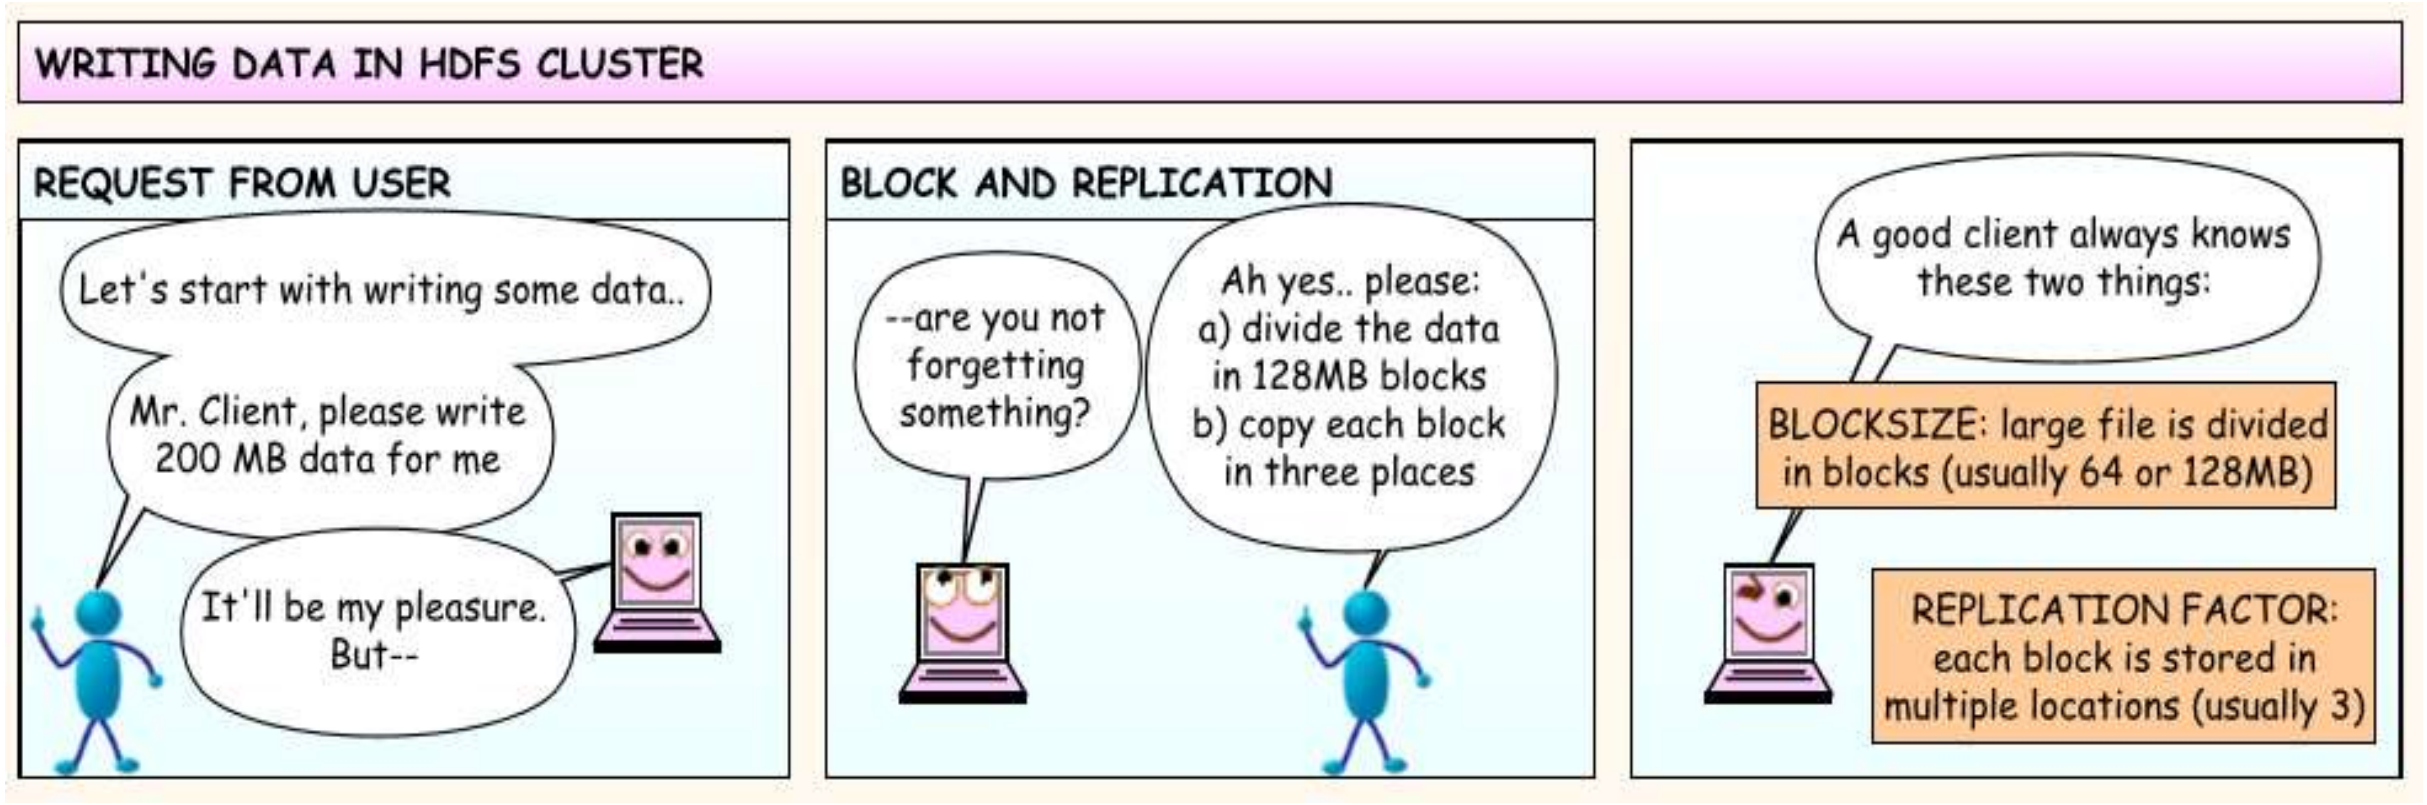

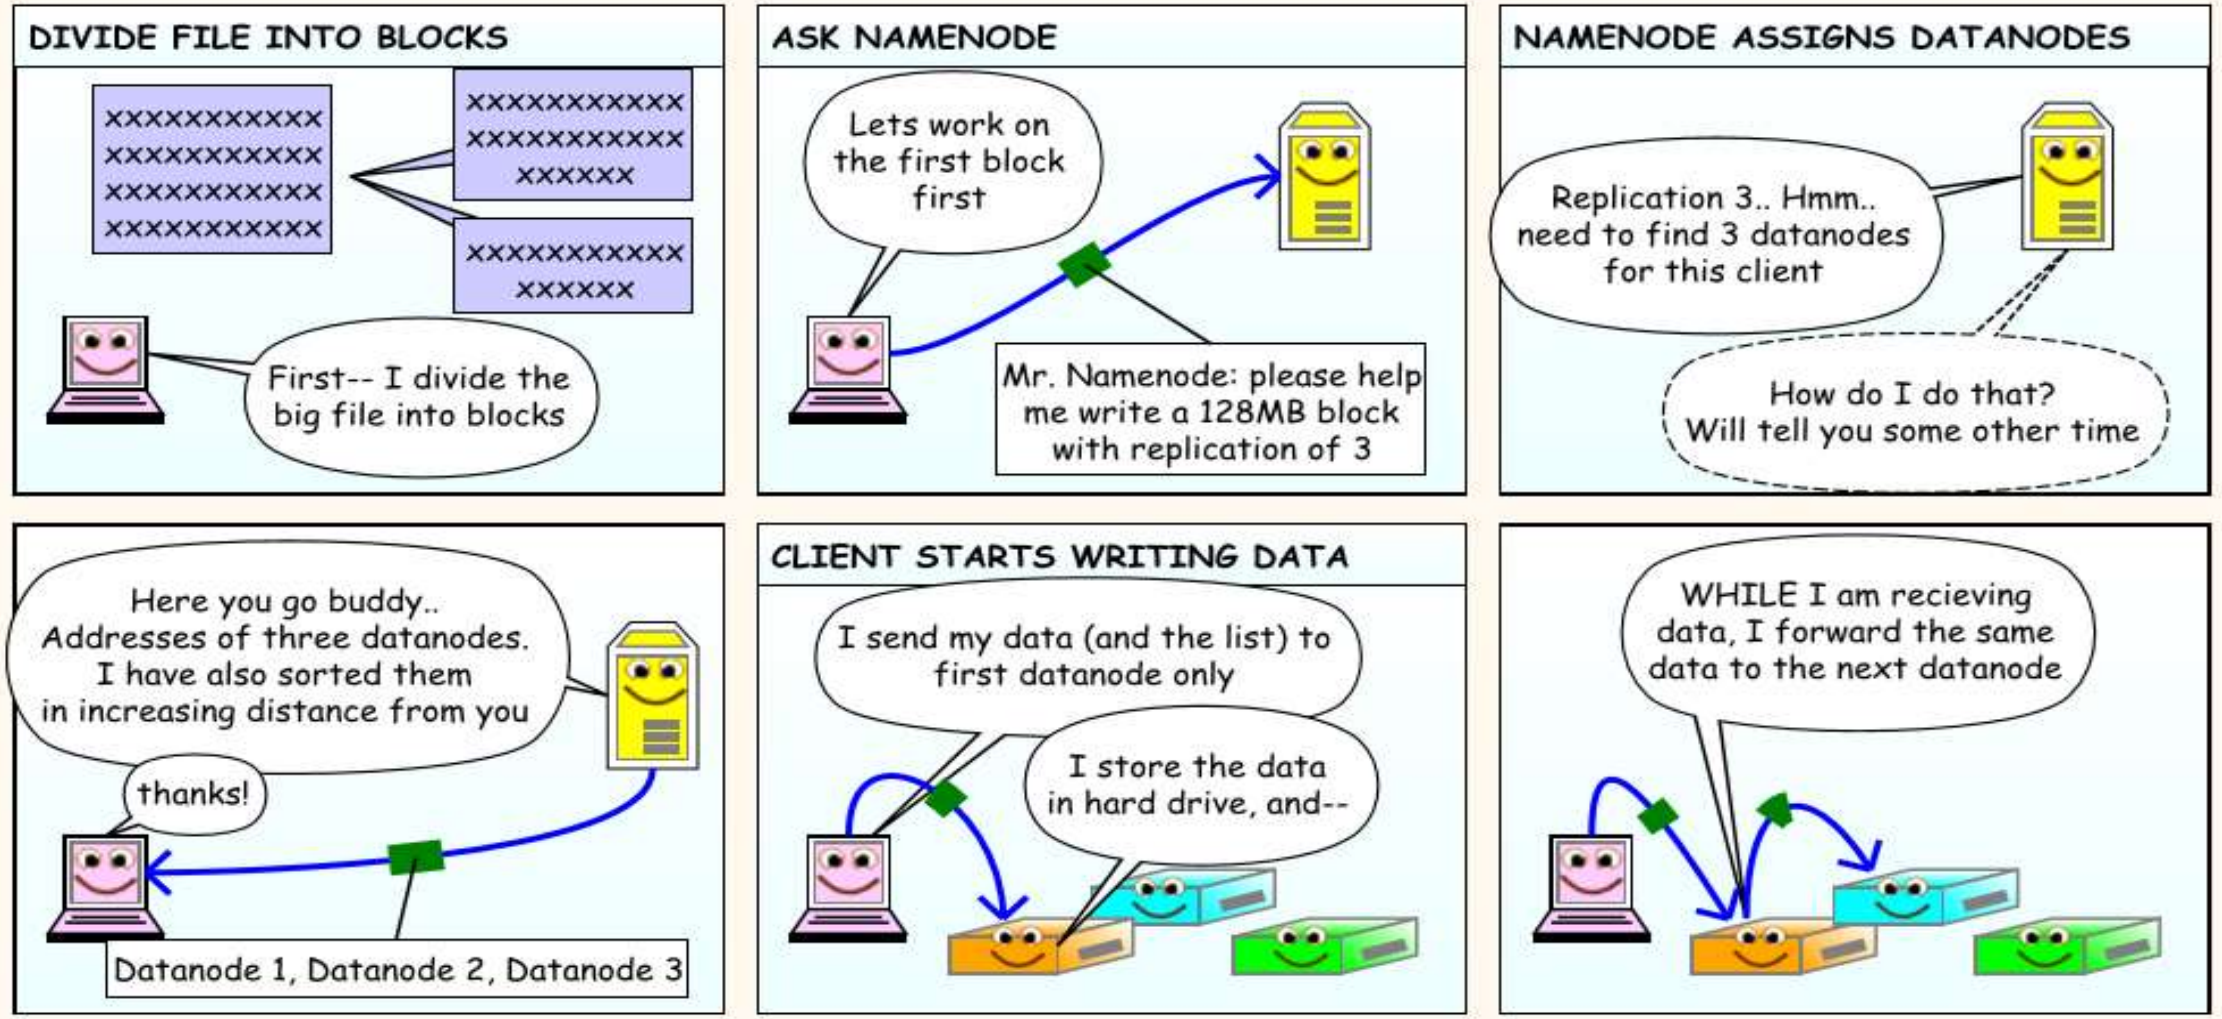
DIVIDE FILE INTO BLOCKS
First-- I divide the big file into blocks
ASK NAMENODE
Lets work on the first block first
Mr. Namenode: please help me write a 128MB block with replication of 3
NAMENODE ASSIGNS DATANODES
Replication 3.. Hmm.. need to find 3 datanodes for this client
How do I do that? Will tell you some other time
Here you go buddy.. Addresses of three datanodes. I have also sorted them in increasing distance from you
thanks!
Datanode 1, Datanode 2, Datanode 3
CLIENT STARTS WRITING DATA
I send my data (and the list) to first datanode only
I store the data in hard drive, and--
WHILE I am recieving data, I forward the same data to the next datanode

Hadoop Write Architecture

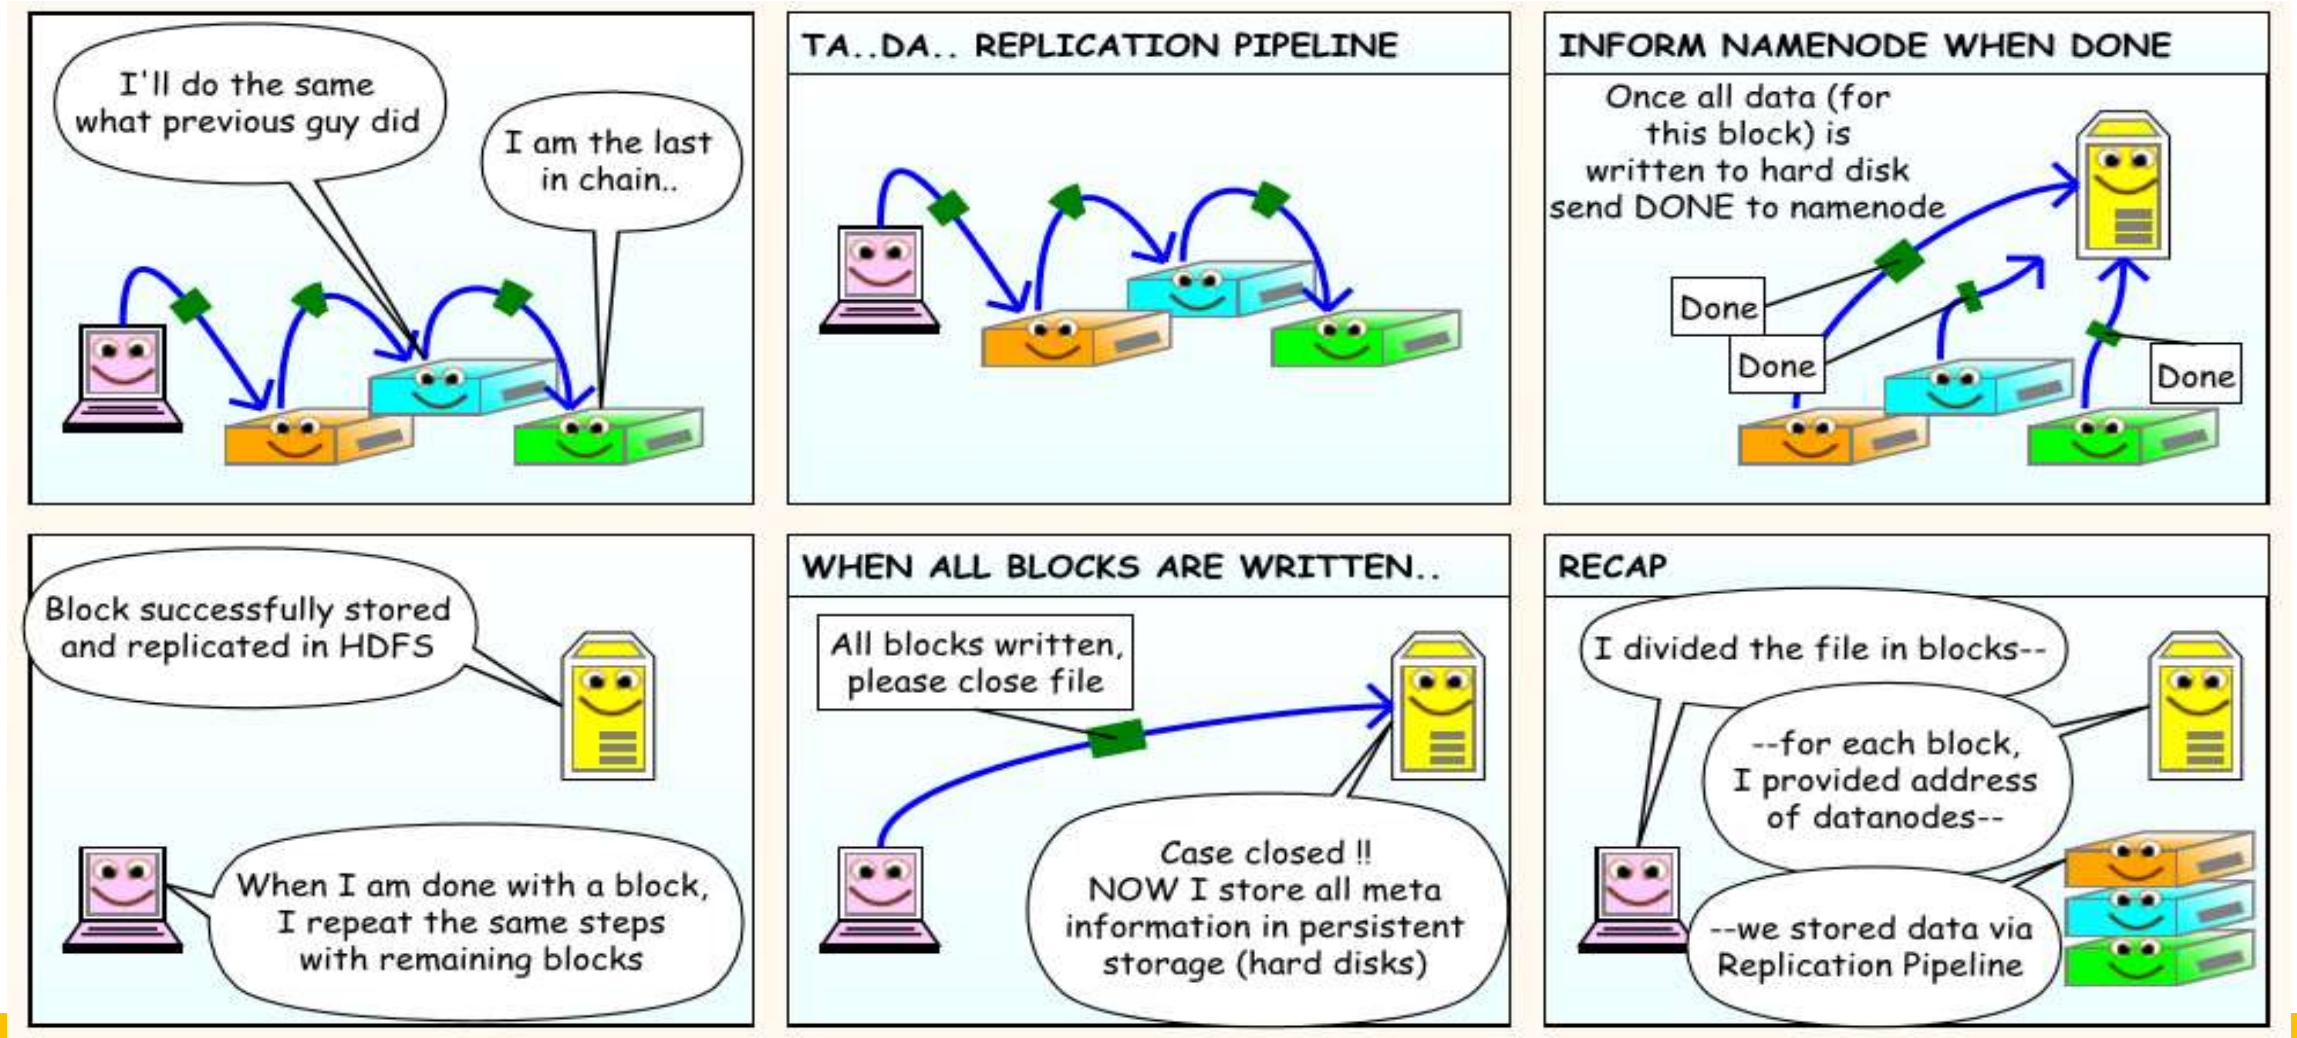

Hadoop Write Architecture

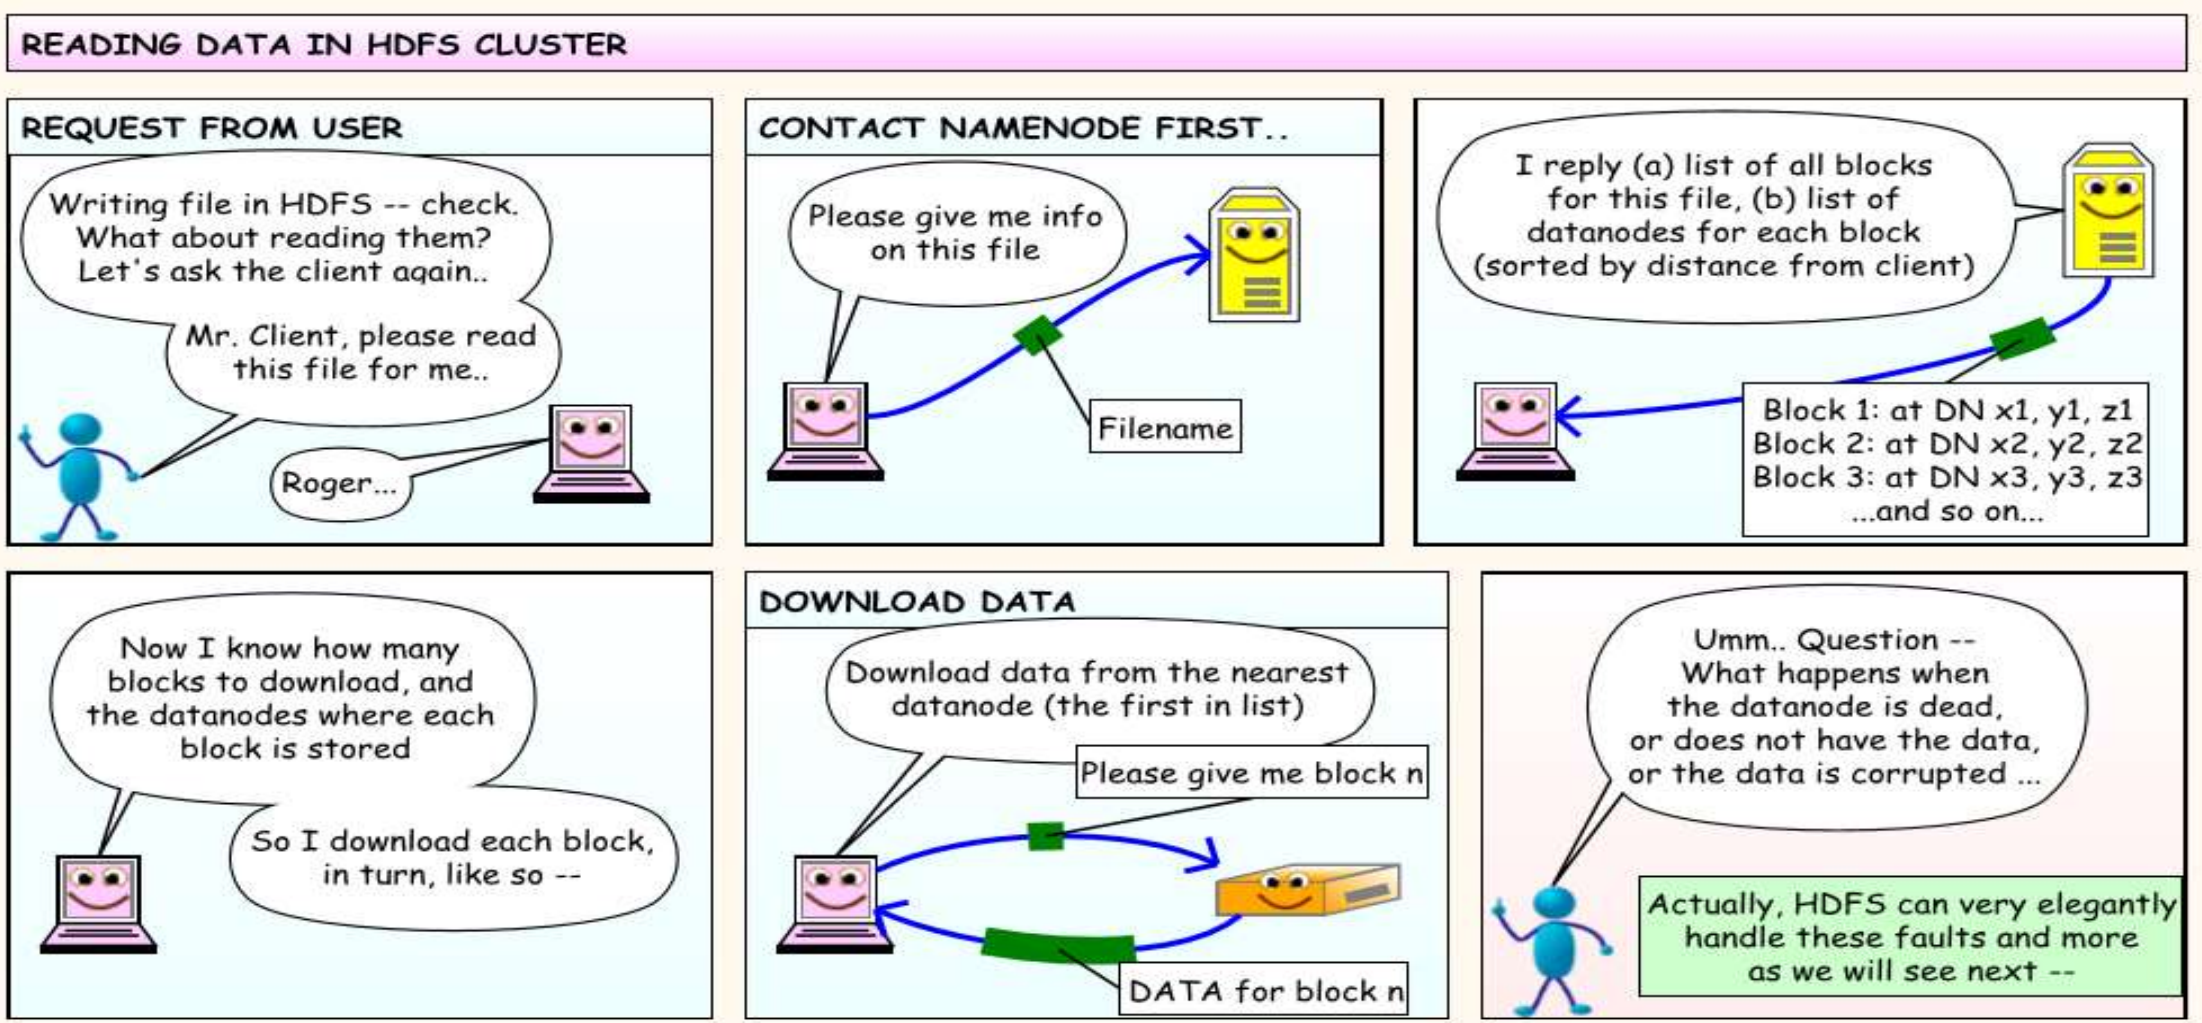

Hadoop Fault Tolerance Architecture

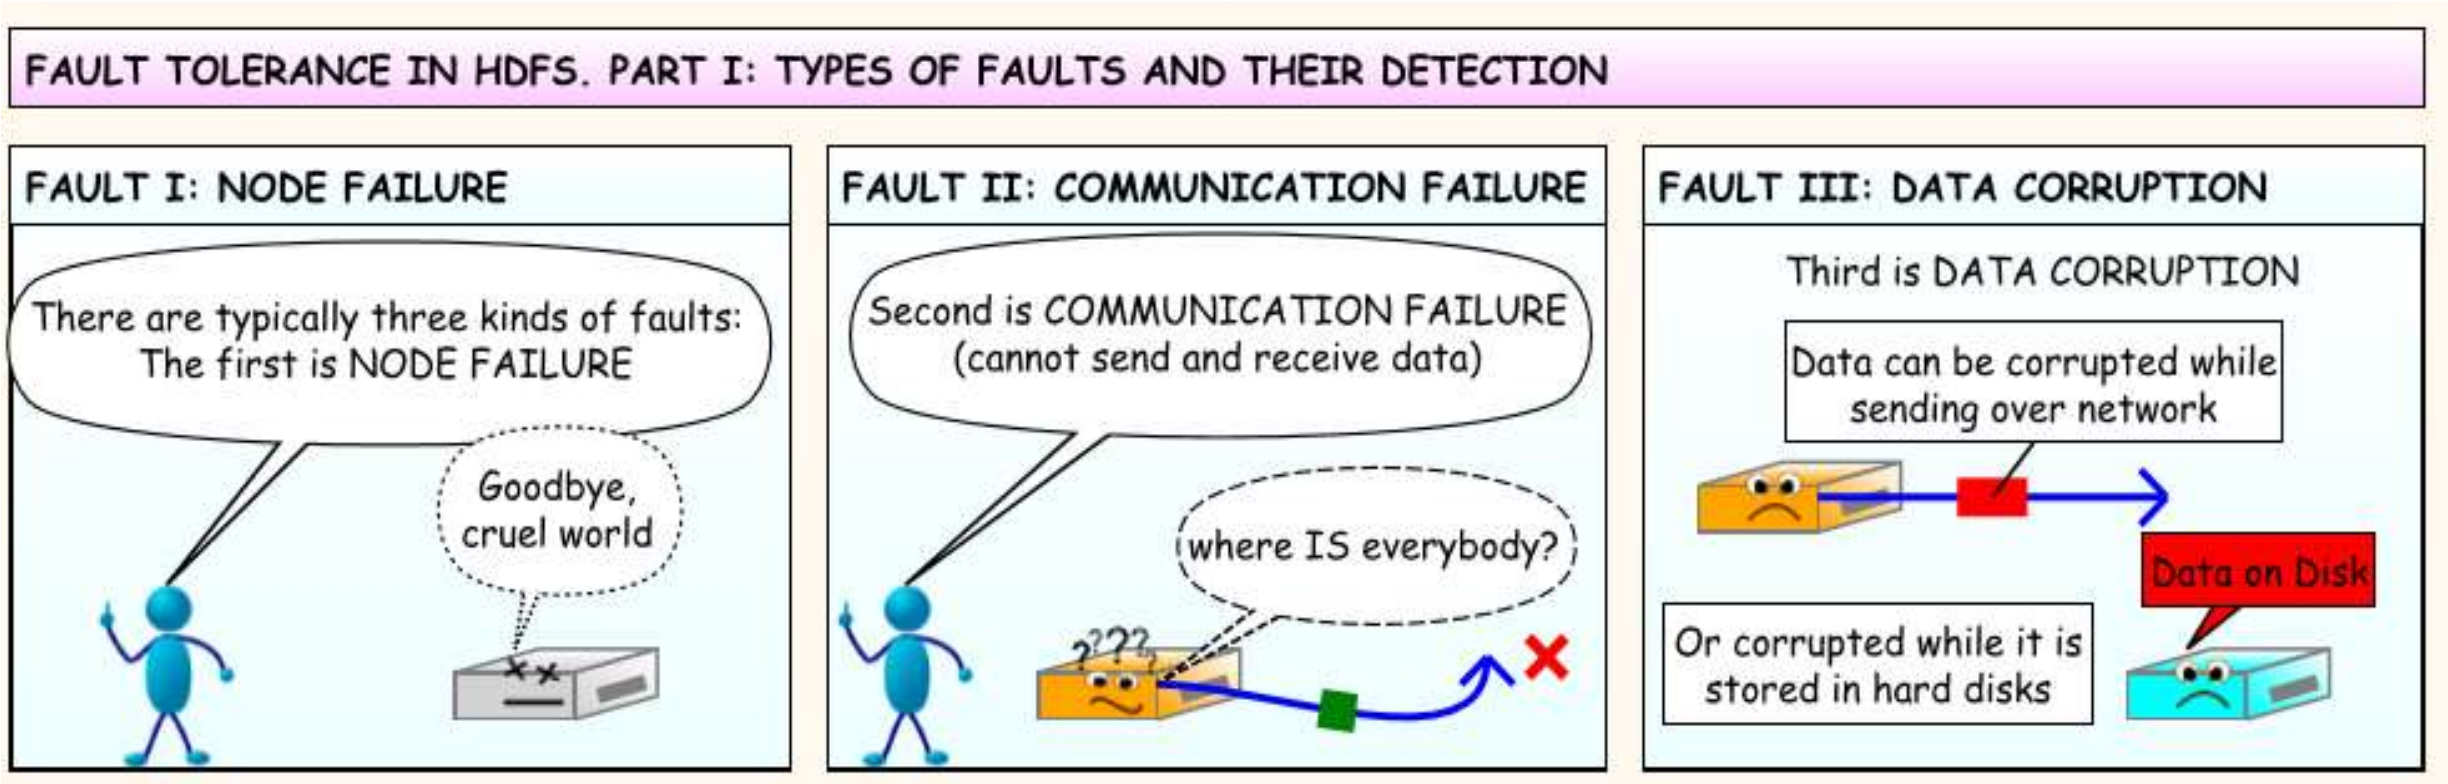

Hadoop Fault Tolerance Architecture

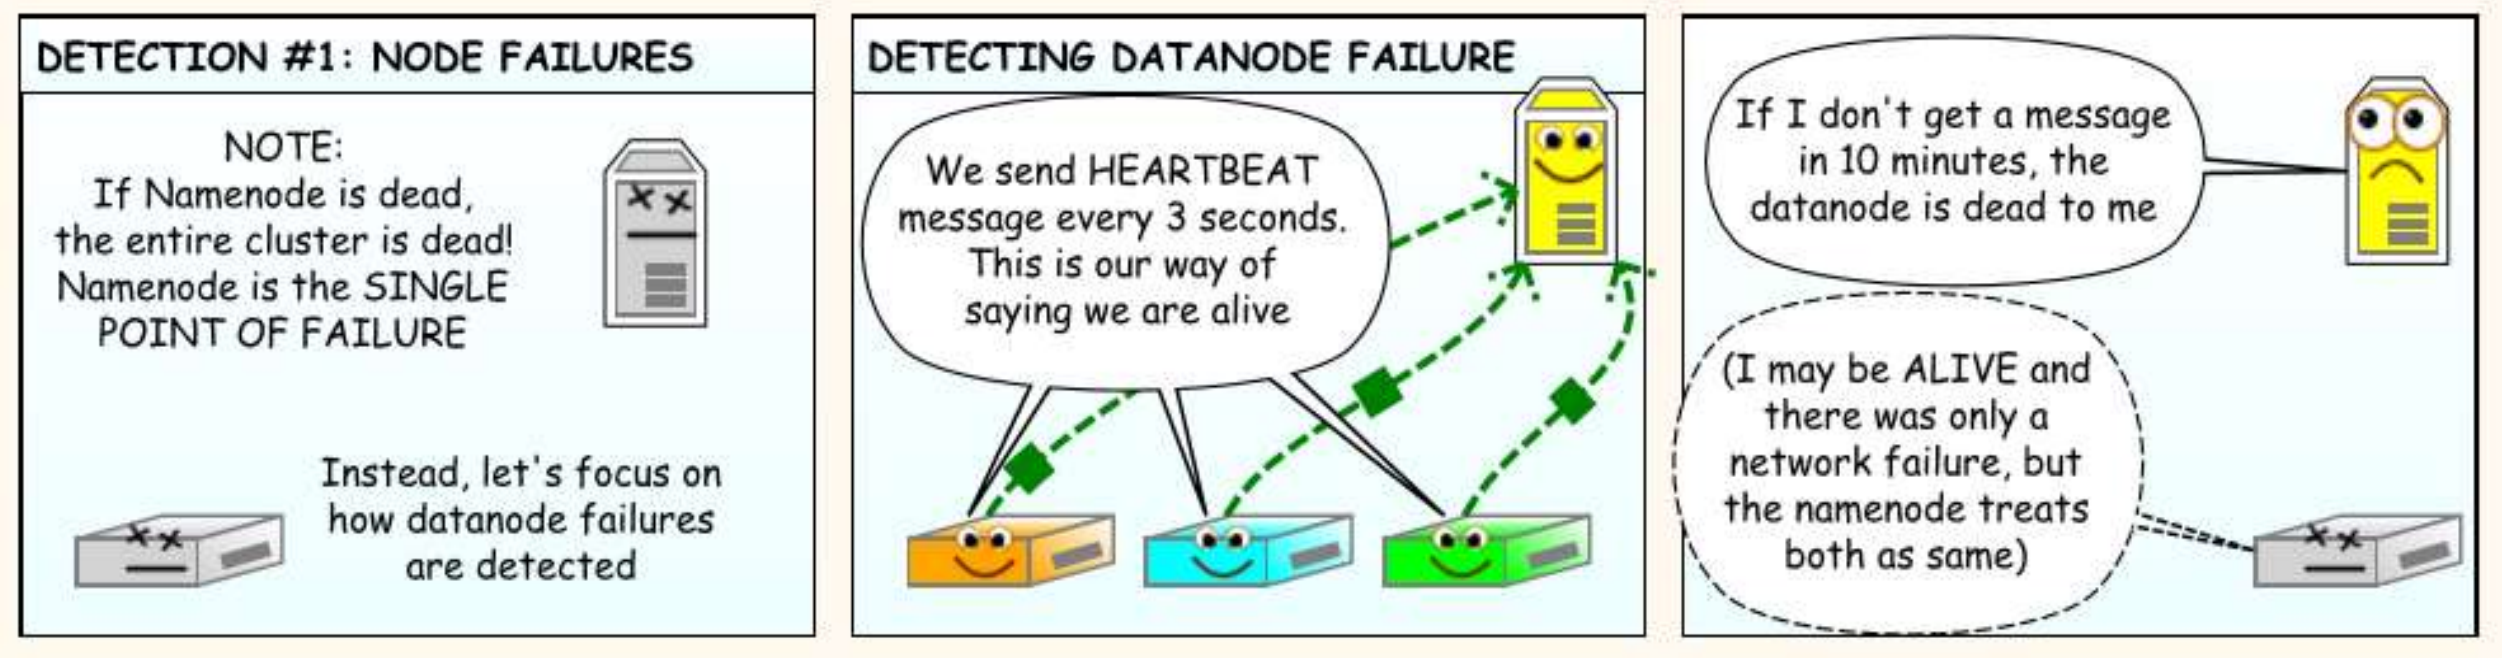

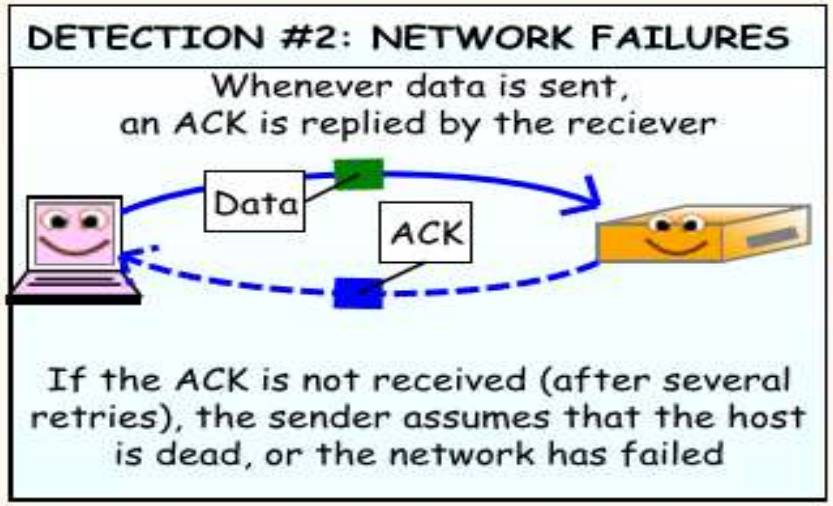
DETECTION #2: NETWORK FAILURES
Whenever data is sent,
an ACK is replied by the reciever
Data
ACK
If the ACK is not received (after several retries), the sender assumes that the host is dead, or the network has failed

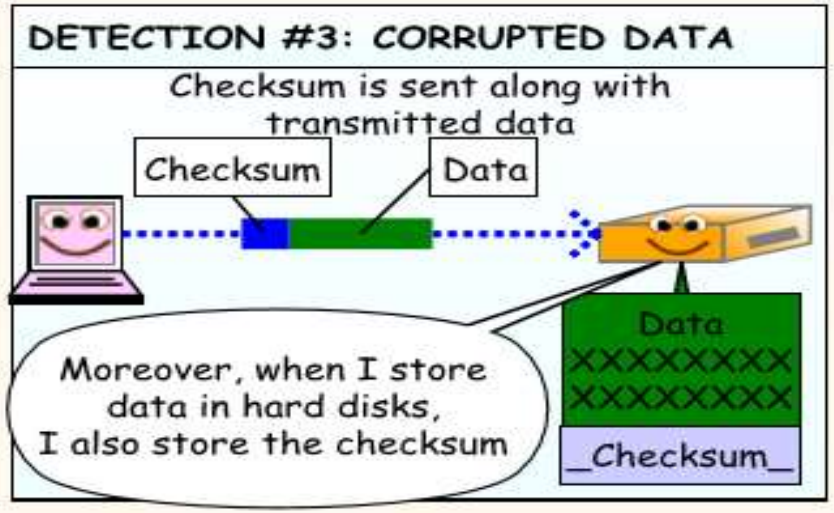
DETECTION #3: CORRUPTED DATA
Checksum is sent along with transmitted data
Checksum
Data
Moreover, when I store data in hard disks, I also store the checksum
Data
XXXXXXXX
XXXXXXXX
Checksum

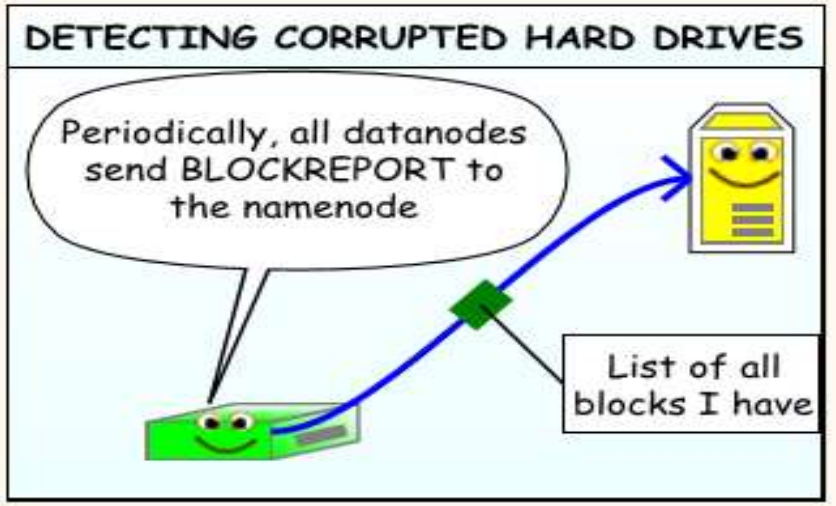
DETECTING CORRUPTED HARD DRIVES
Periodically, all datanodes send BLOCKREPORT to the namenode
List of all blocks I have

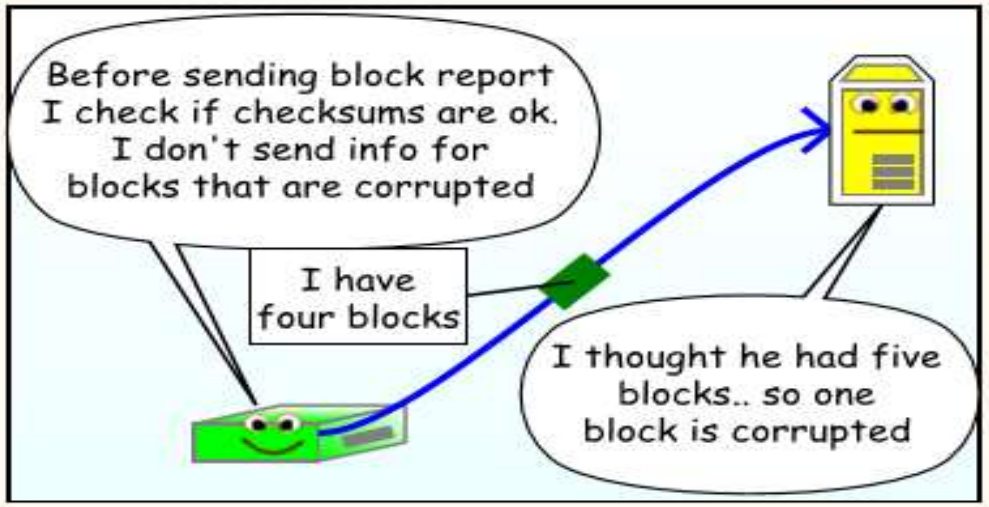
Before sending block report I check if checksums are ok. I don't send info for blocks that are corrupted
I have four blocks
I thought he had five blocks.. so one block is corrupted

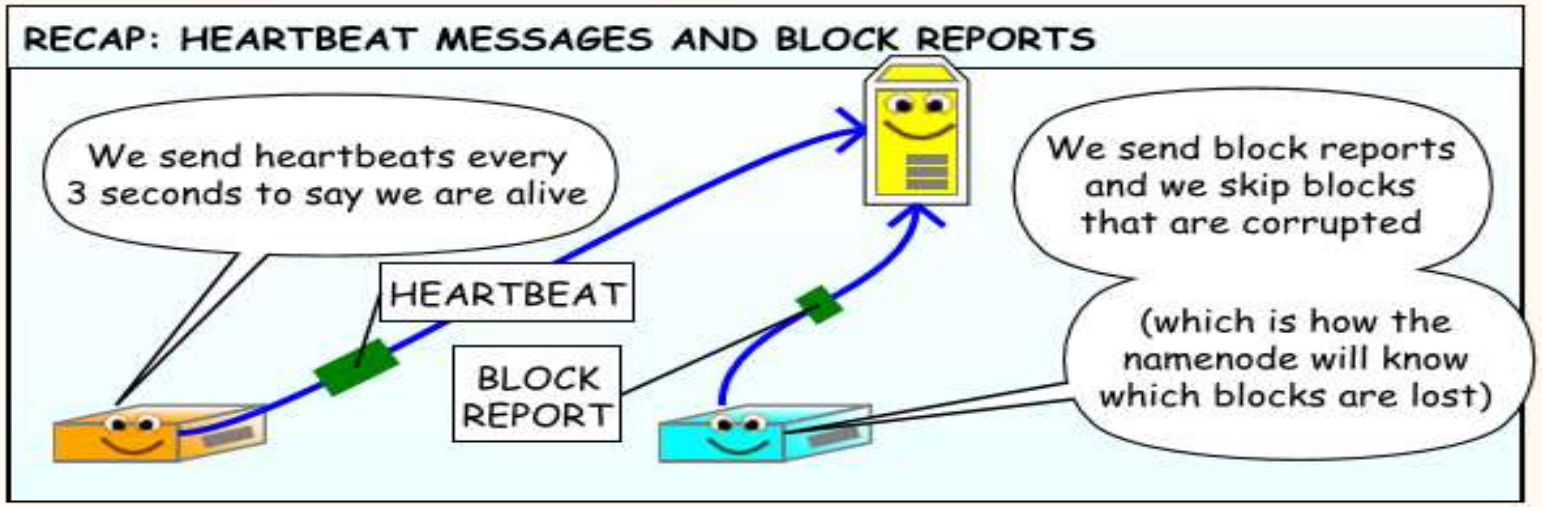
RECAP: HEARTBEAT MESSAGES AND BLOCK REPORTS
We send heartbeats every 3 seconds to say we are alive
HEARTBEAT
BLOCK REPORT
We send block reports and we skip blocks that are corrupted
(which is how the namenode will know which blocks are lost)

Hadoop Write Architecture

Lets Break in size of 256 MB, consider it as Block Size

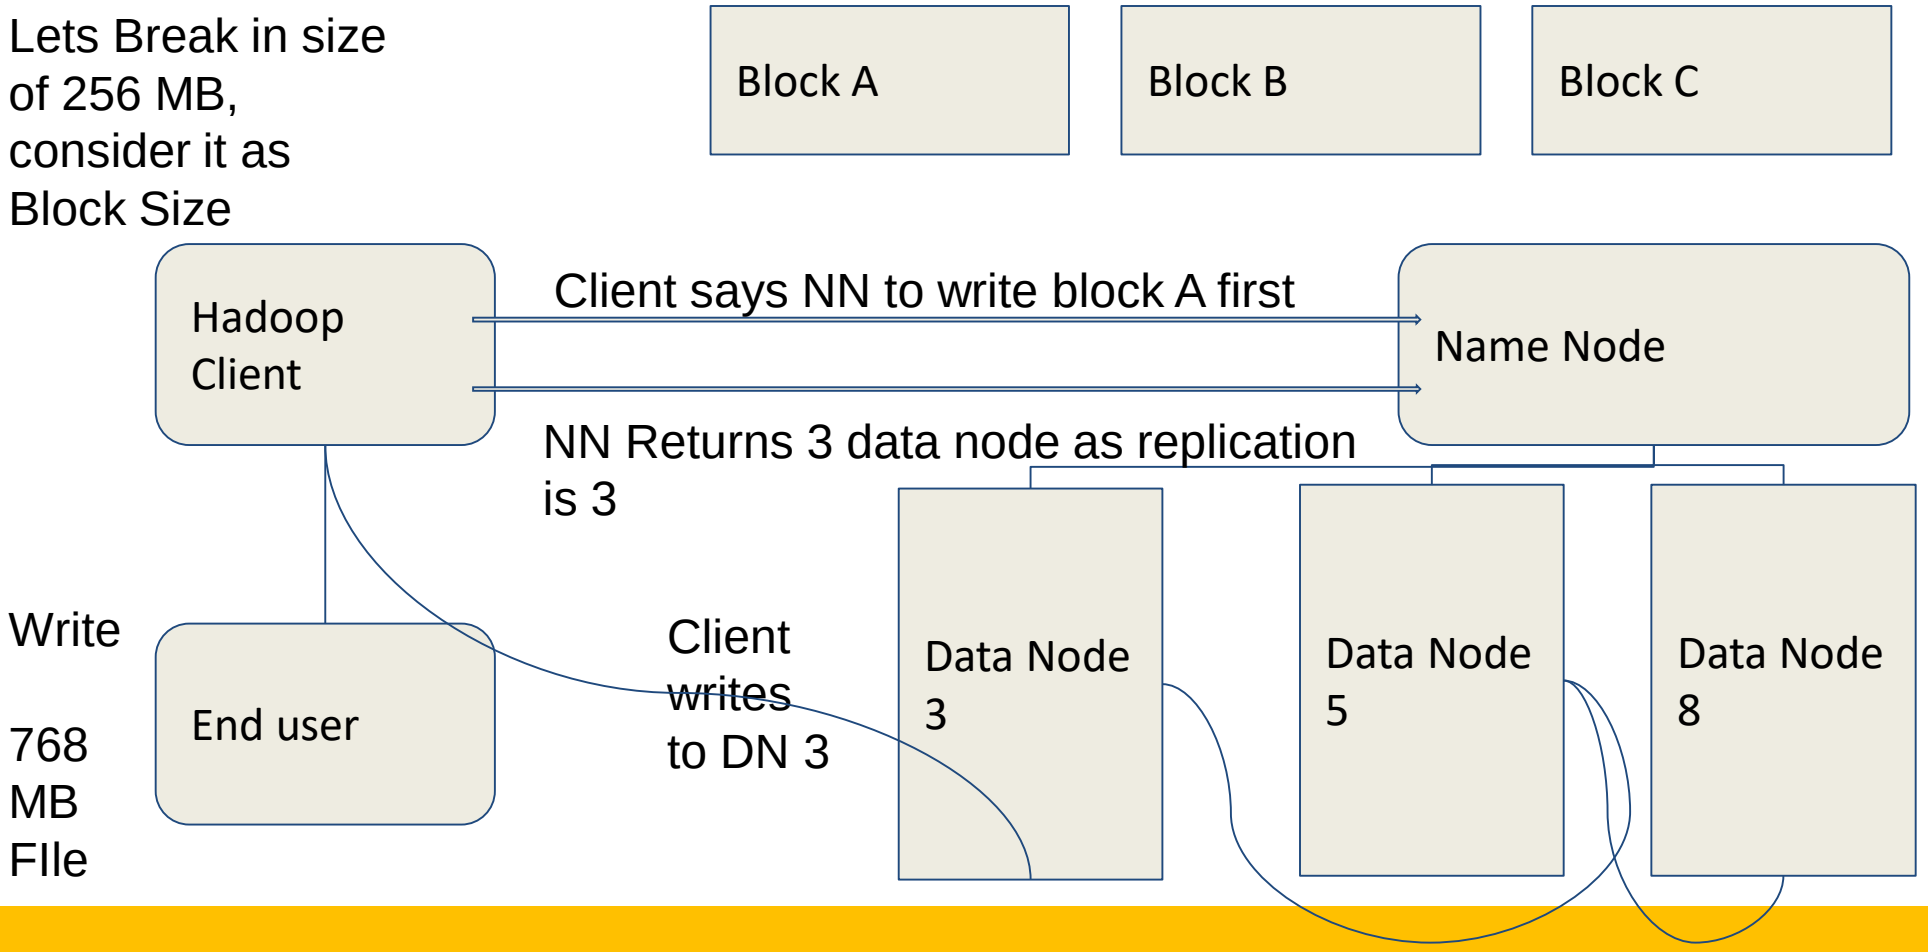

Hadoop Read Architecture

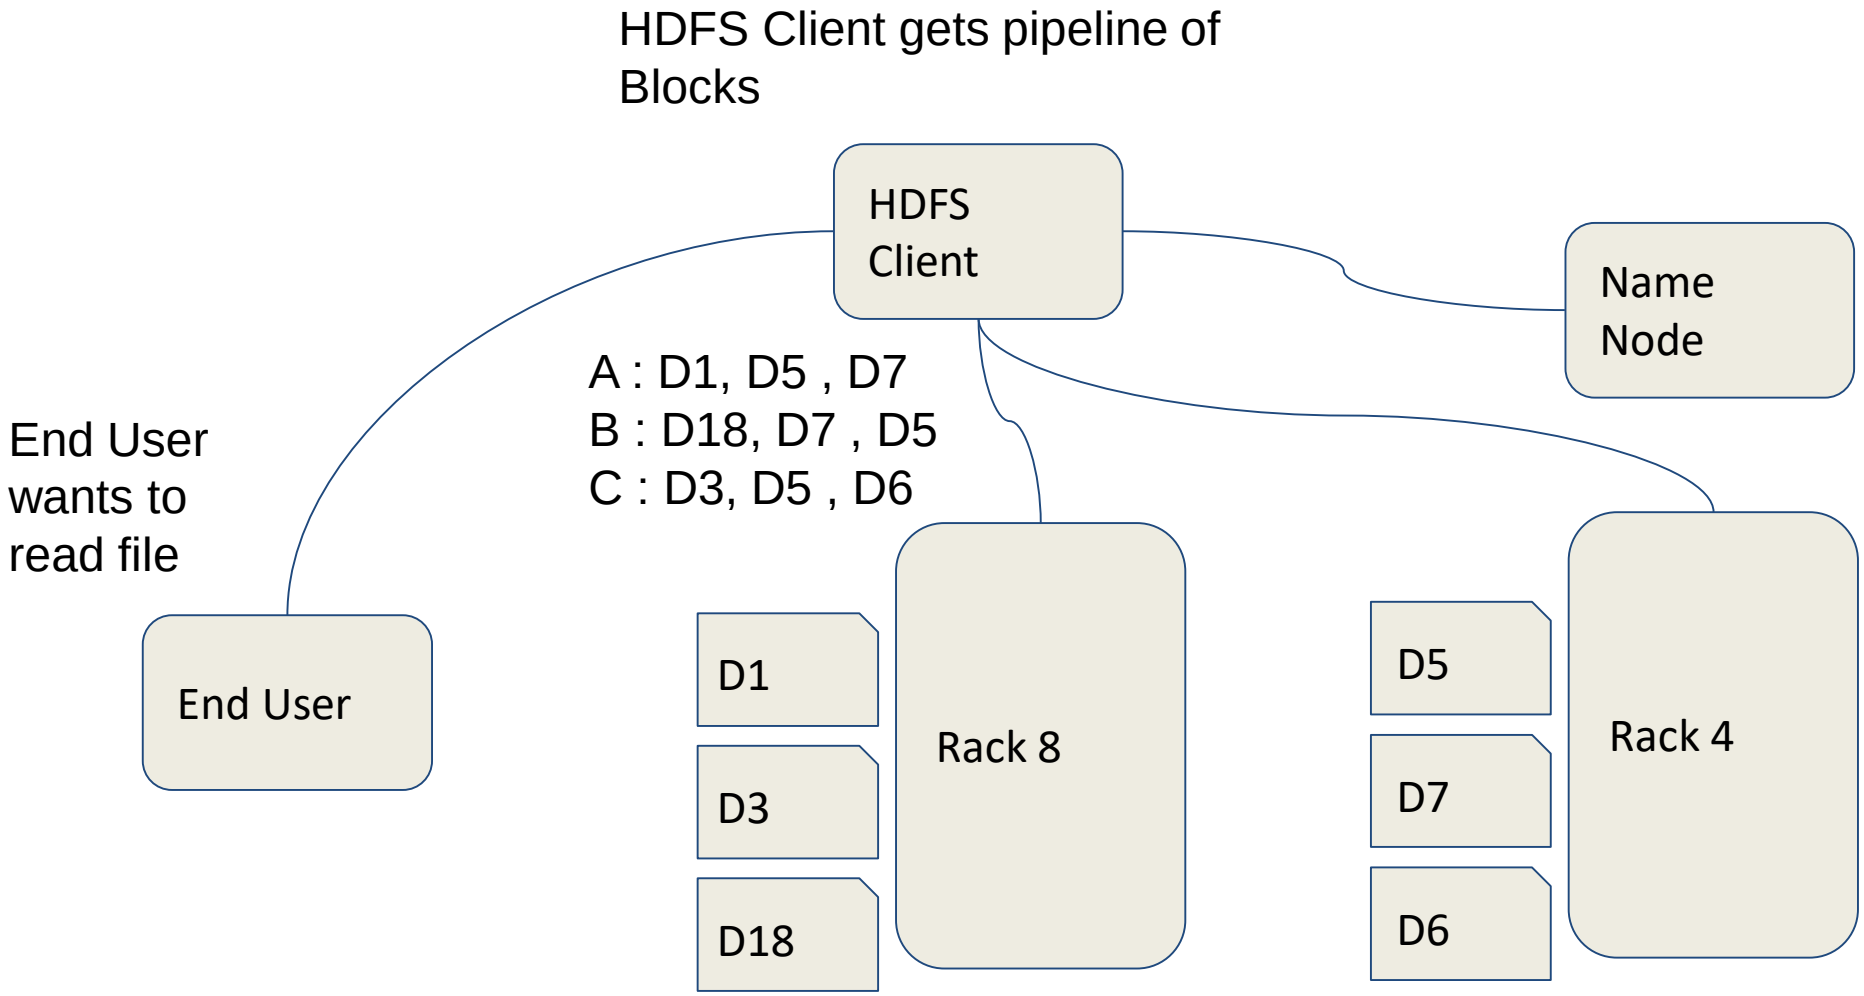

HDFS Commands

a. HDFS Commands
   i. mkdir
   ii. Ls
   iii. Put
   iv. Get
   v. Cat
   vi. Cp
   vii. Mv

Introduction to Map Reduce

MapReduce is a programming model whose libraries have been primarily written in Java

- Record-oriented data processing
- Operates on key and value pairs
- Consists of two phases :
  - ✓ Map
  - ✓ Reduce
- Every Map/Reduce program must specify a Mapper and optionally a Reducer
- Where ever possible, each node processes data stored on that node (Data Locality)
- Provides pluggable APIs and configuration mechanisms for writing applications Map and Reduce functions
  Input formats and splits
  Number of tasks, data types, etc…

MAPREDUCE

- MapReduce provides
  - Automatic parallelization, distribution
  - I/O scheduling
    - ✓ Load balancing
    - ✓ Network and data transfer optimization
- Fault tolerance
  - ✓ Handling of machine failures
- **Need more power: Scale out, not up!**
  - ✓ Large number of **commodity servers(**less processors and less RAM) as opposed to some high end specialized servers, this saves cost.

Technical Prerequisites

proschool
An ims Initiative

- It is expected that the students have knowledge about the following concepts in Java for understanding MapReduce Working

  - ✓ Class and Object
  - ✓ OOPS concepts : Abstraction ,Encapsulation ,Polymorphism and Inheritance
  - ✓ Abstract class and abstract methods
  - ✓ Keyword : Extends and implements meaning
  - ✓ Input formats (int,string,Boolean,double ,etc)
  - ✓ Knowledge of exceptions and exception handling (to understand code)
  - ✓ Use of loops like for,if,while etc in Java
  - ✓ Basic datastrucutres

- Apart from this, knowledge of basic Hadoop shell commands is required.

Motivation for MapReduce

- **Large-Scale Data Processing**
  - ✓ Need to manage large number of machines(CPU's) but the system should be hassle free.
- **MapReduce Architecture provides**
  - ✓ Automatic parallelization & distribution
  - ✓ Fault tolerance
  - ✓ I/O scheduling
  - ✓ Monitoring & status updates
  - ✓ Security and administration
  - ✓ Flexibility

Ways to use MAPREDUCE

Libraries

- HBase
- Hive
- Pig
- Sqoop
- Oozie
- Mahout

Languages

- Java
- HiveQL
- PigLatin
- Python
- JavaScript
- R

Keys and Values

- Keys are objects which implement WritableComparable
  - ✓ A Writable Comparable is a Writable which is also Comparable.
  - ✓ Two WritableComparables can be compared against each other to determine their order.
  - ✓ Keys must be WritableComparables because they are passed to the reducer in sorted order.

- Values are objects which implement Writable interface
  - ✓ IntWritable for ints
  - ✓ LongWritable for longs
  - ✓ FloatWritable for floats
  - ✓ DoubleWritable for doubles
  - ✓ Text for strings

Map Reduce Steps

MAPREDUCE working

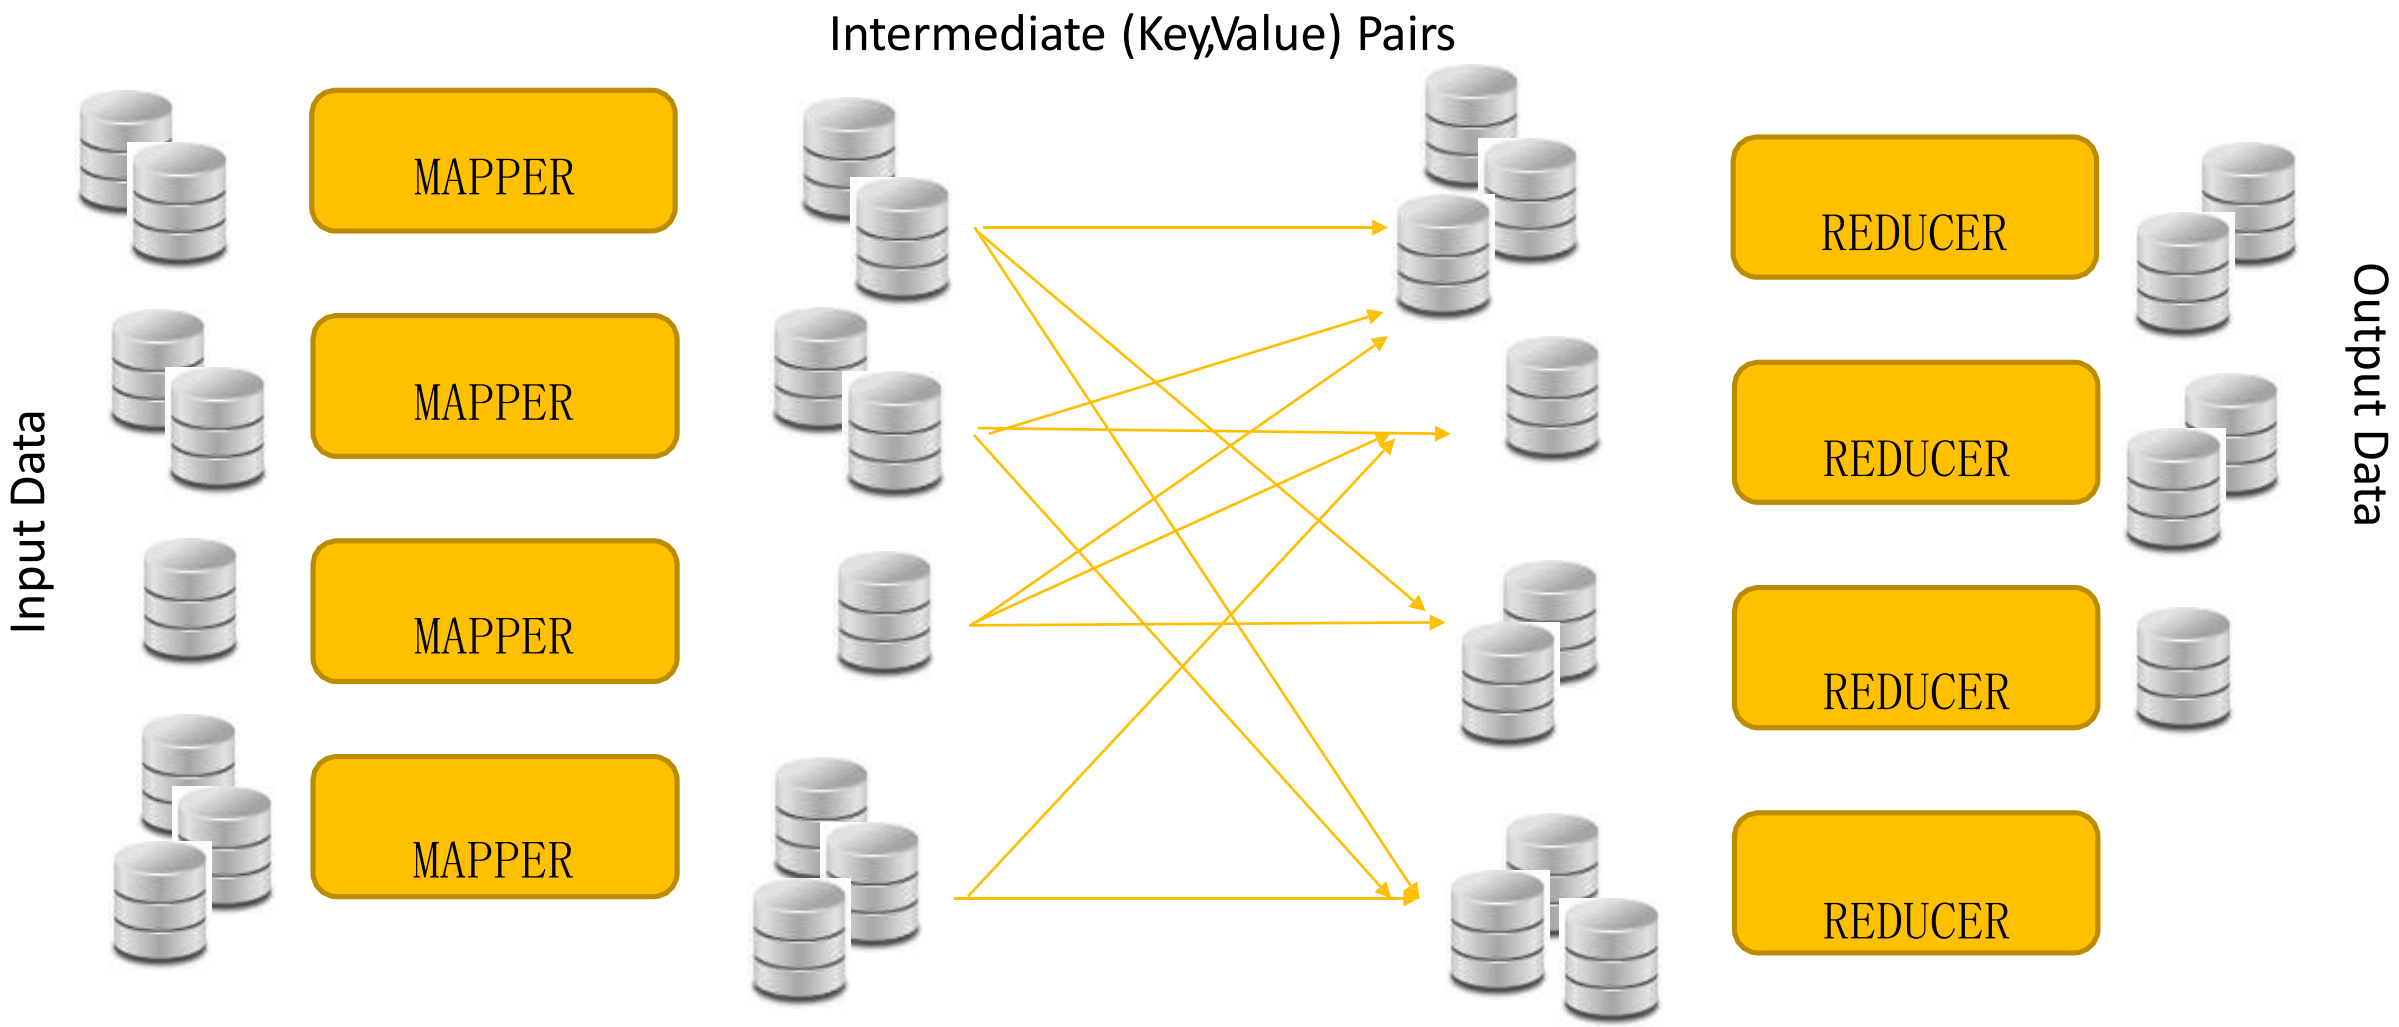

MAPREDUCE 1/3

Map

1) Key,Value(K1,V1)

Input (Data)

Input Split

2) list (K2, V2)

Key / Value out (intermediate values)

One list per local node

Can implement local Reducer (or Combiner)

MAPREDUCE 2/3

proschool
An ims Initiative

Map

1) Key,Value(K1,V1)

Input(Data)

Input Split

2) list (K2, V2)

Key / Value out
(intermediate values)

One list per local node

Can implement local

Reducer (or Combiner)

Shuffle/Sort

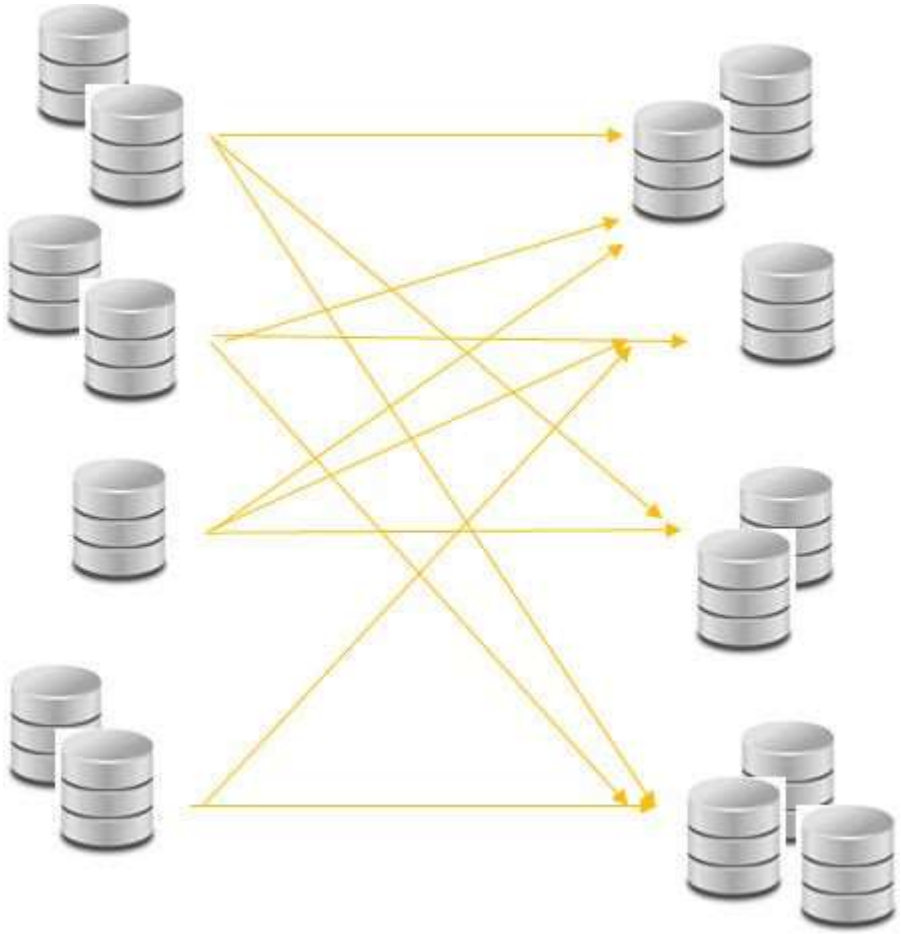

MAPREDUCE 3/3

proschool
An ims Initiative

Map

1) Key,Value(K1,V1)

Input (Data)

Input Split

2) list (K2, V2)

Key / Value out (intermediate values)

One list per local node

Can implement local Reducer (or Combiner)

Shuffle/Sort

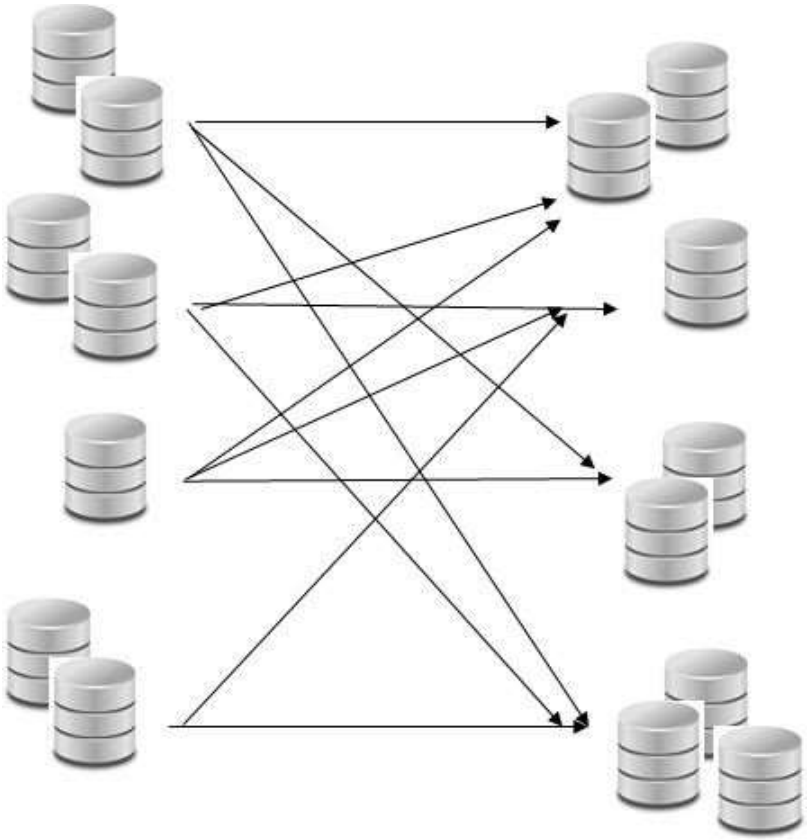

Reducer

1) (K2, list(V2))→

Shuffle / Sort phase precedes Reduce phase

Combines Map output into a list

2) list (K3, V3)

Usually aggregates intermediate values

Mapper

```
package mapper;

import java.io.IOException;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class WordCountMapper extends Mapper<LongWritable,Text,Text,LongWritable>  {
        @Override
        protected void map(LongWritable key,Text value,Context context)
        throws IOException,InterruptedException
        {
                LongWritable one = new LongWritable(1);

                String line = value.toString();
                String [] words = line.split(" ");

                for(String word : words)
                {
                        context.write(new Text(word),one);
                }}}
```

Reducer

```
package reducer;

import java.io.IOException;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class WordCountReducer extends Reducer<Text,LongWritable,Text,LongWritable> {

        @Override
        protected void reduce(Text key,Iterable<LongWritable> values,Context context) throws IOException,
InterruptedException{
                long count = 0;

                for(@SuppressWarnings("unused") LongWritable value : values)
                {
                        count++;
                }

                LongWritable finalCount = new LongWritable(count);
                context.write(key,finalCount);
        }
}
```

Driver

```
package driver;

import java.io.IOException;
import mapper.WordCountMapper;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.Job; import
org.apache.hadoop.mapreduce.lib.input.*;
import org.apache.hadoop.mapreduce.lib.output.*;
import reducer.WordCountReducer;

public class WordCountDriver {

@SuppressWarnings("deprecation")
public static void main(String[] args) throws IOException, ClassNotFoundException, InterruptedException{

        Configuration conf = new Configuration(); // standard to be followed
        Job job= new Job(conf);                   //standard to be followed
        job.setJarByClass(WordCountDriver.class);      //starting point
        job.setMapperClass(WordCountMapper.class);
        job.setReducerClass(WordCountReducer.class);
```

```
job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(LongWritable.class);

job.setOutputValueClass(LongWritable.class);
job.setOutputKeyClass(Text.class);


job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);


Path input = new Path(args[0]);
Path output = new Path(args[1]);

FileInputFormat.addInputPath(job,input);
FileOutputFormat.setOutputPath(job,output);

boolean isJobRunning = job.waitForCompletion(true);      //start the job and
                                                          wait for completion

System.exit(isJobRunning ? 0:1); // return functions for exit

}
}
```

Hadoop - Shell commands

- hadoop fs –cat file:///file2 (copies source paths to stdout)
- hadoop fs –mkdir /user/hadoop/dir1 /user/hadoop/dir2 (creates directory at the specified path).
- hadoop fs –copyFromLocal <fromDir> <toDir>
- hadoop fs –put <localfile> hdfs://nn.example.com/hadoop/hadoopfile
- sudo hadoop jar <jarFileName> <method> <fromDir> <toDir>
- hadoop fs –ls /user/hadoop/dir1
- Hadoop fs –cat hdfs://nn1.example.com/file1
- hadoop fs –get /user/hadoop/file <localfile>
- FOR OTHER COMMANDS REFER : https://hadoop.apache.org/docs/r2.7.2/hadoop-project-dist/hadoop-common/FileSystemShell.html
- **sudo refers to ‘super user’ i.e. run as super user(administrator)

Click on Icon for more commands

Appendix

1) https://hadoop.apache.org/docs/r1.2.1/mapred_tutorial.html
2) Hadoop – The Definitive Guide
3) https://hadoop.apache.org/docs/r2.7.2/hadoop-yarn/hadoop-yarn site/YARN.html