



- Hive Introduction
- Hive vs SQL
- Hive Architecture Overview
- Interacting HDFS using HIVE
- Hive Hands on Queries
  - Hive Tables
  - Importing Data
  - Querying Data
  - Joins in Hive



# Introduction and Uses

# Need for Hive



Hadoop is great for cost, but MapReduce is too difficult.

I'm deleting important data because it's too expensive to store it.



SQL on Hadoop makes Hadoop real and gives me scale that traditional SQL can't offer.



# Introduction to Hive



- Developed Hive to address traditional RDBMS limitations. 300+ PB of data under management<sup>(1)</sup>.
- 600+ TB of data loaded daily. 60,000+ Hive queries per day<sup>(2)</sup>. More than 1,000 users per day. Initial Apache release in April 2009.



HIVE	SQL
Built with an analytical focus in consideration	General purpose database language used for all types of
Doesn't support update and delete functions	SQL supports update and delete
Faster than SQL in reading & processing huge volumes of data	Slow in processing huge volumes of data
Hive has less number of built in functions as compared to SQL	More number of built in functions

## Hive Classic: Strengths and Challenges

 Familiar SQL Interface

 Economical Processing of Petabytes

 Hive Classic tied to MapReduce, leading to latency





# Hive Architecture and Components

# Physical Layout of Hive

- Warehouse directory in HDFS
  - e.g., **/user/hive/warehouse**
- Table row data stored in subdirectories of warehouse
- Partitions form subdirectories of table directories
- Actual data stored in flat files
  - Control char-delimited text, or Sequence Files



# Flow of Hive

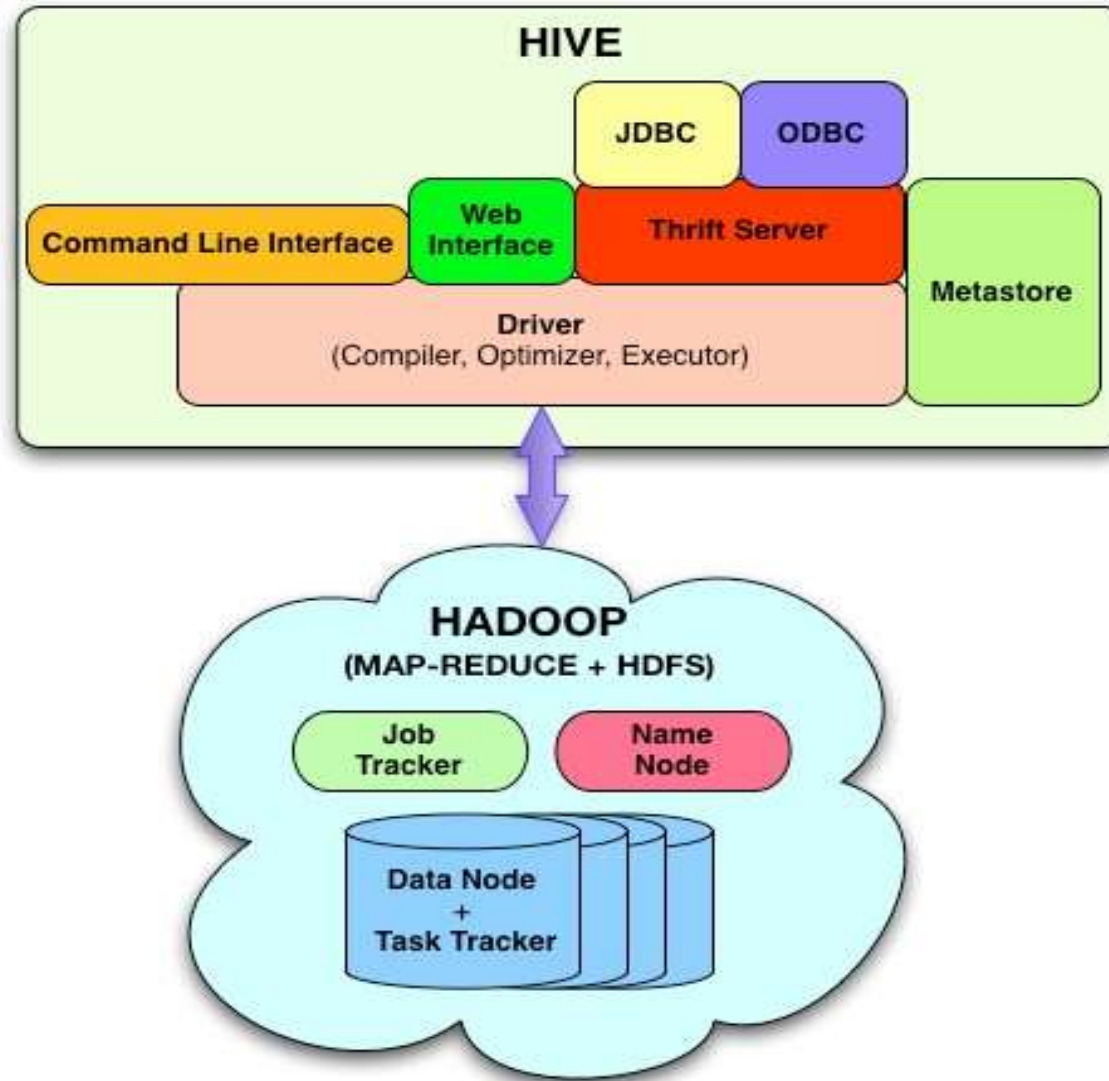


User ask a data in HIVE ( Kind of SQL)





# Hive Architecture



Hive Client




Hive Client


**Hive Client** supports different types of client applications in different languages for performing queries



**Thrift** is a software framework. Hive Server is based on Thrift, so it can support the request from all programming languages that support thrift

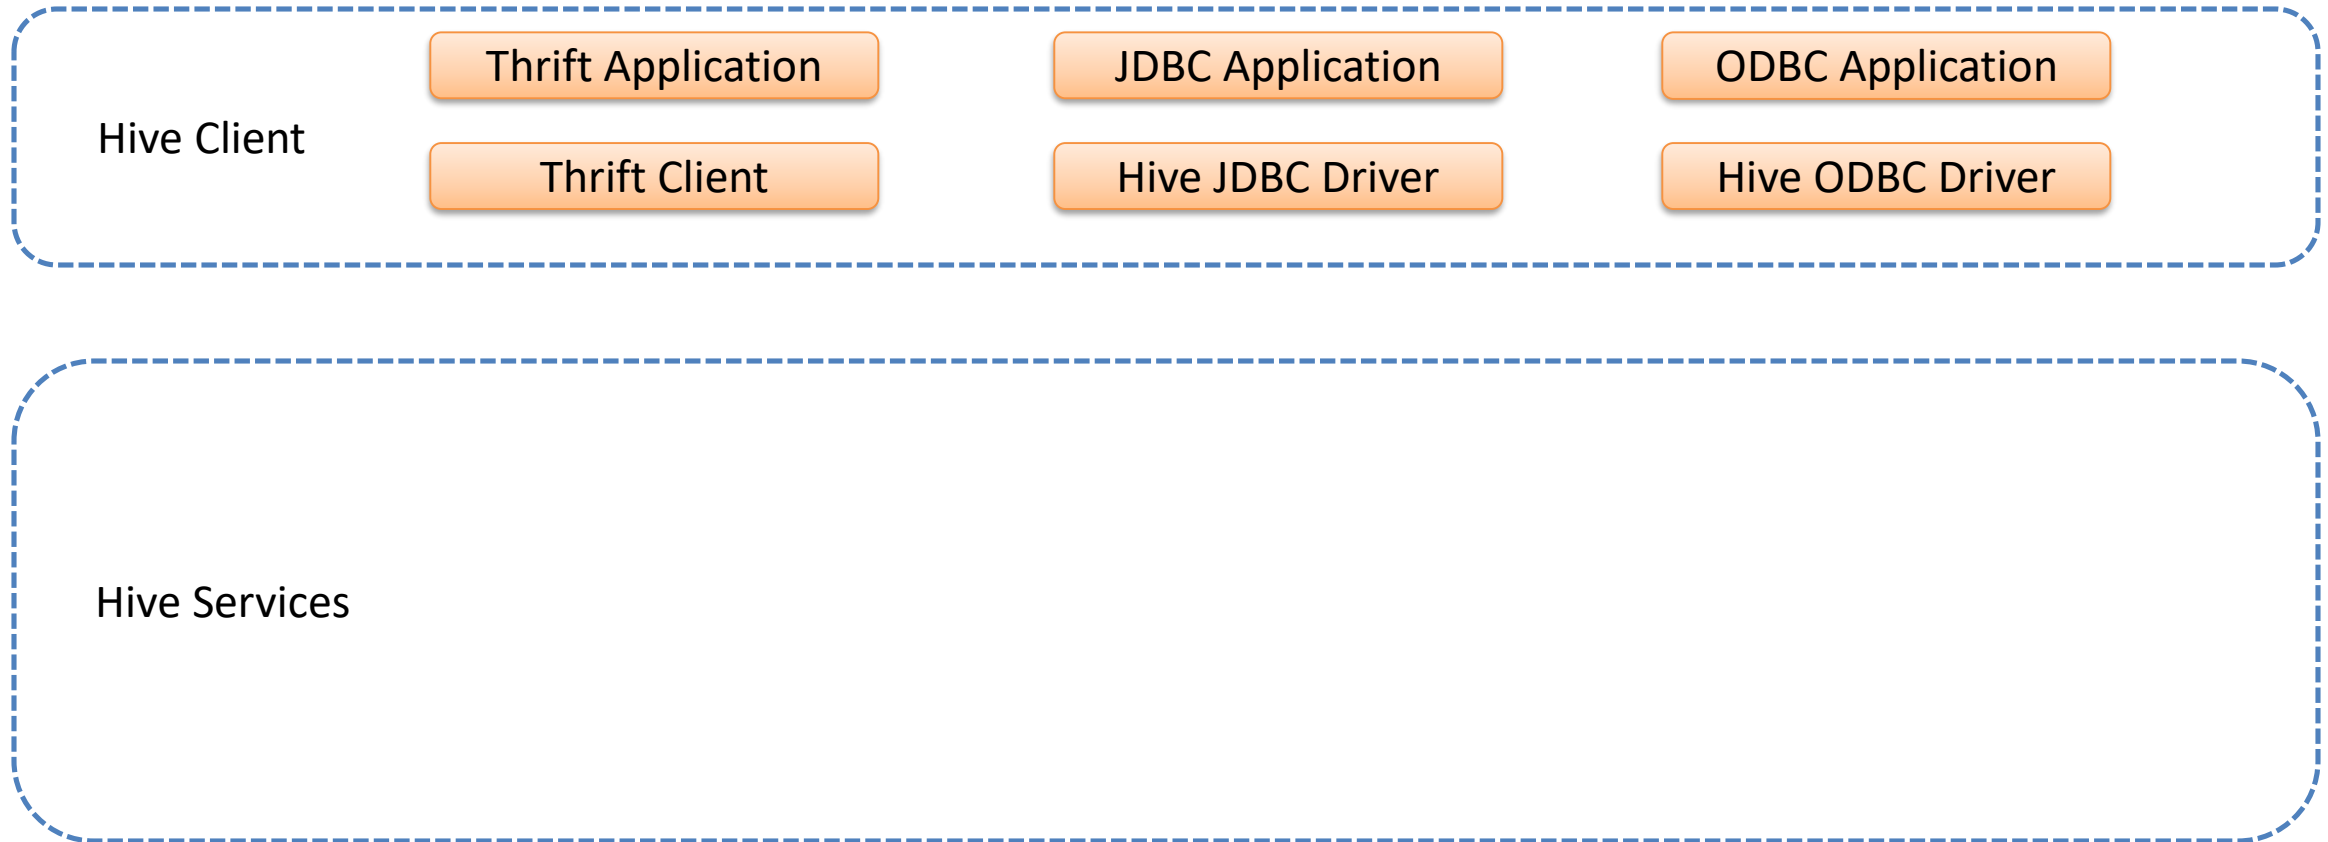


**JDBC** – Java Database Connectivity, helps to connect JDBC applications



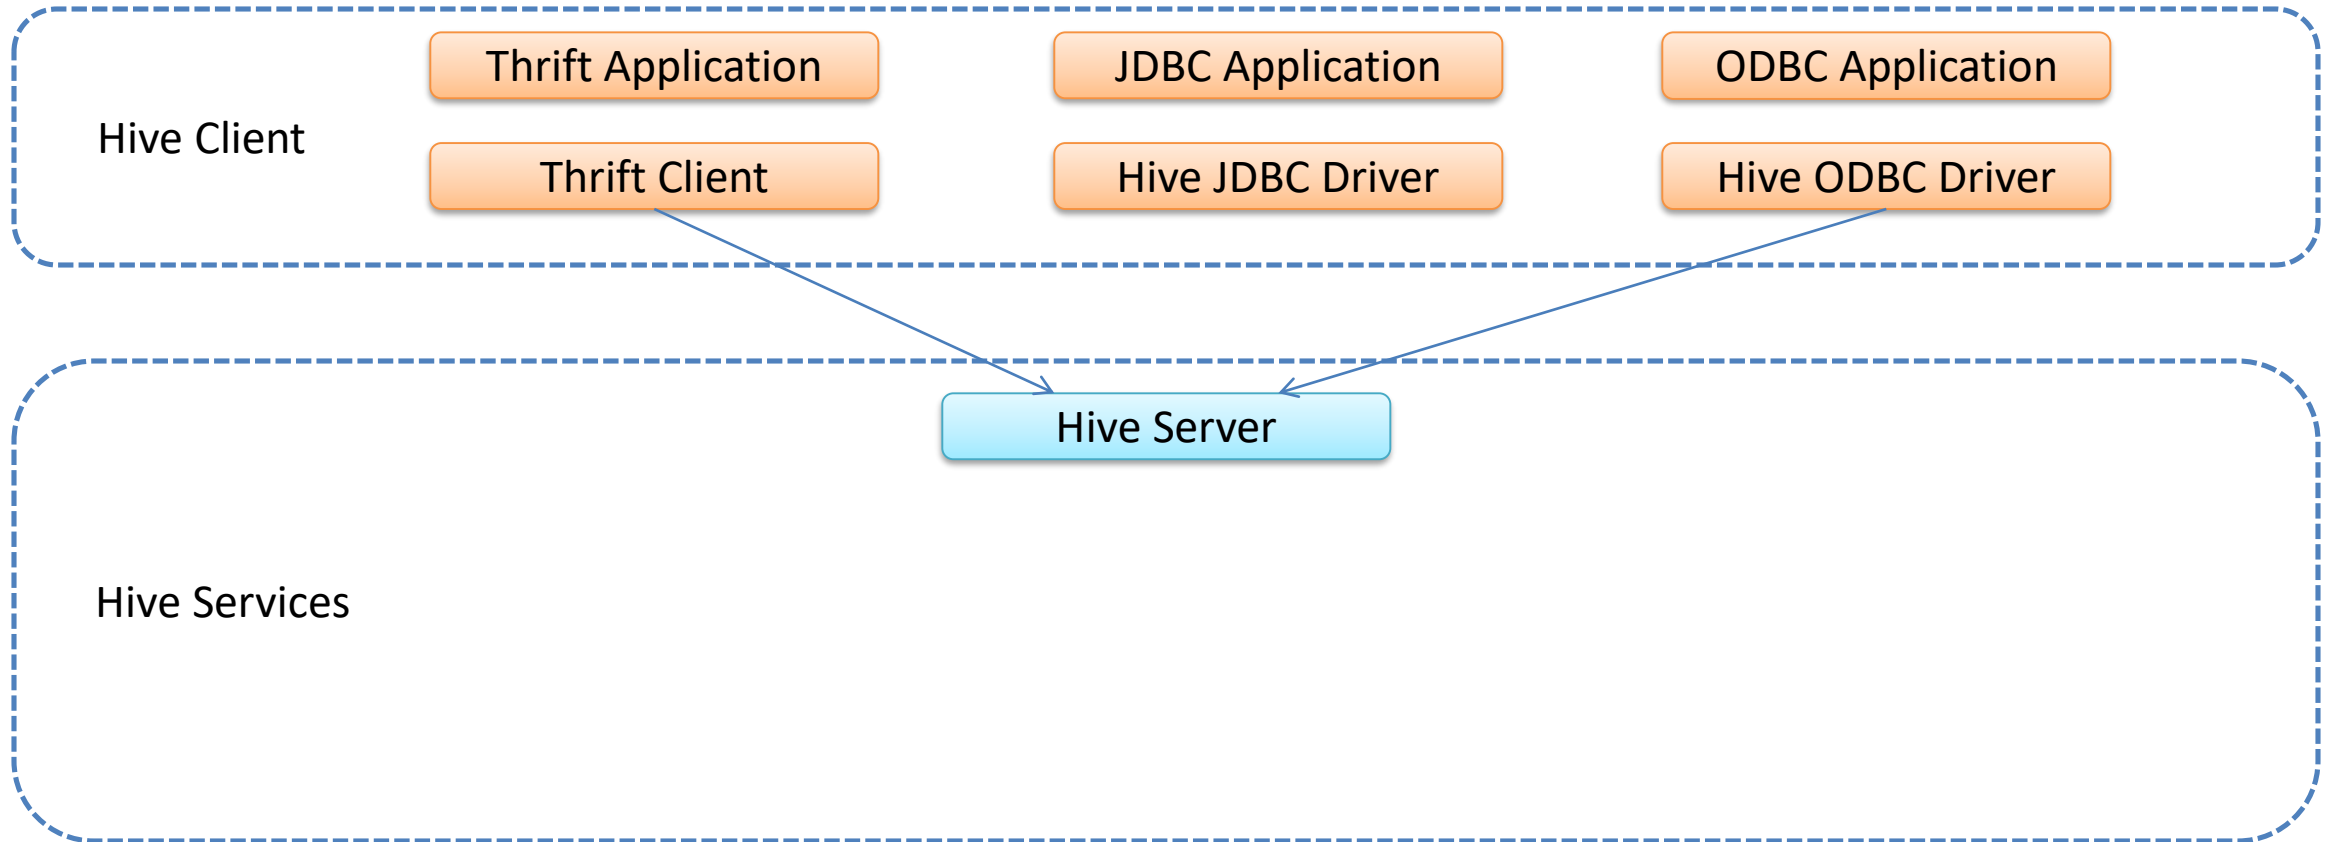
**ODBC**– Open Database Connectivity, helps to connect ODBC applications

# Hive Architecture



**Hive Supports various services**

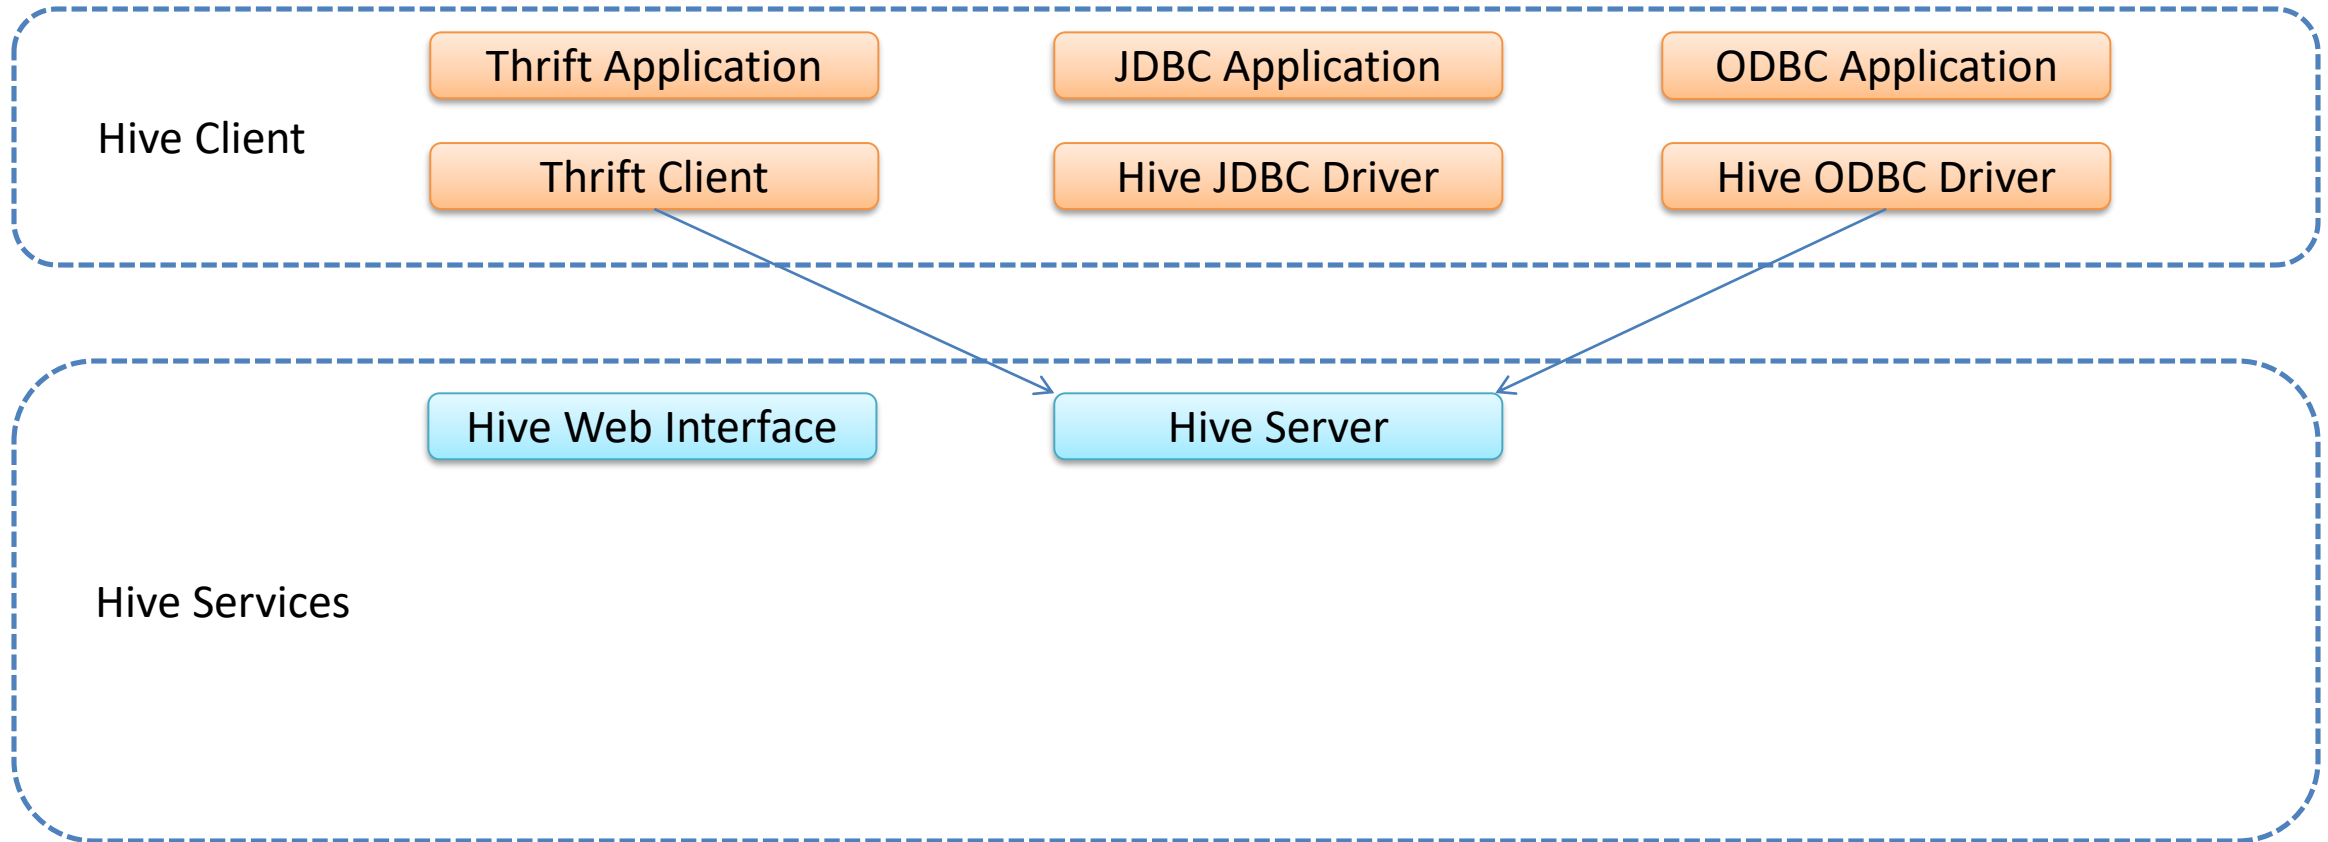
# Hive Architecture



All the client requests are submitted to **Hive Server**

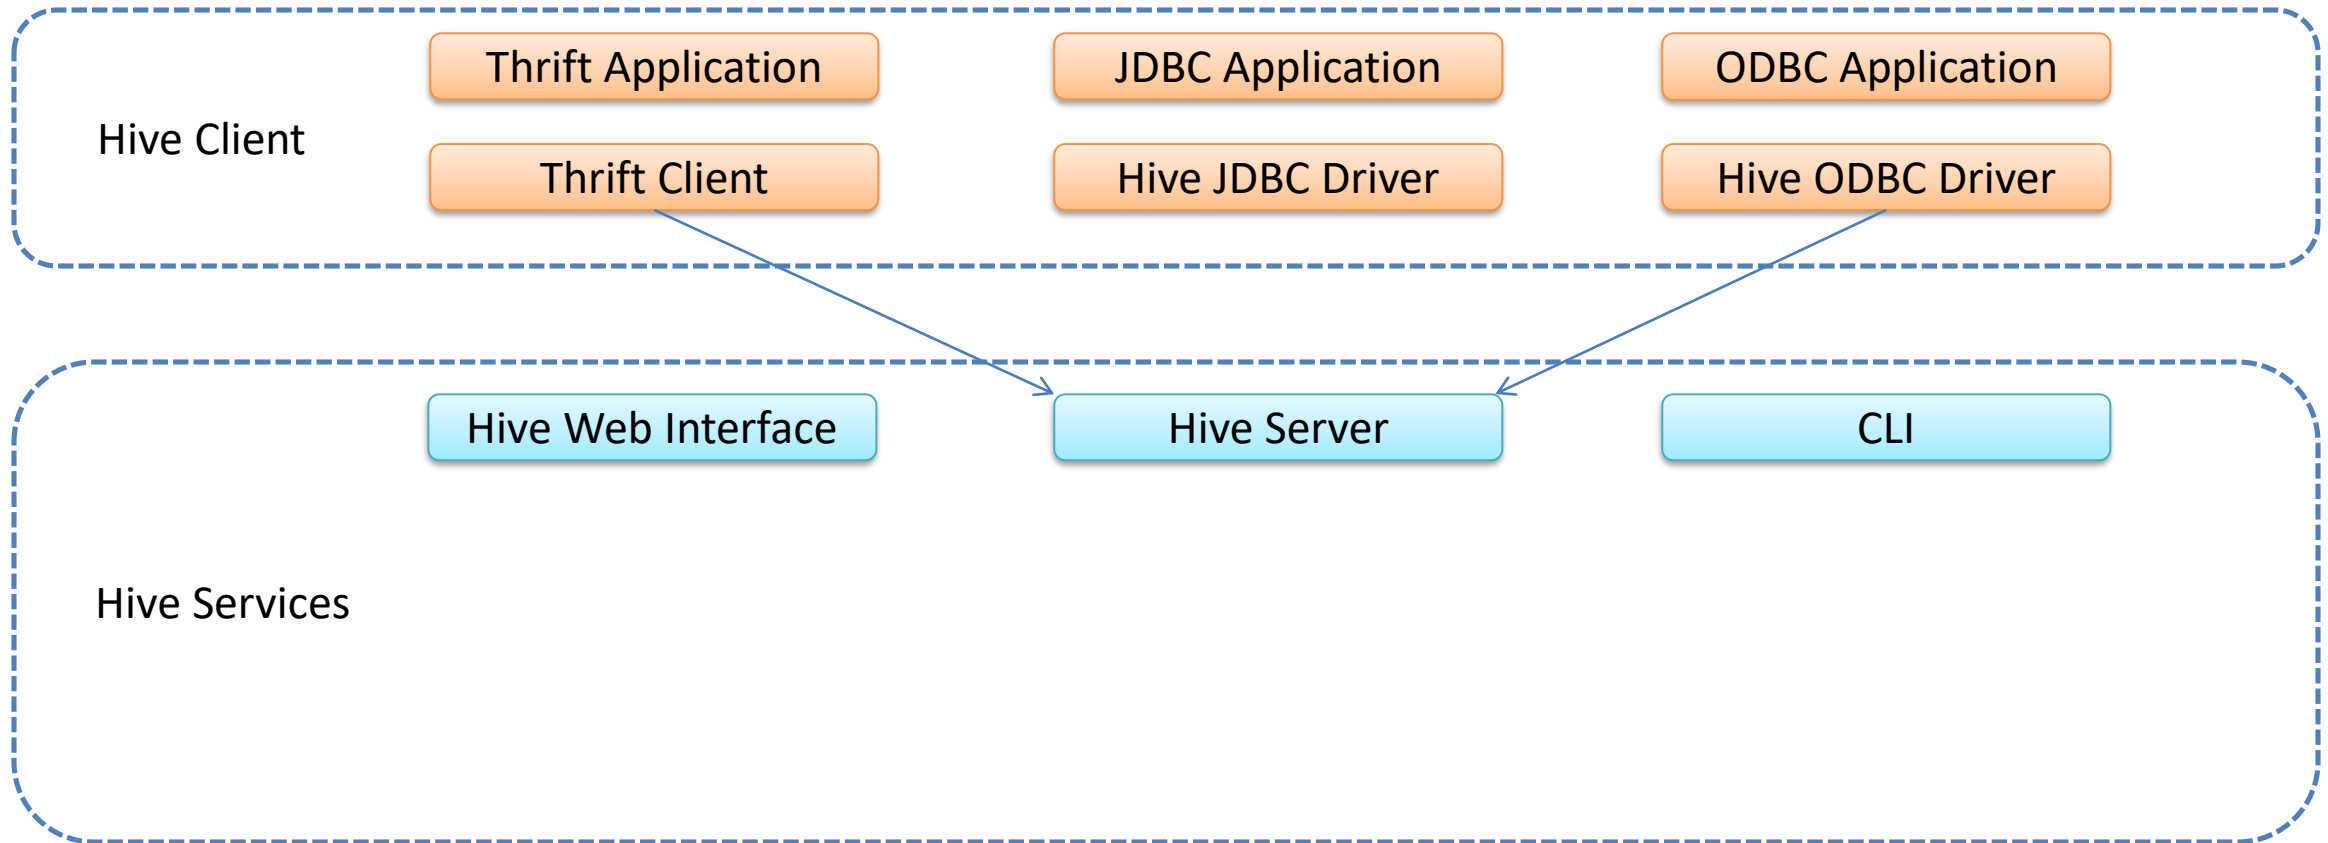


# Hive Architecture



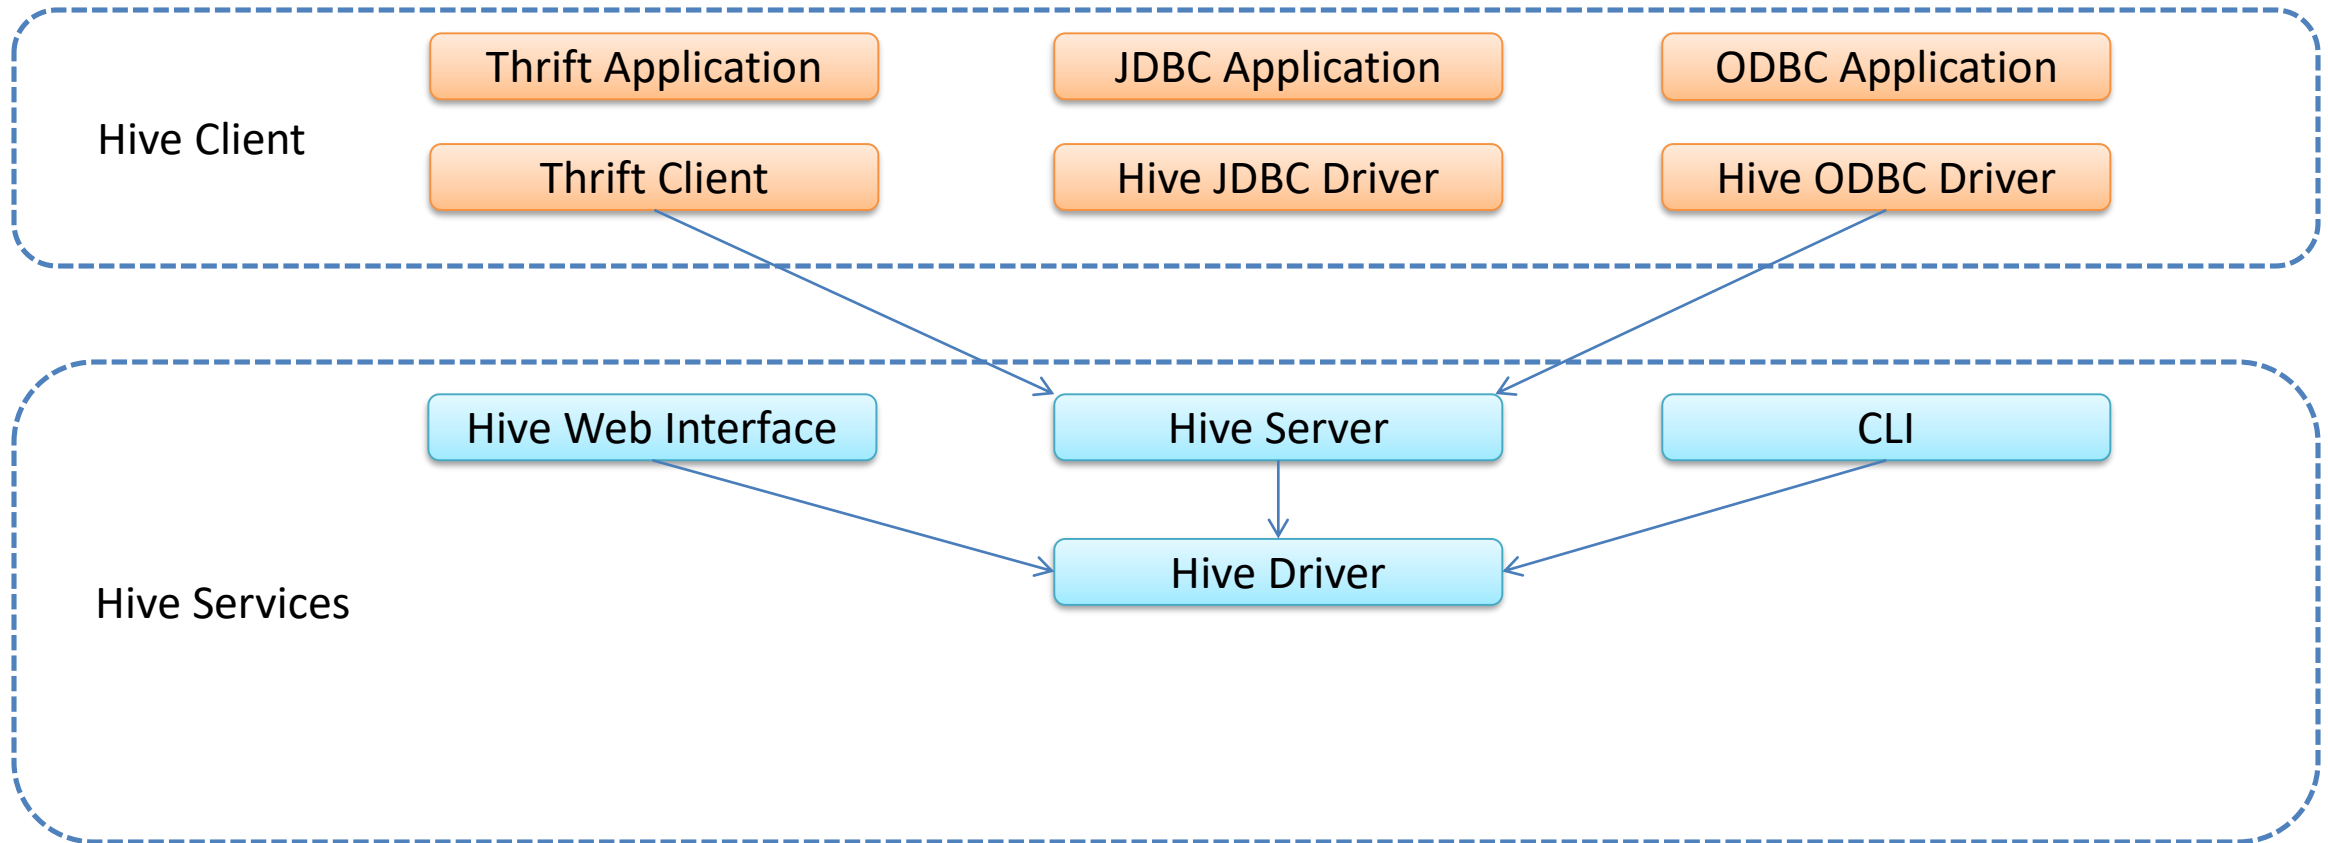
Web interface or GUI is provided to execute Queries

# Hive Architecture



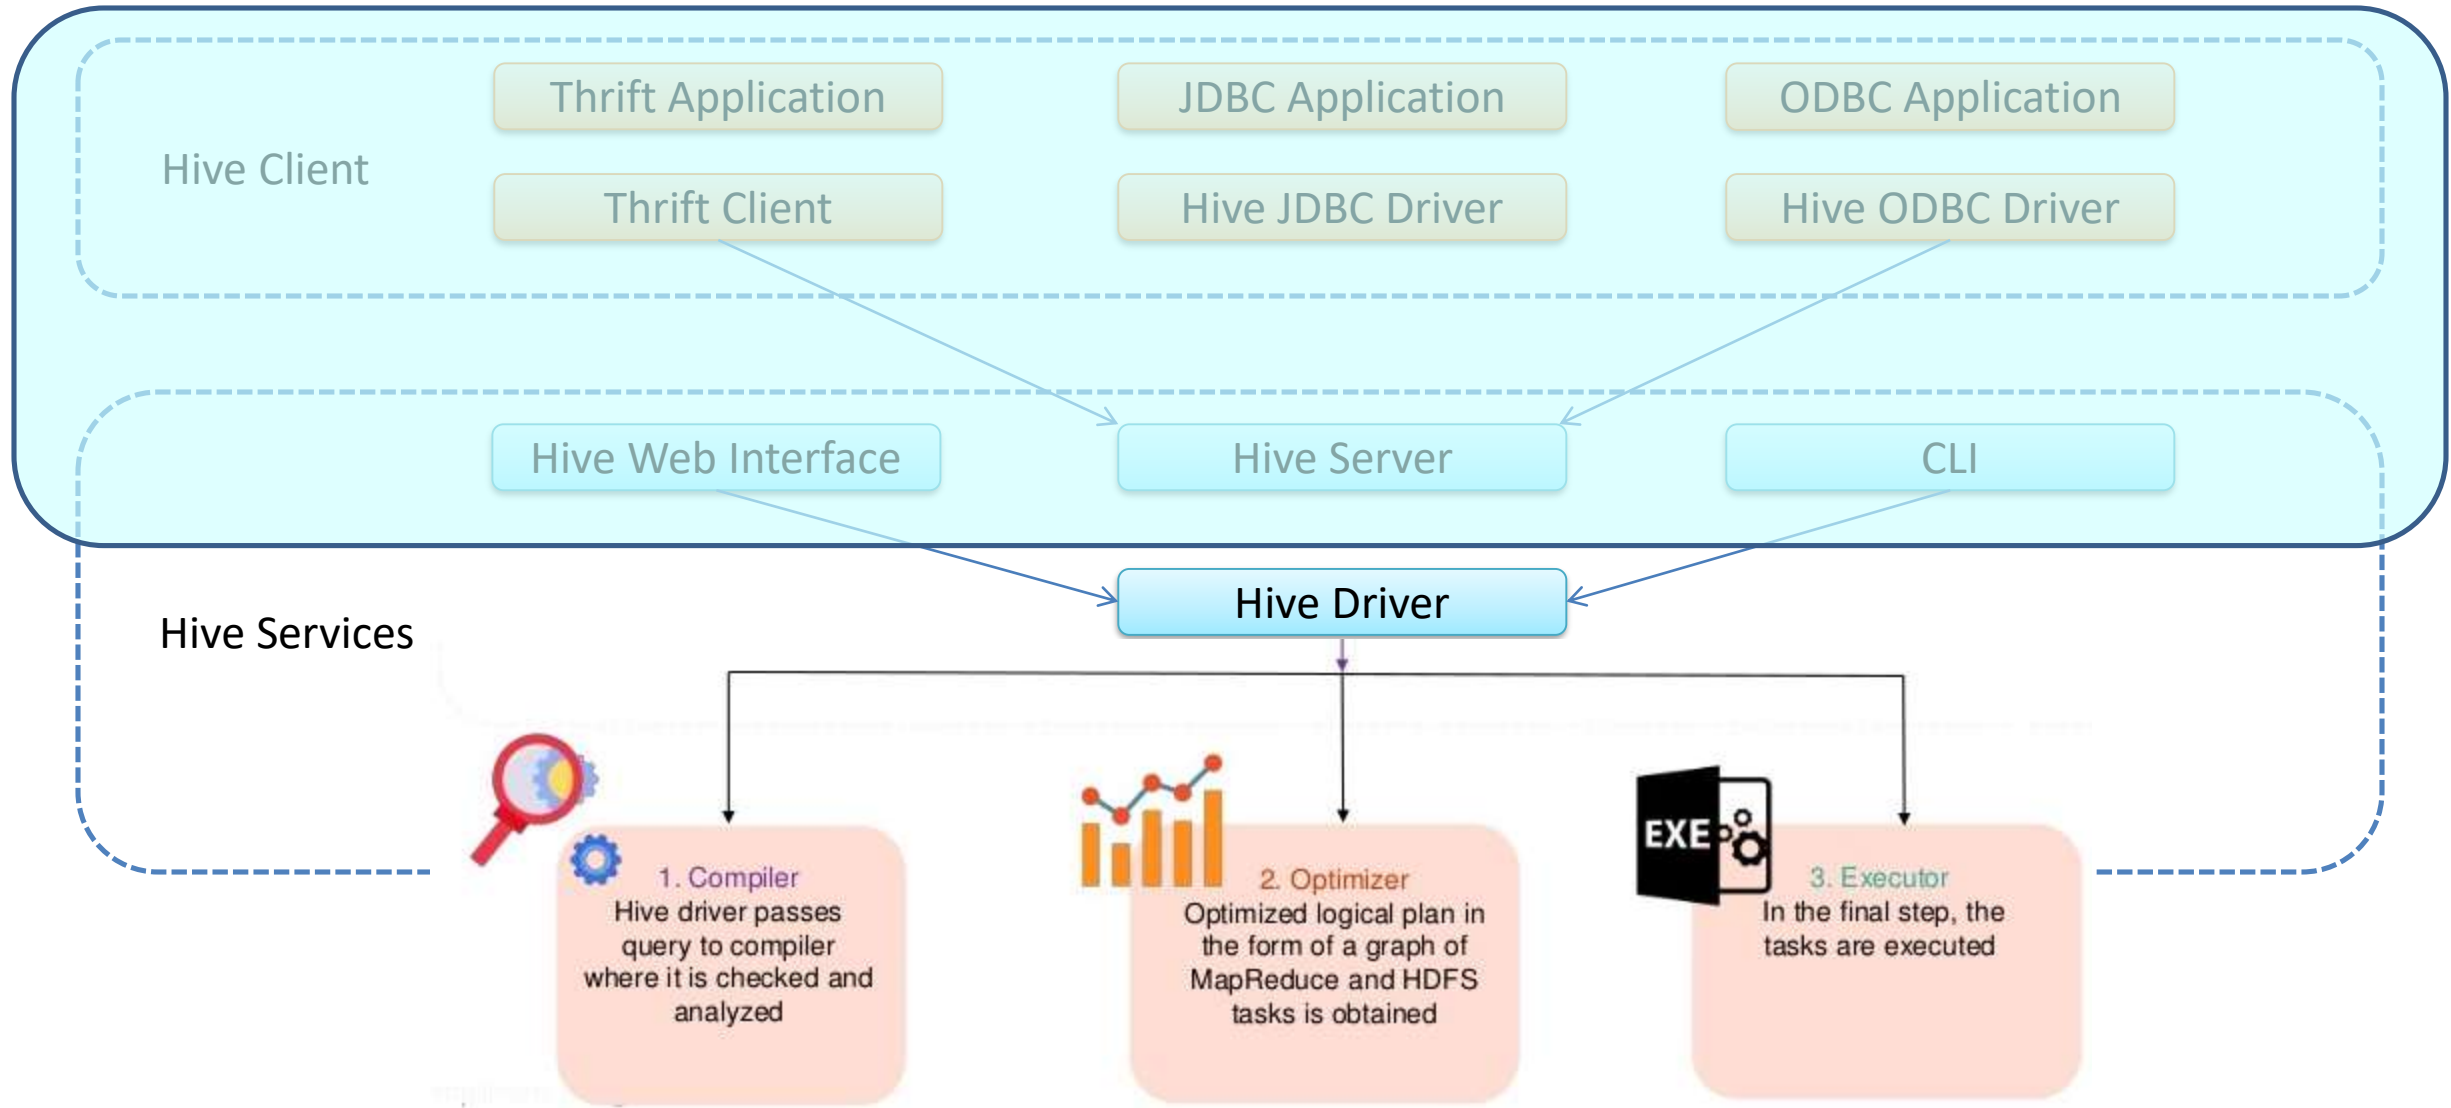
Commands are executed directly in CLI

# Hive Architecture



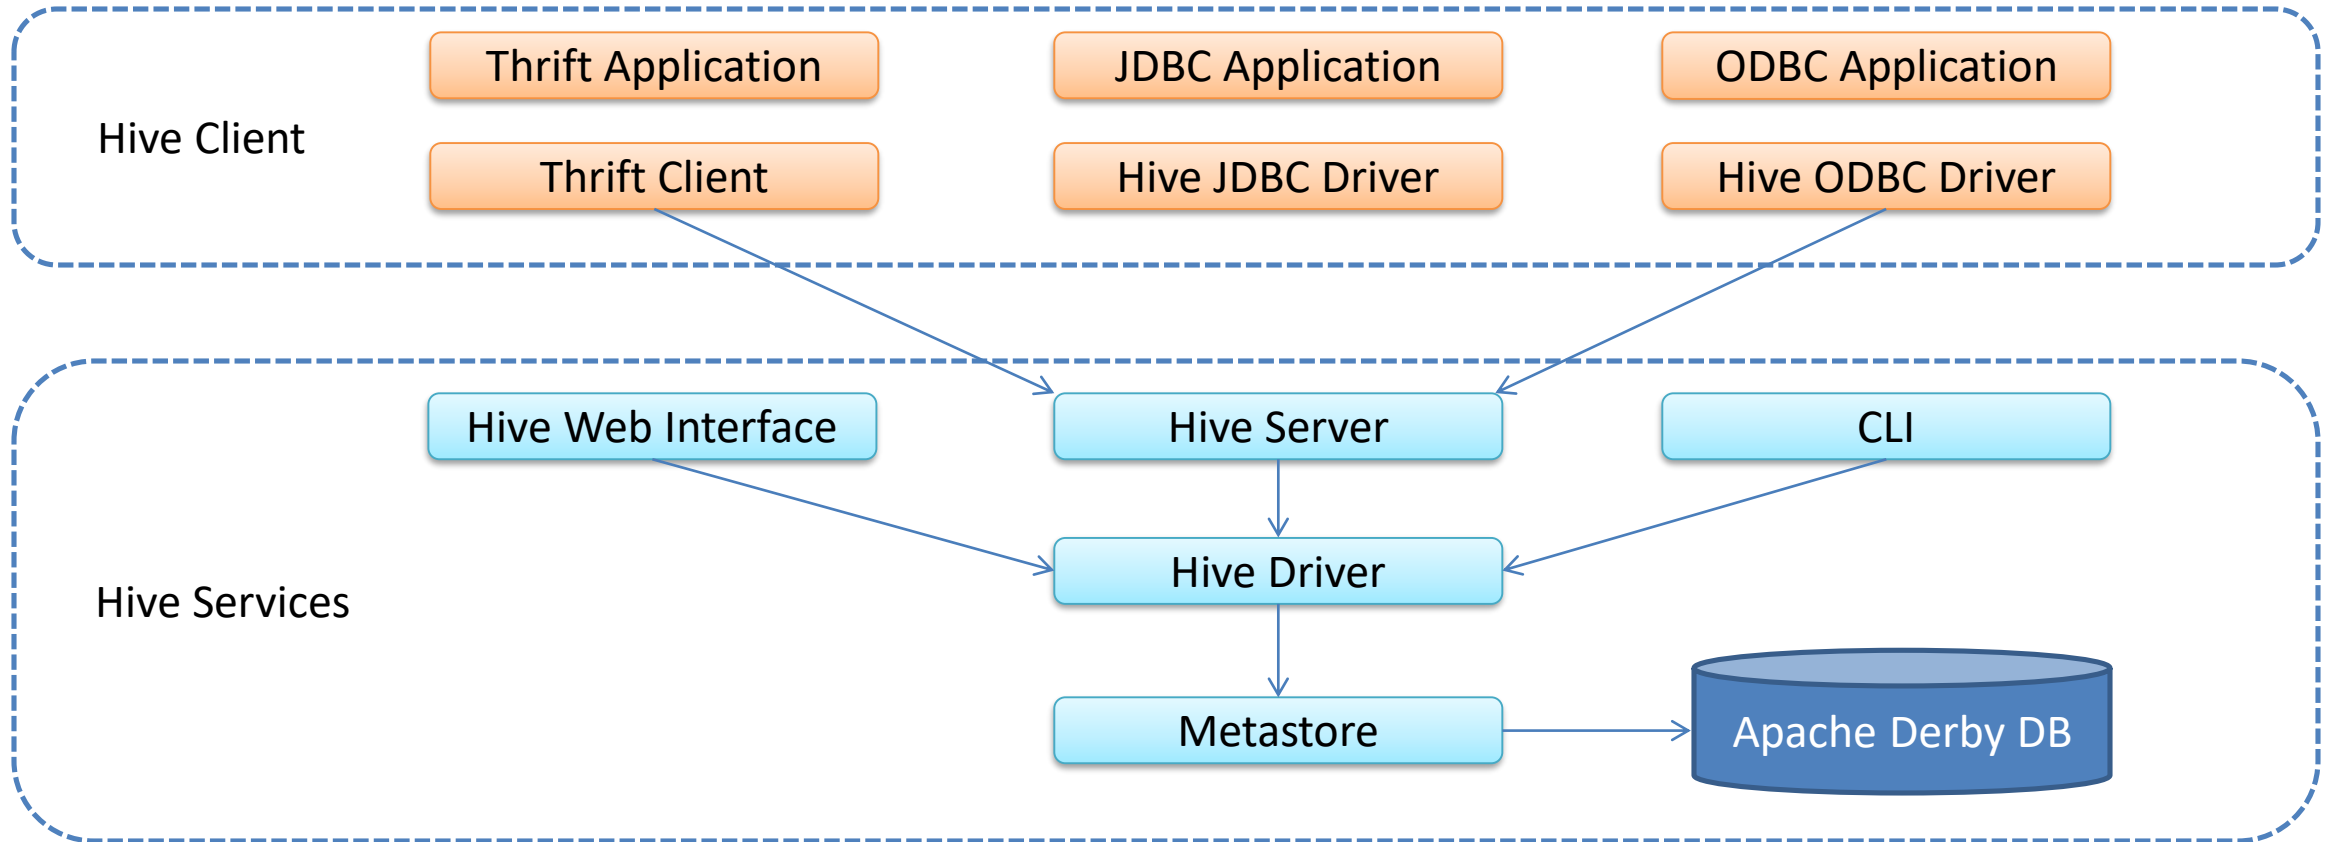
Hive Driver is responsible for all the queries submitted

# Hive Architecture



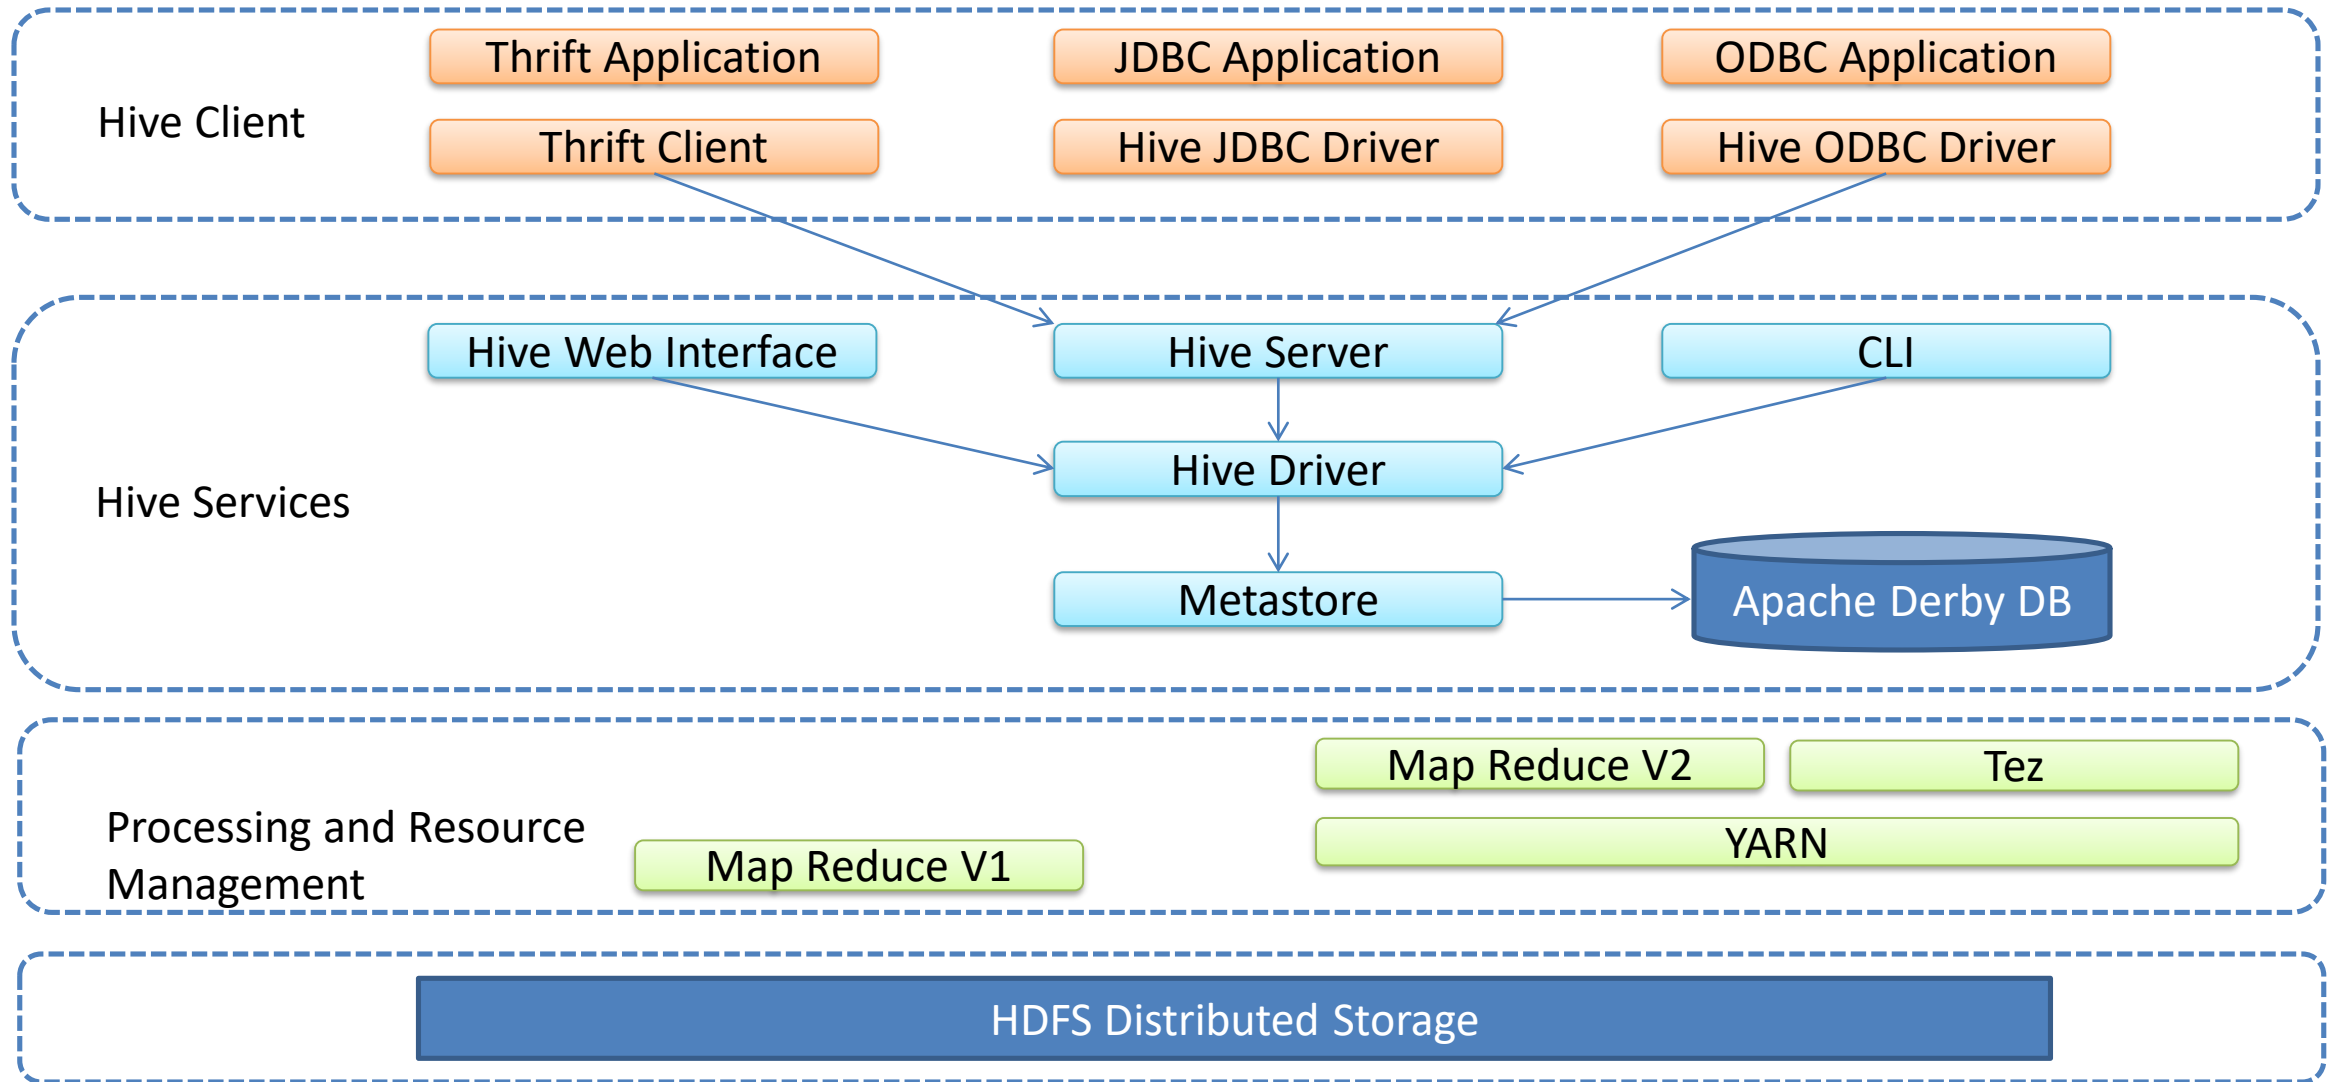
Hive Driver Performs 3 steps internally

# Hive Architecture



**Metastore** is repository for Hive Metadata. Stores metadata for Hive Tables


# Hive Architecture



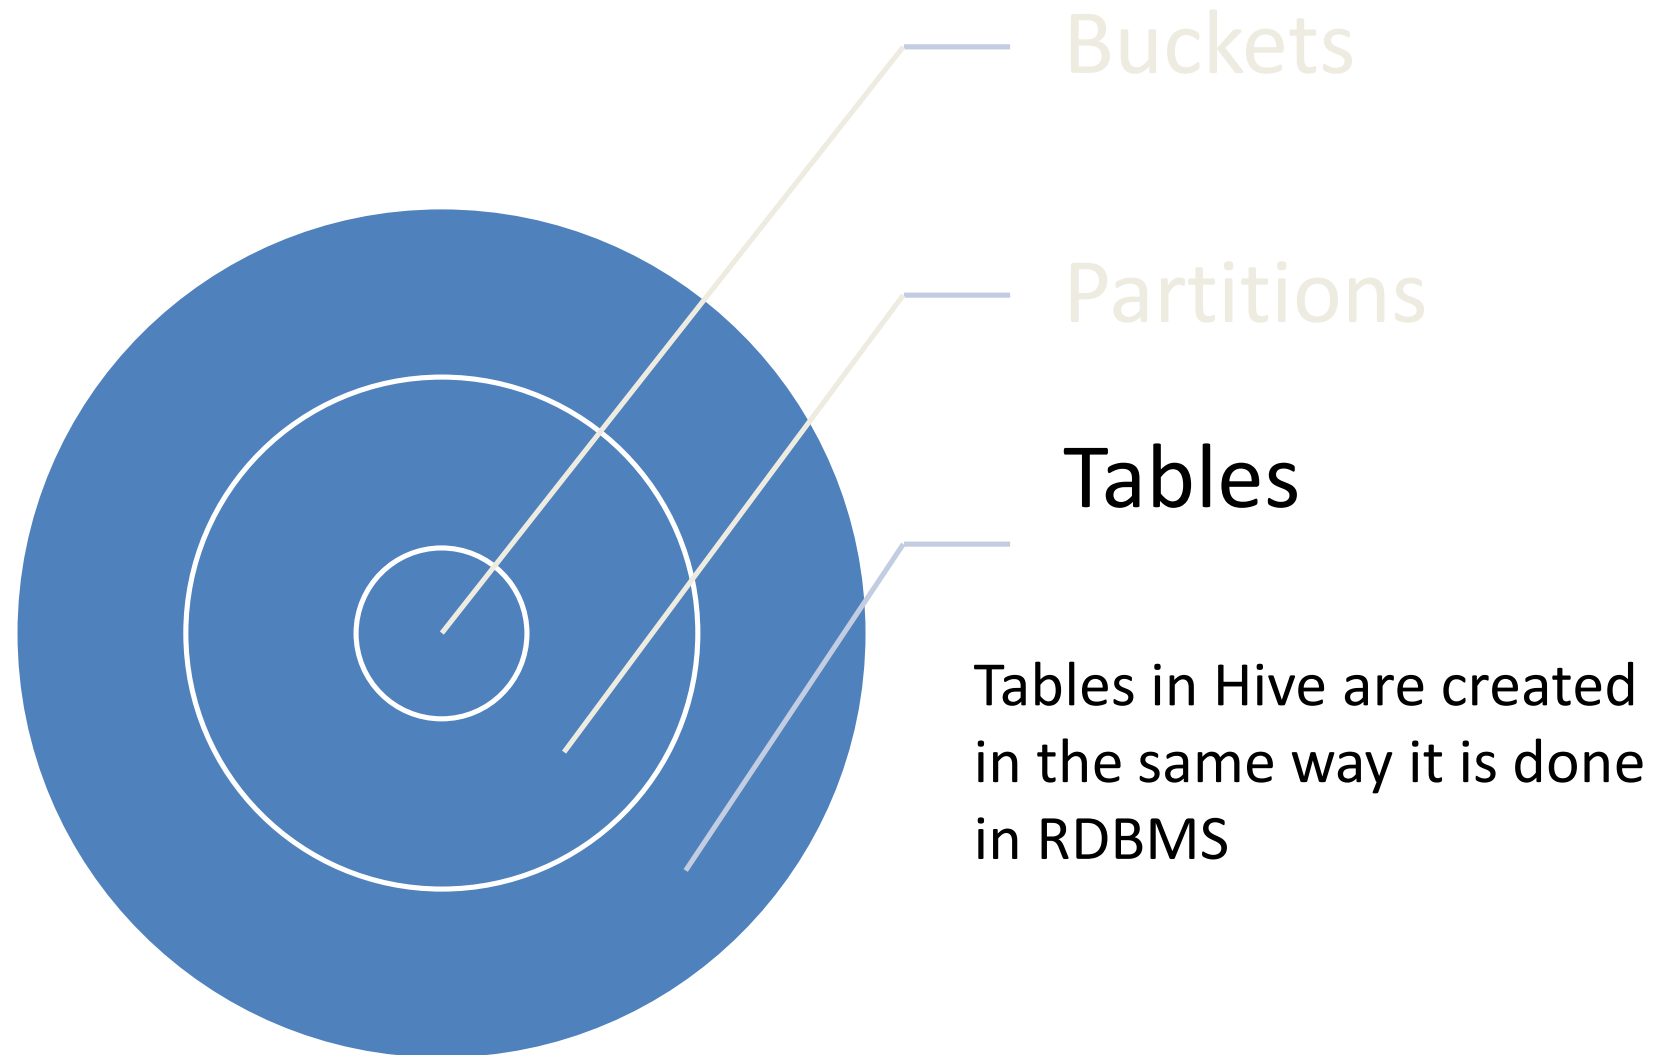
We will not go in details of Processing and Storage Steps

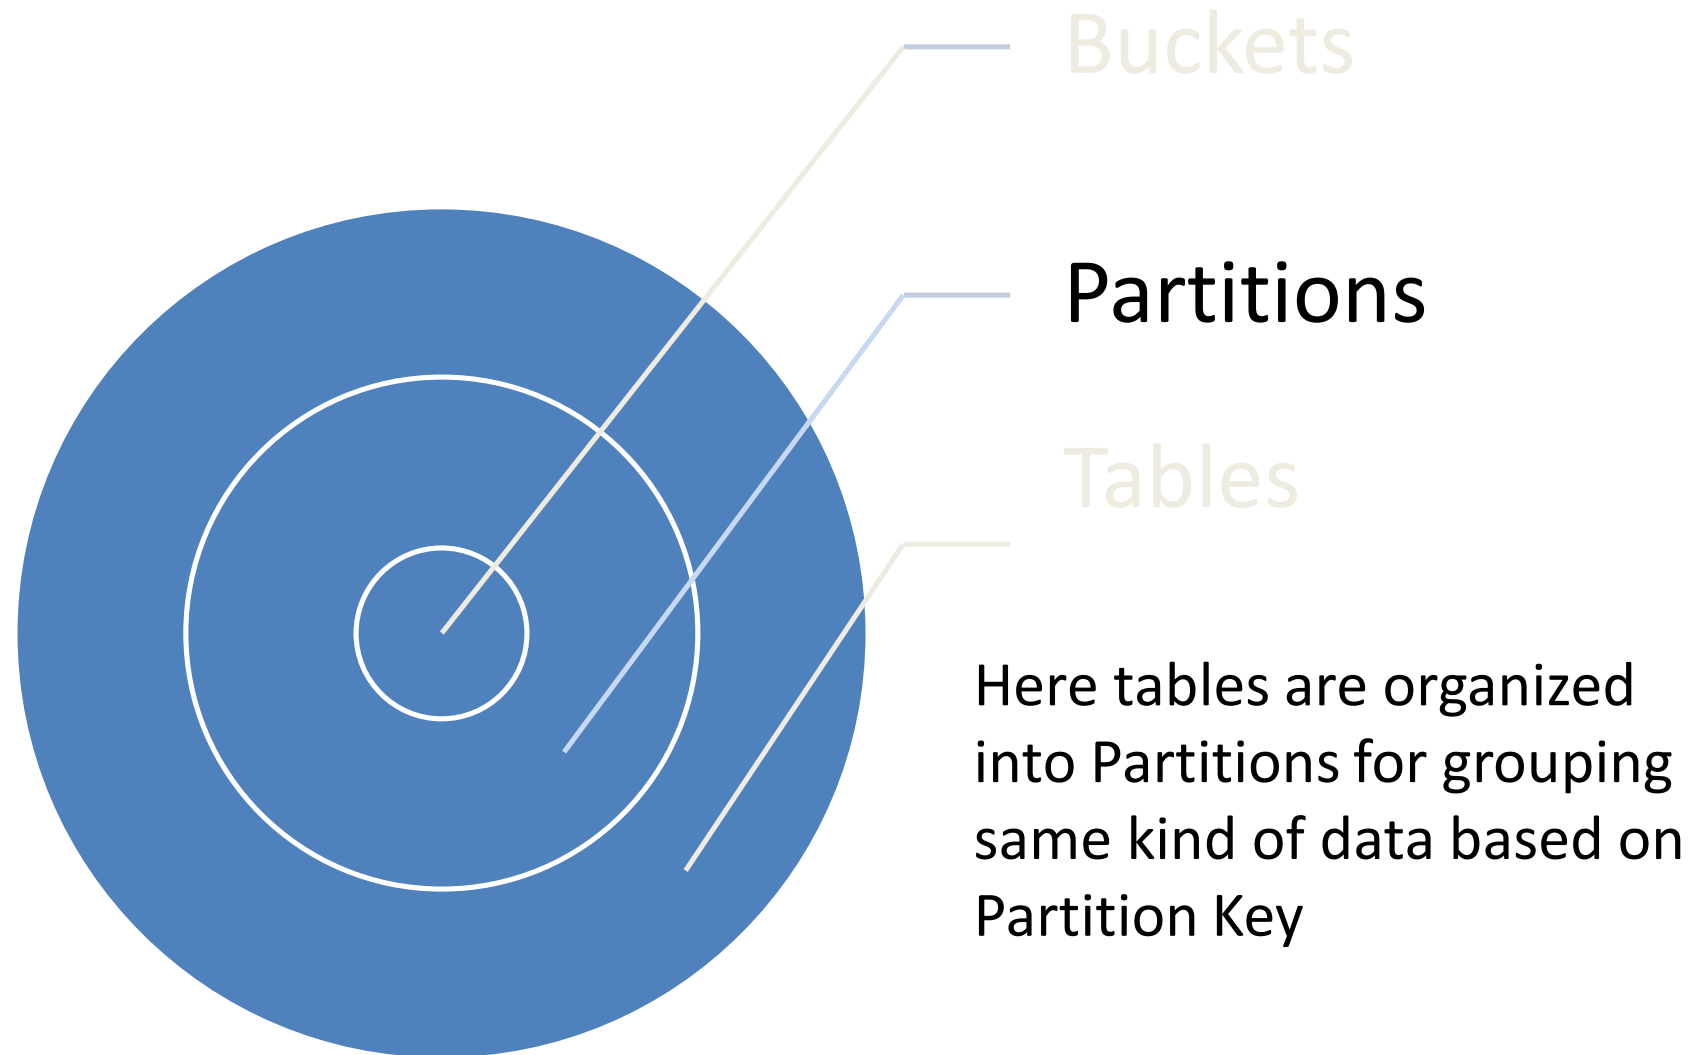


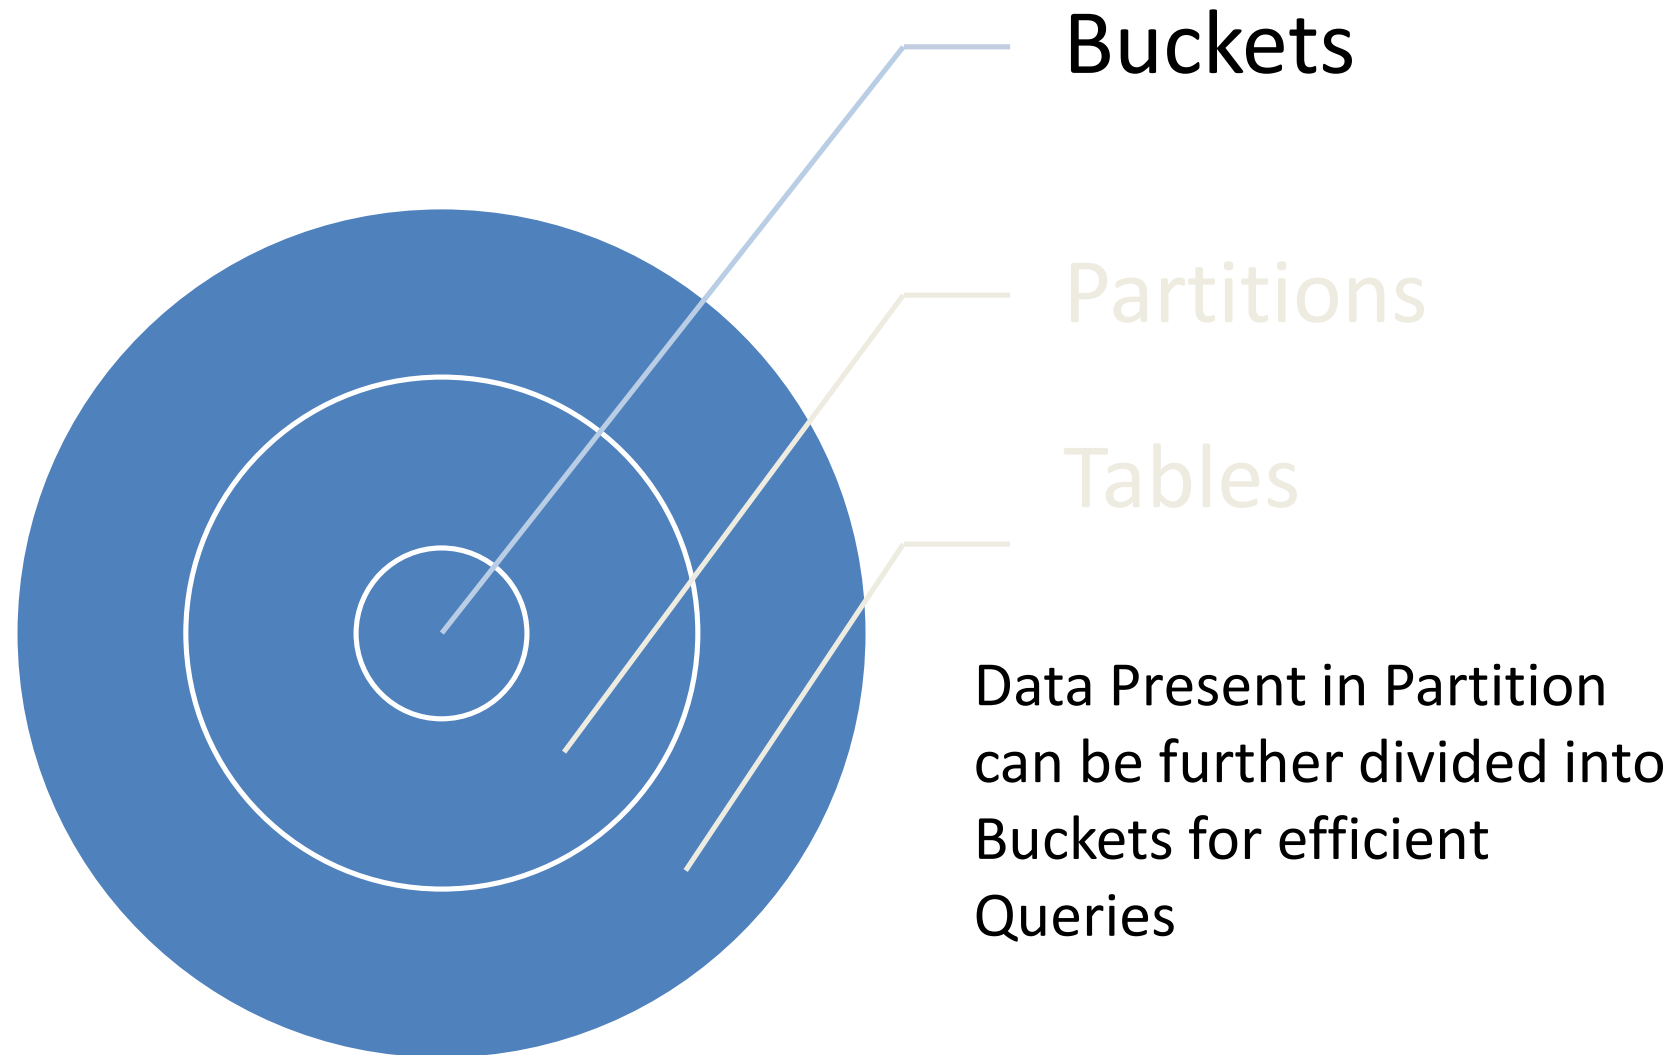
# Hive Data Modeling













# Beginning of HQL

# Create Database in Hive

First Make sure we have a database created, if not then follow below steps

Then check there is database Created

Following are the commands to create database in Hive

```
hive> CREATE DATABASE [IF NOT EXISTS] userdb;
```

```
hive> SHOW DATABASES;
```

Output is as below:

```
default  
userdb
```



# Create Table in Hive

Create table of following data type in Hive

```
Hive>CREATE TABLE IF NOT EXISTS employee ( eid int,  
      name String, salary String, destination String)  
COMMENT 'Employee details'  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '\t'  
LINES TERMINATED BY '\n'  
STORED AS TEXTFILE;
```

## Load Data Statement

Generally, after creating a table in SQL, we can insert data using the Insert statement. But in Hive, we can insert data using the LOAD DATA statement.

While inserting data into Hive, it is better to use LOAD DATA to store bulk records. There are two ways to load data: one is from local file system and second is from Hadoop file system.

Field Name	Data Type
Eid	int
Name	String
Salary	Float
Designation	string



# Managed Table Vs External Table

## Managed Table

- When we create a table in Hive by default it is a managed table
- The table is managed by hive
- When a data is loaded in a managed table its moved to warehouse directory
- The table data will be deleted when table is dropped

## External Table

- Managed by user
- Referring the data outside warehouse
- Applies only schema to the external file
- When table is dropped only the metadata/schema is deleted the file being referred is not deleted

# Load Data in Table

- Generally, after creating a table in SQL, we can insert data using the Insert statement.
- But in Hive, we can insert data using the LOAD DATA statement.
- While inserting data into Hive, it is better to use LOAD DATA to store bulk records.
- There are two ways to load data
  - one is from local file system
  - second is from Hadoop file system.

We are assuming the file is in local folder of following path

```
Hive>LOAD DATA LOCAL INPATH '/home/user/sample.txt'  
      OVERWRITE INTO TABLE employee;
```





# Adding Partition

- Consider file Contains data

```
id, name, dept, yoj  
1, gopal, TP, 2012  
2, kiran, HR, 2012  
3, kaleel, SC, 2013  
4, Prasanth, SC, 2013
```

- Consider file is divided in the 2 files by year of joining

/tab1/employeedata/2012/file2

```
1, gopal, TP, 2012  
2, kiran, HR, 2012
```

/tab1/employeedata/2013/file3

```
3, kaleel, SC, 2013  
4, Prasanth, SC, 2013
```

- Following command is used to add a partition, give a try to rename also

```
hive> ALTER TABLE employee > ADD PARTITION  
(year='2012') > location '/2012/part2012';
```





# Partitioning and Bucketing

- What is Partitions?
  - ✓ A way of dividing a table into coarse-grained parts based on the value of a partition column, such as a date.
  - ✓ Using partitions can make it faster to do queries on slices of the data.
- A Table may be partitioned in multiple dimensions
  - ✓ it by location so as to get to know which location has what all profiles.
- Partitions are defined at table creation time using the PARTITIONED BY clause, which takes a list of column definitions



# Partitions

- At the filesystem level ,partitions are nested subdirectories of the table directories.
- For example we want to partition data of sales of an company by date and city :

The files created after using **PARTITIONED BY**(dat STRING,city STRING) are :

```
| --dat = 2016-12-12/  
| | --city=Pune/  
| | | --file1  
| | | -file2....
```

```
hive>SHOW PARTITIONS sales ;  
dat=2016-12-12/city=Pune  
dat=2016-12-12/city=Mumbai  
dat=2016-12-13/city=Pune  
dat=2016-12-13/city=Mumbai
```



- **Partition columns** : Columns defined in the PARTITIONED BY clause are full fledged table columns known as Partition columns
- Datafiles do not contain values for these columns ,since they are derived from directory names.
- ✓ Partition columns can be used in SELECT statements in usual way .
- ✓ Hive performs input pruning to scan only the relevant partitions. For eg : `SELECT dat, company_name FROM sales WHERE city = 'Pune'`
- ✓ Only scans the files containing city = Pune
- ✓ The query returns the values of the dat partition columns ,which Hive reads from the directory names since they are not in datafiles



- Bucketing is a technique for decomposing data sets into more manageable parts
- It imposes extra structure on the table.
- For example :

```
hive> CREATE TABLE weblog (url STRING, source_ip STRING)  
> PARTITIONED BY (dt STRING, user_id INT);  
hive> FROM raw_weblog  
> INSERT OVERWRITE TABLE page_view PARTITION(dt='2016-06-08', user_id)  
> SELECT server_name, url, source_ip, dt, user_id;
```
- If use dynamic partitioning, above commands might fail because by default Hive limits the maximum number of dynamic partitions that may be created to prevent the extreme case that overwhelm the filesystem's ability.



- Instead, if we bucket the weblog table and use user\_id as the bucketing column, the value of this column will be hashed by a user-defined number into buckets.

```
hive > CREATE TABLE weblog (user_id INT, url STRING, source_ip  
STRING)
```

```
> PARTITIONED BY (dt STRING)
```

```
> CLUSTERED BY (user_id) INTO 96 BUCKETS;
```

- Records with the same user\_id will always be stored in the same bucket, each bucket will have many users



- Creating Bucketed Table :  
CREATE TABLE bucketed\_users (emp\_id INT, name STRING)  
CLUSTERED BY (id) INTO 8 BUCKETS;
- Inserting Data :  
INSERT OVERWRITE TABLE  
bucketed\_users SELECT \* FROM  
users;
- To populate the bucketed table, we need to set the hive.enforce.bucketing property to true





# Bucket : working

- Hive bucketing works by hashing the value and reducing modulo number of buckets
- EG : hive>**SELECT \* FROM STUDENTS**
  - 1 Ram
  - 2 Rajiv
  - 3 Rohit
  - 4 Nikhil

Each bucket is a file in the table or partition

directory Hive > dfs -ls

/user/hive/training\_hive/bucketed\_users Output :

00000\_0

00000\_1

00000\_2

..



# Hive Partition : Features

- Hive Partitioning refers to dividing large amount of data into number pieces of folders based on table columns value.
- Hive Partition is often used for distributing load horizontally, this has performance benefit, and helps in organizing data in a logical fashion.
- We can perform partition on any number of columns in a table by using hive partition concept.
- Partitioning works better when the cardinality of the partitioning field is not too high .
- Partition is not solving responsiveness problem in case of data skewing towards a particular partition value.
  - ✓ Grouping of customers of Walmart will take long time compared to grouping of customers of Big Bazaar.



# Hive Bucketing : Features

- Hive bucketing is responsible for dividing the data into number of equal parts
- We can perform Hive bucketing optimization only on one column only not more than one.
- The value of this column will be hashed by a user-defined number into buckets.
- Bucketing works well when the field has high cardinality and data is evenly distributed among buckets
- Bucketing on the other hand, will result with a fixed number of files, since you do specify the number of buckets.
  - ✓ But, if you use let's assume 512 buckets and the field you're bucketing on has a low cardinality (for instance, it's a Australia state, so can be only 100 different values). Then we will have 100 buckets with data, and 412 buckets with no data.
- Due to equal volumes of data in each partition, joins at Map side will be quicker.



- When querying a partition table when a query is issued only the required partitions are scanned.
- For Bucketed tables ,we need to provide an hint to the query using TABLESAMPLE clause ,if we want to scan few particular buckets else the whole set of files would be scanned.
- For example :
  - If we want to choose data from only 5 buckets out of 50 buckets then we use
  - SELECT \* FROM table1 TABLESAMPLE (5 out of 50) WHERE dat = '2016-12-11'



# Comparing Queries with SQL and Hive

Function	MySQL	Hive QL
Retrieving information	SELECT from_columns FROM table WHERE conditions;	SELECT from_columns FROM table WHERE conditions;
All values	SELECT * FROM table;	SELECT * FROM table;
Some values	SELECT * FROM table WHERE rec_name = "value";	SELECT * FROM table WHERE rec_name = "value";
Multiple criteria	SELECT * FROM table WHERE rec1="value1" AND rec2="value2";	SELECT * FROM TABLE WHERE rec1 = "value1" AND rec2 = "value2";
Selecting specific columns	SELECT column_name FROM table;	SELECT column_name FROM table;
Retrieving unique output records	SELECT DISTINCT column_name FROM table;	SELECT DISTINCT column_name FROM table;
Sorting	SELECT col1, col2 FROM table ORDER BY col2;	SELECT col1, col2 FROM table ORDER BY col2;
Sorting backward	SELECT col1, col2 FROM table ORDER BY col2 DESC;	SELECT col1, col2 FROM table ORDER BY col2 DESC;
Counting rows	SELECT COUNT(*) FROM table;	SELECT COUNT(*) FROM table;
Grouping with counting	SELECT owner, COUNT(*) FROM table GROUP BY owner;	SELECT owner, COUNT(*) FROM table GROUP BY owner;
Maximum value	SELECT MAX(col_name) AS label FROM table;	SELECT MAX(col_name) AS label FROM table;

Following table named CUSTOMERS

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00
5	Hardik	27	Bhopal	8500.00
6	Komal	22	MP	4500.00
7	Muffy	24	Indore	10000.00

Table ORDERS as follows

OID	DATE	CUSTOMER_ID	AMOUNT
102	2009-10-08 00:00:00	3	3000
100	2009-10-08 00:00:00	3	1500
101	2009-11-20 00:00:00	2	1560
103	2008-05-20 00:00:00	4	2060

Default Join is outer Join and query as follows

```
hive> SELECT c.ID, c.NAME, c.AGE, o.AMOUNT FROM  
CUSTOMERS c JOIN ORDERS o ON (c.ID = o.CUSTOMER_ID);
```



Default Join is outer Join and query as follows

```
hive> SELECT c.ID, c.NAME, c.AGE, o.AMOUNT FROM  
CUSTOMERS c JOIN ORDERS o ON (c.ID = o.CUSTOMER_ID);
```

Output on successful completion is

ID	NAME	AGE	AMOUNT
3	kaushik	23	3000
3	kaushik	23	1500
2	Khilan	25	1560
4	Chaitali	25	2060



Left outer Join in Hive as follows

```
hive> SELECT c.ID, c.NAME, o.AMOUNT, o.DATE FROM  
CUSTOMERS c LEFT OUTER JOIN ORDERS o ON (c.ID =  
o.CUSTOMER_ID);
```

Output on successful completion is

ID	NAME	AMOUNT	DATE
1	Ramesh	NULL	NULL
2	Khilan	1560	2009-11-20 00:00:00
3	kaushik	3000	2009-10-08 00:00:00
3	kaushik	1500	2009-10-08 00:00:00
4	Chaitali	2060	2008-05-20 00:00:00
5	Hardik	NULL	NULL
6	Komal	NULL	NULL
7	Muffy	NULL	NULL





Right outer Join in Hive as follows

```
hive> SELECT c.ID, c.NAME, o.AMOUNT, o.DATE FROM  
CUSTOMERS c RIGHT OUTER JOIN ORDERS o ON (c.ID =  
o.CUSTOMER_ID);
```

Output on successful completion is

ID	NAME	AMOUNT	DATE
3	kaushik	3000	2009-10-08 00:00:00
3	kaushik	1500	2009-10-08 00:00:00
2	Khilan	1560	2009-11-20 00:00:00
4	Chaitali	2060	2008-05-20 00:00:00

Please also explore the full outer join on your own



# Introduction to NOSQL

# NOSQL

Not

Only

SQL



- NOSQL – **Not Only SQL** (Structured Query Language)
- Non relational data storage system
- They do not require tabular relations used in relational databases
- Used in big data and real-time web applications
- They may support SQL like language
- *Horizontal scaling* of clusters of machines is achievable using NoSQL
- Data structures used in NoSQL databases are different from those in relational databases making some operations faster in NoSQL


- HBase began as a project by the company Powerset in order to process
- massive amounts of data for the purpose of natural language search.
- Facebook selected HBase to implement its new messaging platform in 2010
- The first HBase release was bundled as part of Hadoop 0.15.0 in October 2007.
- HBase is basically a database ,it's a Hadoop database .

- HBase is often described as a sparse, distributed, persistent, multidimensional sorted map, which is indexed by rowkey, columnkey, and timestamp.
- It's a platform for storing and retrieving data with random access .
- HBase stores structured and semi structured data naturally so you can load it with tweets and parsed log files and a catalogue of all .
- HBase can store unstructured data too, as long as it's not too large. It doesn't care about types and allows for a dynamic and flexible data model that doesn't constrain the kind of data you store.
- It is not a relational database like RDBMS ,hence doesn't include SQL like syntax
- It is designed to run on cluster of computers , it scales horizontally with the addition of new machines.

HDFS	HBase
It is a distributed file storage system	It is a database built on top of HDFS
Doesn't support fast individual record lookups	Provides fast record lookups
Sequential access of data	Uses Hash tables and provide random access of data
Provides high latency batch processing	Provides low latency access to single rows

HBase	RDBMS
It is schema-less ,doesn't have the concept of fixed columns schema.	RDBMS has well defined schema which describes the whole structure of tables
No transactions are there in HBase	RDBMS is transactional
It is built for wide tables and is horizontally scalable	It is built for small tables and is hard to scale
It is good for semi-structured as well as structured data	It is good for structured data

**Sqoop = SQL-to-Hadoop**





- JDBC-based interface (MySQL, Oracle, PostgreSQL, etc...)
- Automatic Datatype generation
  - Reads column info from table and generates Java classes
  - Can be used in further MapReduce processing passes
- Uses MapReduce to read tables from database
  - Can select individual table (or subset of columns)
  - Can read all tables in database
- Supports most JDBC standard types and null values

# Example of Sqoop

```
mysql> use corp;  
Database changed
```

```
mysql> describe employees;
```

Field	Type	Null	Key	Default	Extra
id	int(11)	NO	PRI	NULL	auto_increment
firstname	varchar(32)	YES		NULL	
lastname	varchar(32)	YES		NULL	
jobtitle	varchar(64)	YES		NULL	
start_date	date	YES		NULL	
dept_id	int(11)	YES		NULL	

# Import table in HDFS

```
sqoop --connect  
jdbc:mysql://db.foo.com/corp \  
--table employees
```

- Imports “employees” table into HDFS directory

# Sqoop Hive Integration

```
$ sqoop --connect jdbc:mysql://db.foo.com/ corp --hive-import --  
table employees
```

- Auto-generates CREATE TABLE / LOAD/ INPATH statements for Hive
- After data is imported to HDFS, auto-executes Hive script
- Follow-up step: Loading into partitions

# Appendix

- Programming Hive : O'Reilly
- <https://hive.apache.org/>
- <http://hortonworks.com/apache/hive/>
- <https://cwiki.apache.org/confluence/display/Hive/Home>
- <http://www.cloudera.com/products/apache-hadoop/apache-hive.html>

